

# TRASH TALK



EXPLORING THE MEMORY MANAGEMENT IN THE JVM

# ABOUT ME.



Gerrit Grunwald | Developer Advocate | Azul

# MEMORY MANAGEMENT

IN THE JVM...

IS  
AUTOMATIC  
RIGHT...?

SO...WHY  
CARE...?

# MEMORY MANAGEMENT

Why you should care...

🗑️ Impact on application performance

# MEMORY MANAGEMENT

Why you should care...

- 🗑 Impact on application performance
- 🗑 Impact on application responsiveness



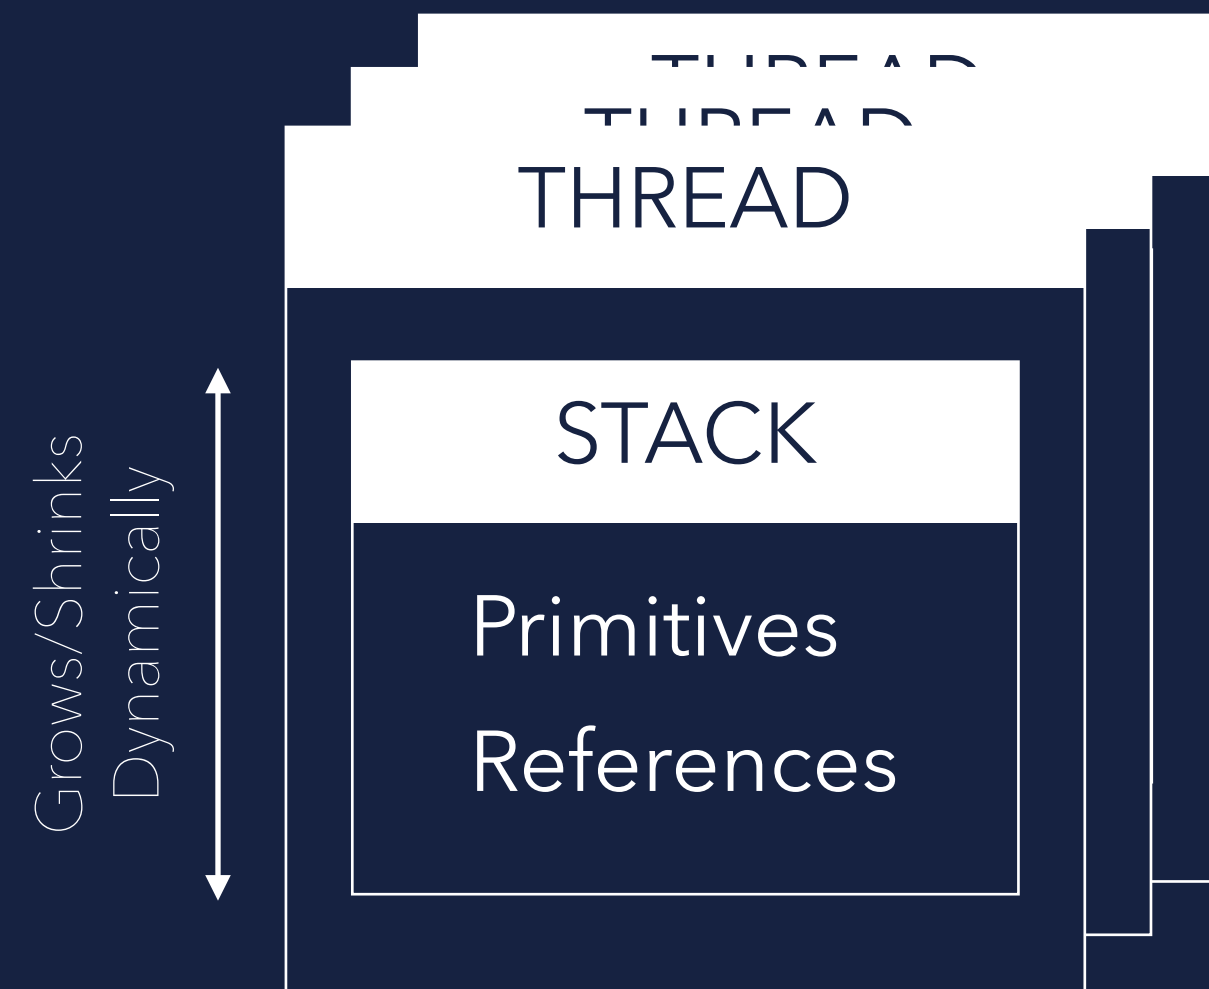
# MEMORY MANAGEMENT

Why you should care...

- 🗑 Impact on application performance
- 🗑 Impact on application responsiveness
- 🗑 Impact on system requirements

# MEMORY MANAGEMENT

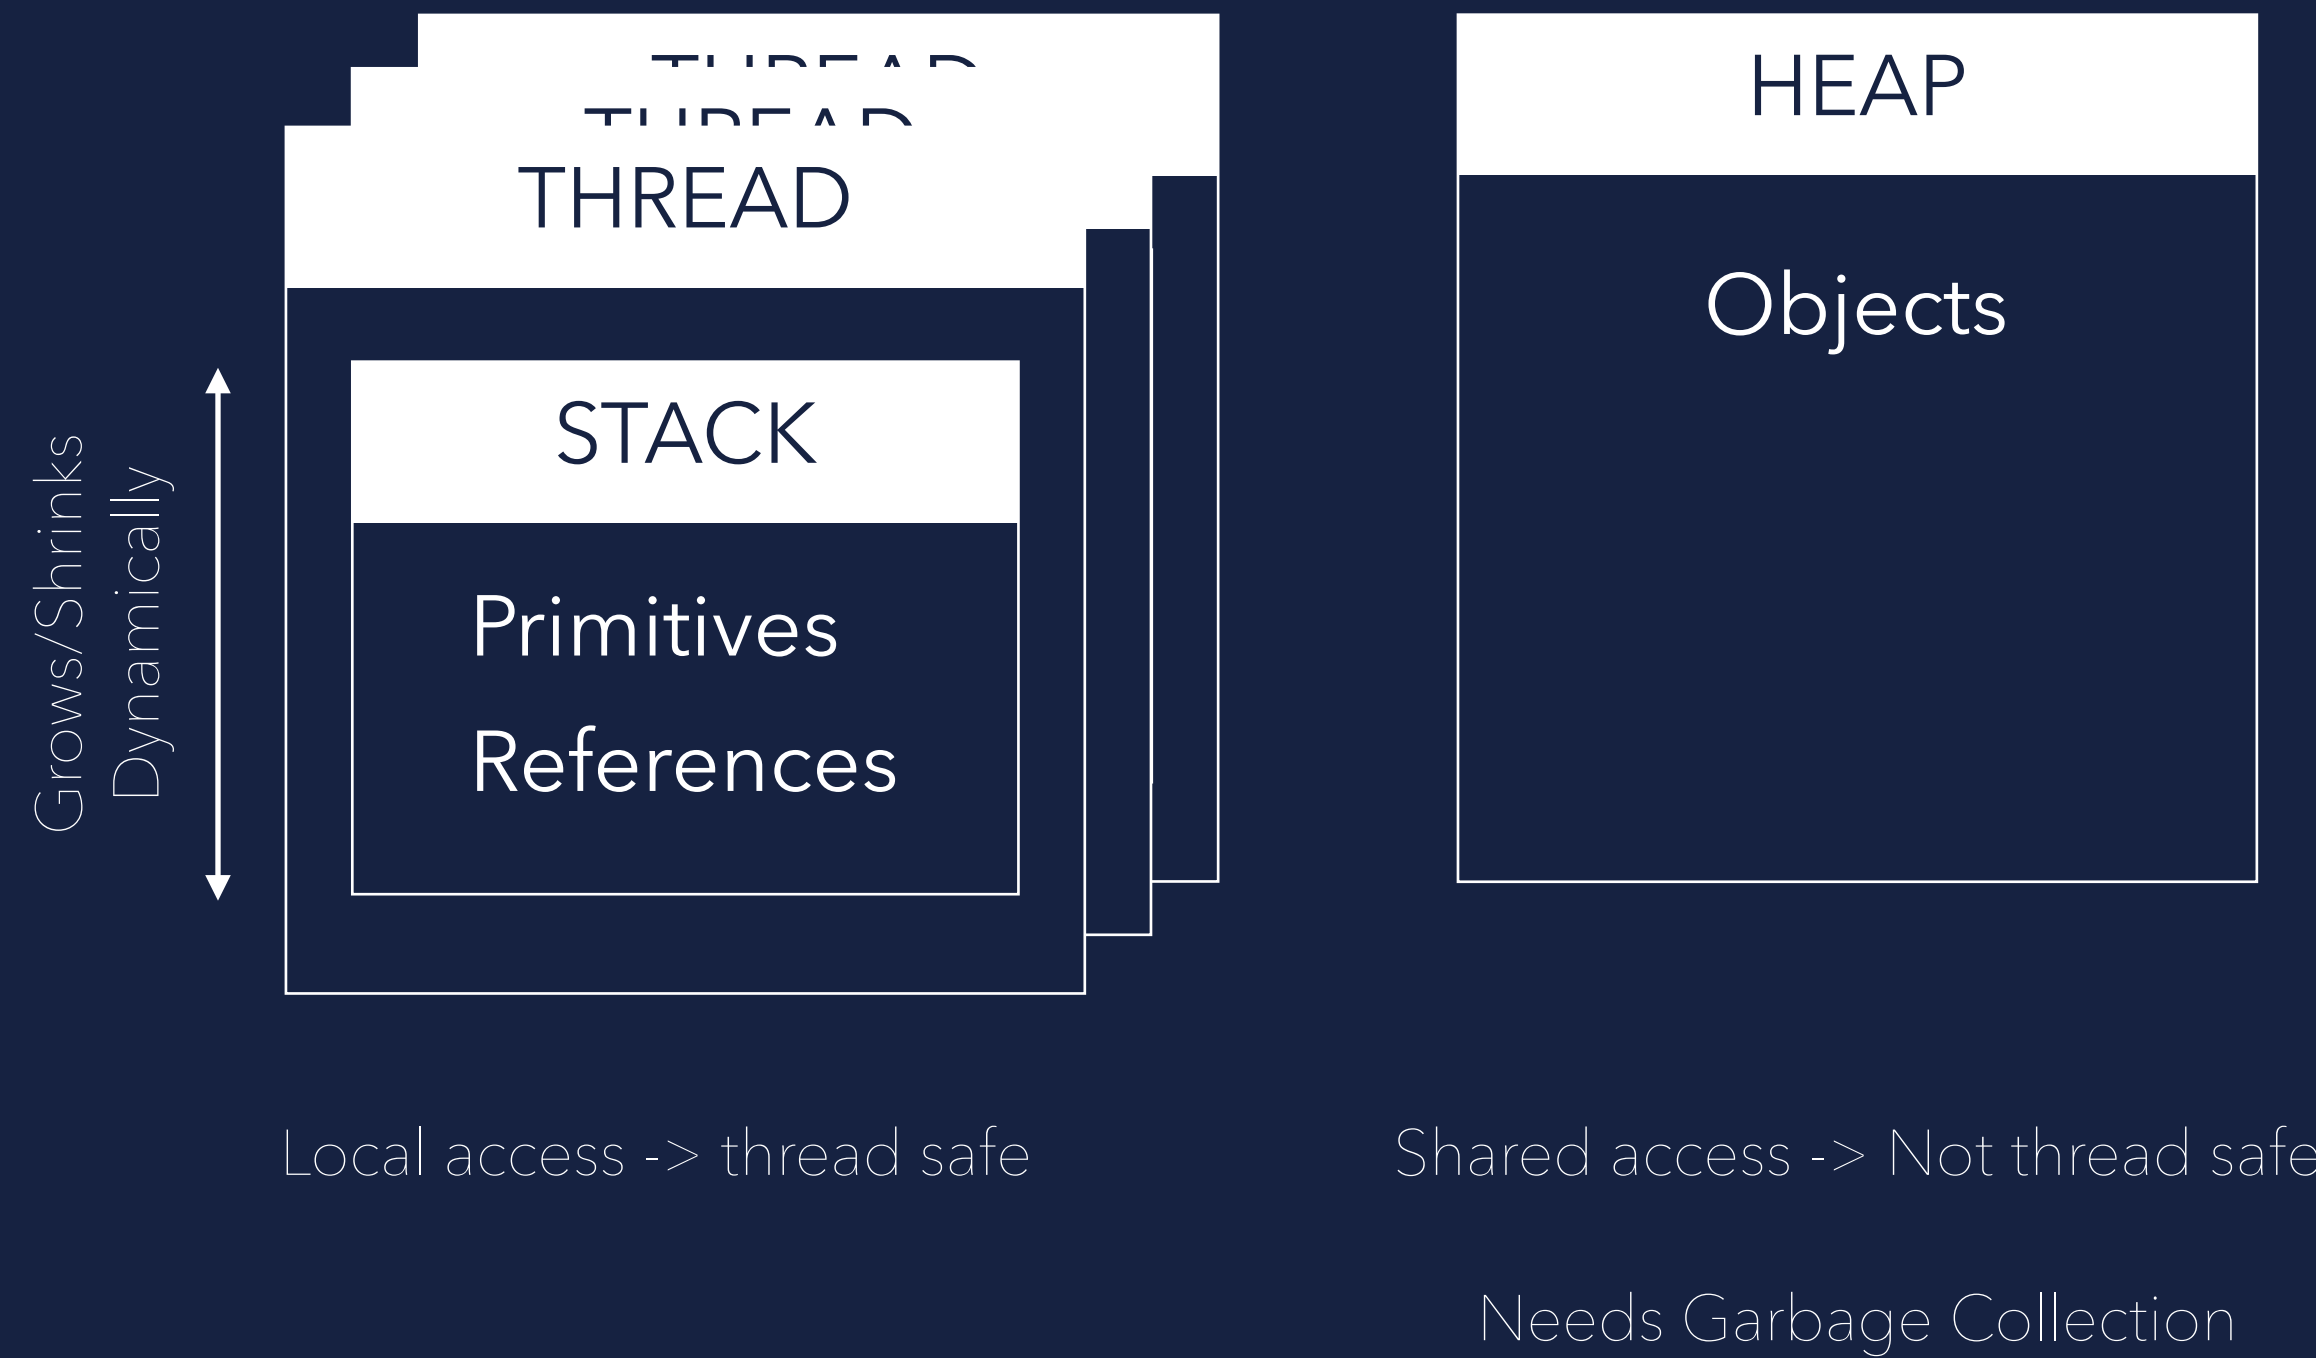
Stack, Heap and Metaspace



Local access -> thread safe


# MEMORY MANAGEMENT

## Stack, Heap and Metaspace



# MEMORY MANAGEMENT

## Stack, Heap and Metaspace



# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
}
```

# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
}
```

Stack area for thread 1

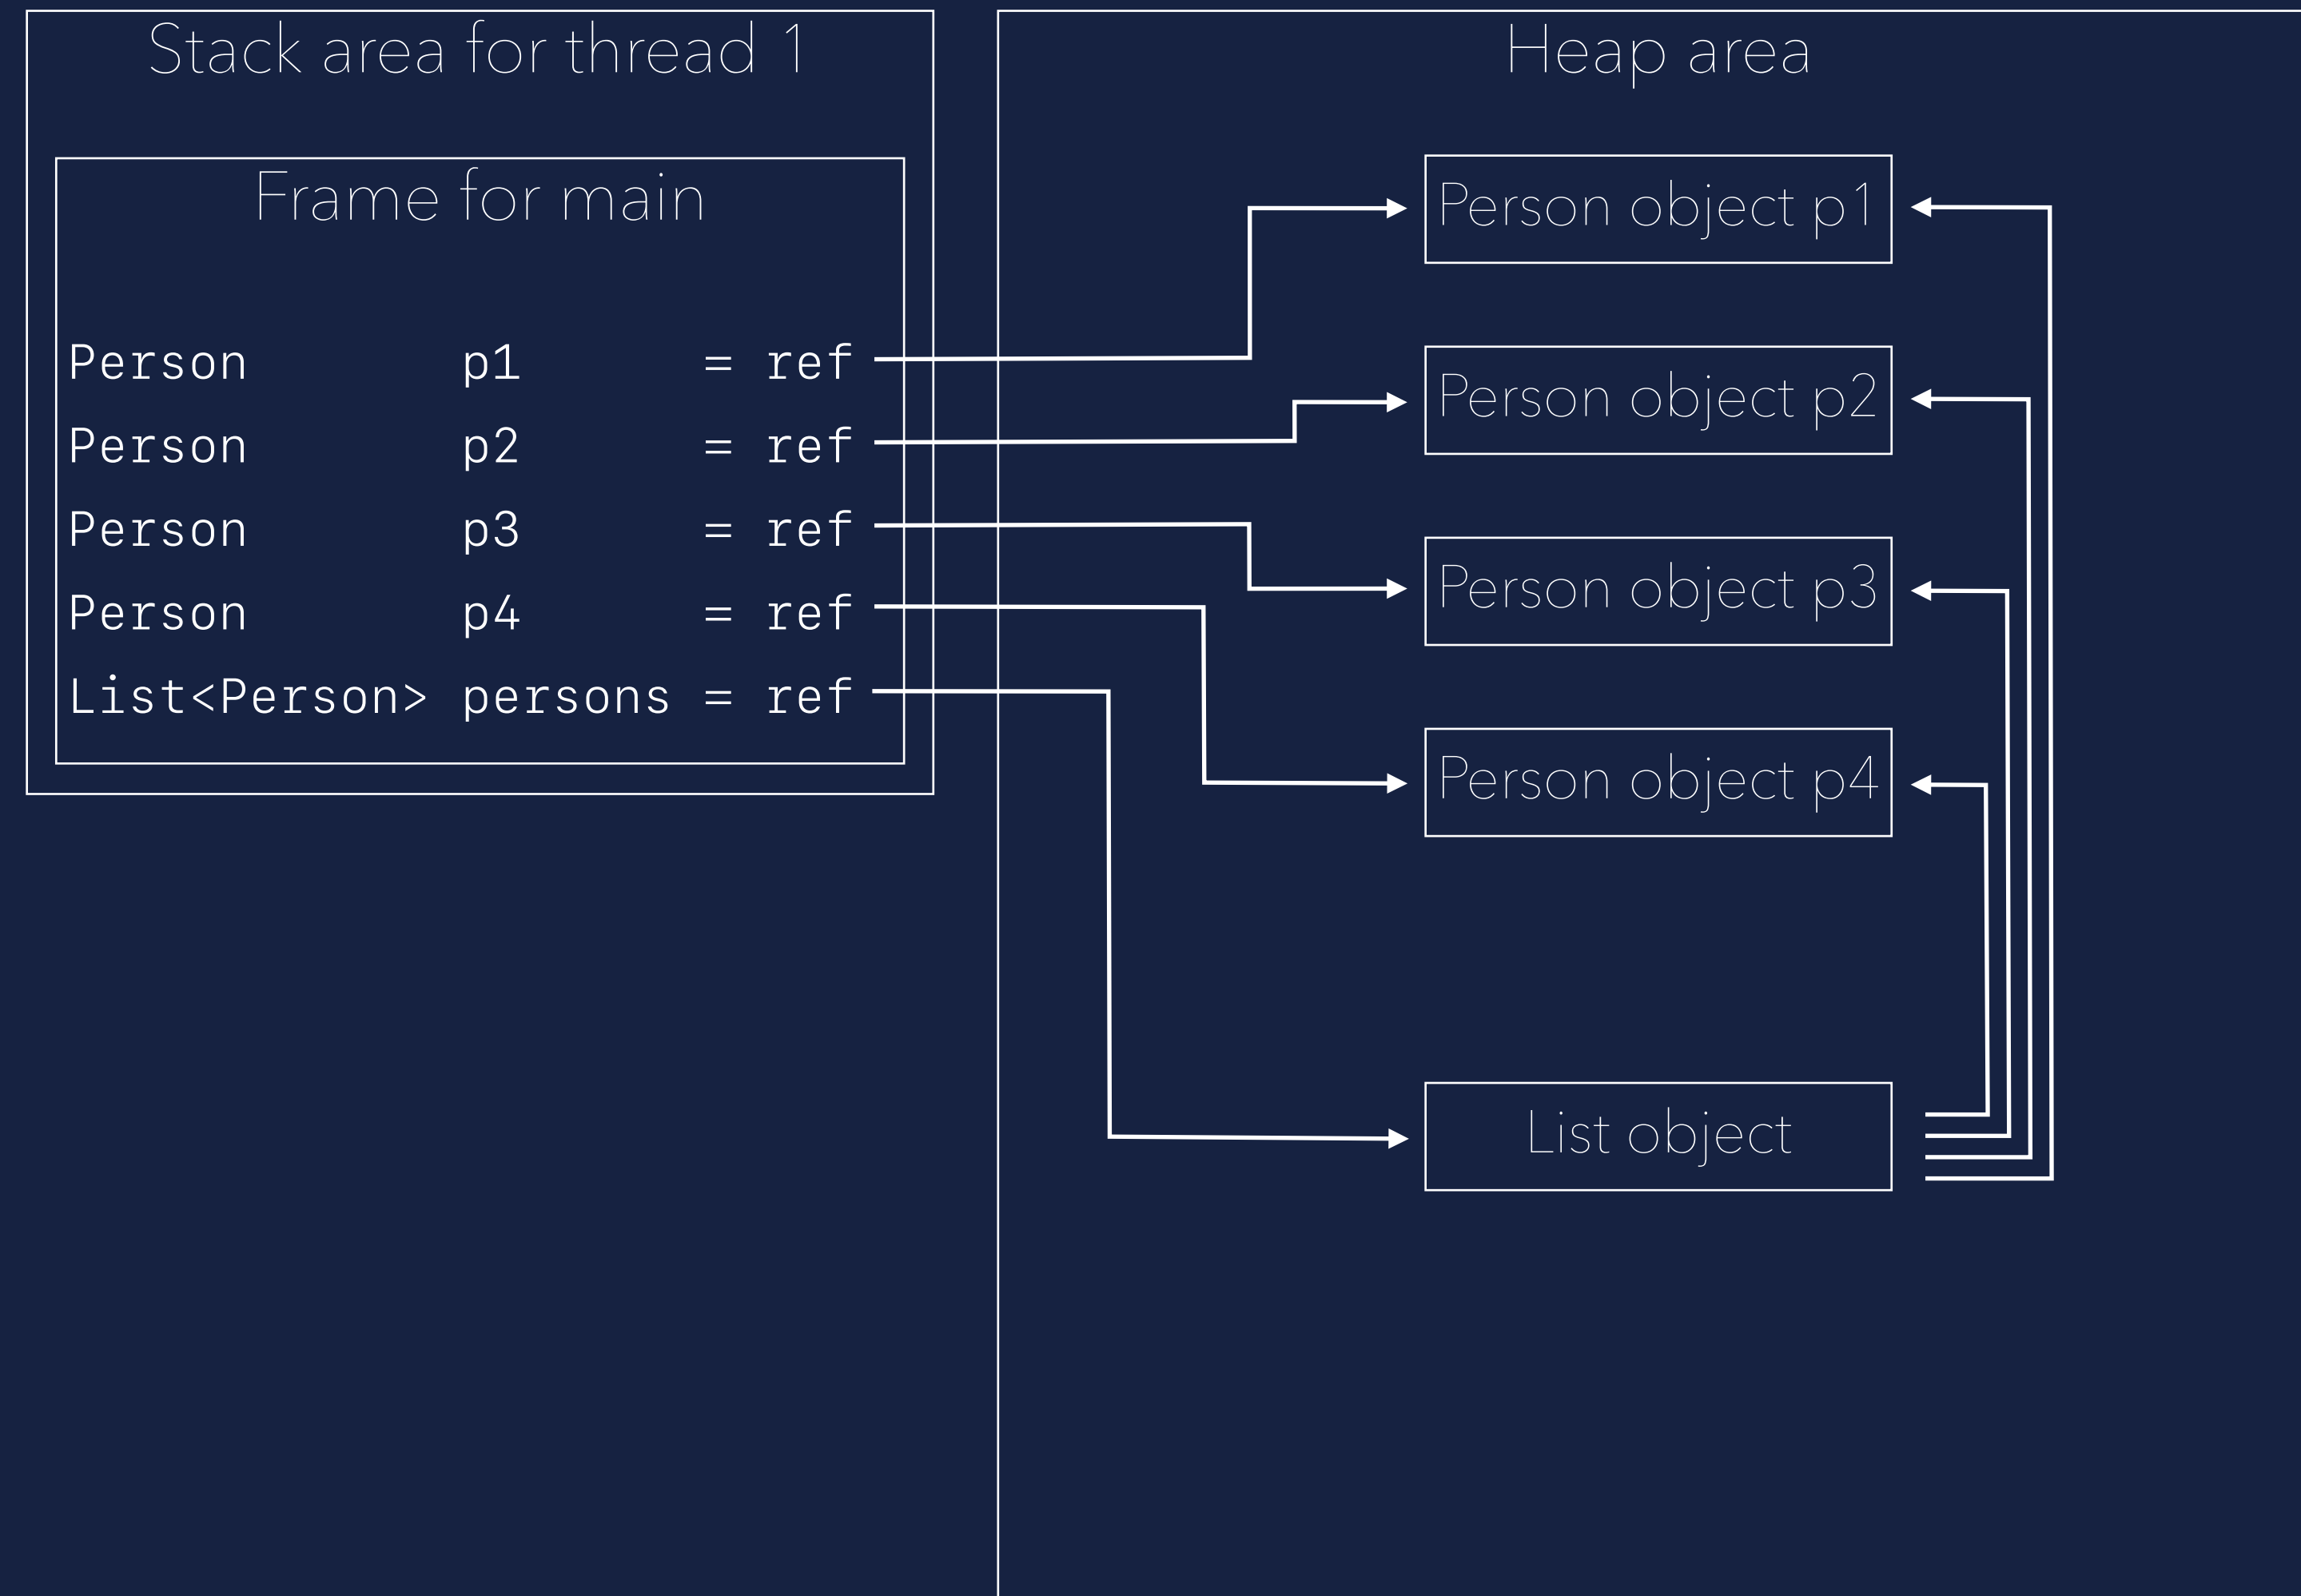
Frame for main

Person	p1	= ref
Person	p2	= ref
Person	p3	= ref
Person	p4	= ref
List<Person>	persons	= ref

# MEMORY MANAGEMENT

In the JVM...

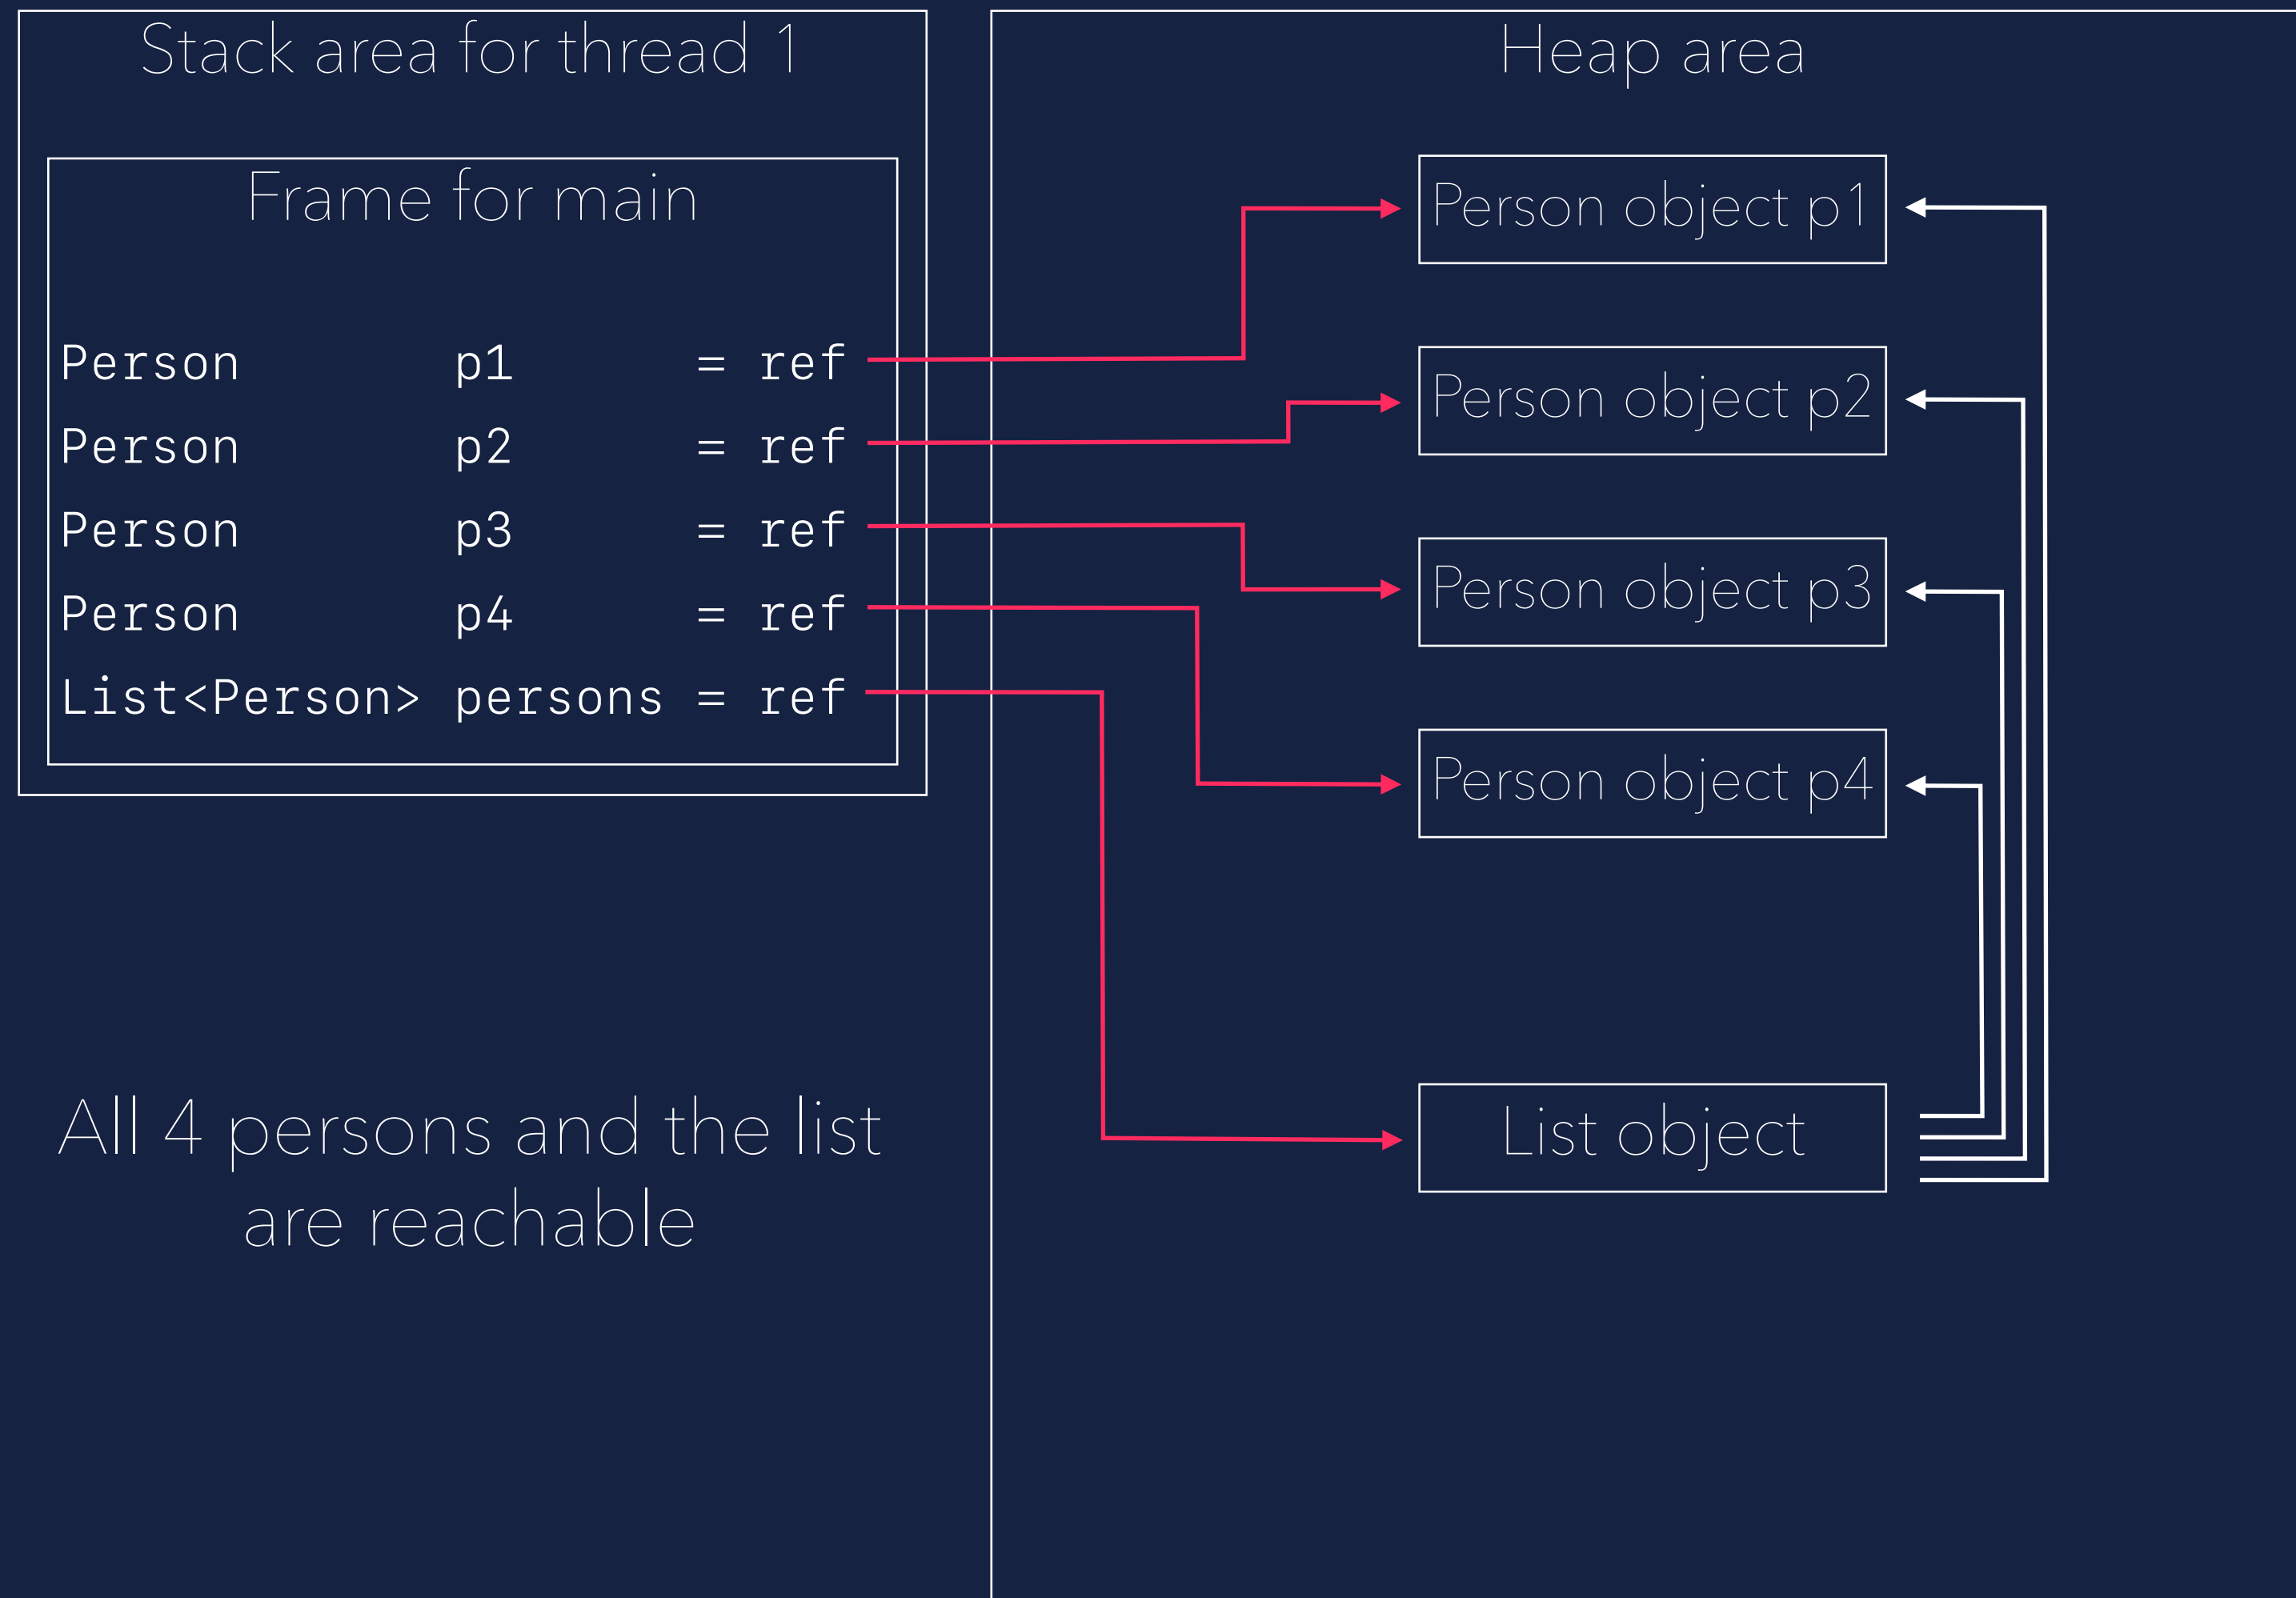
```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
}
```



# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
}
```





# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
}
```

Stack area for thread 1

Frame for main

Person	p1	= null
Person	p2	= ref
Person	p3	= ref
Person	p4	= ref
List<Person>	persons	= ref

Setting p1 = null

Heap area

Person object p1

Person object p2

Person object p3

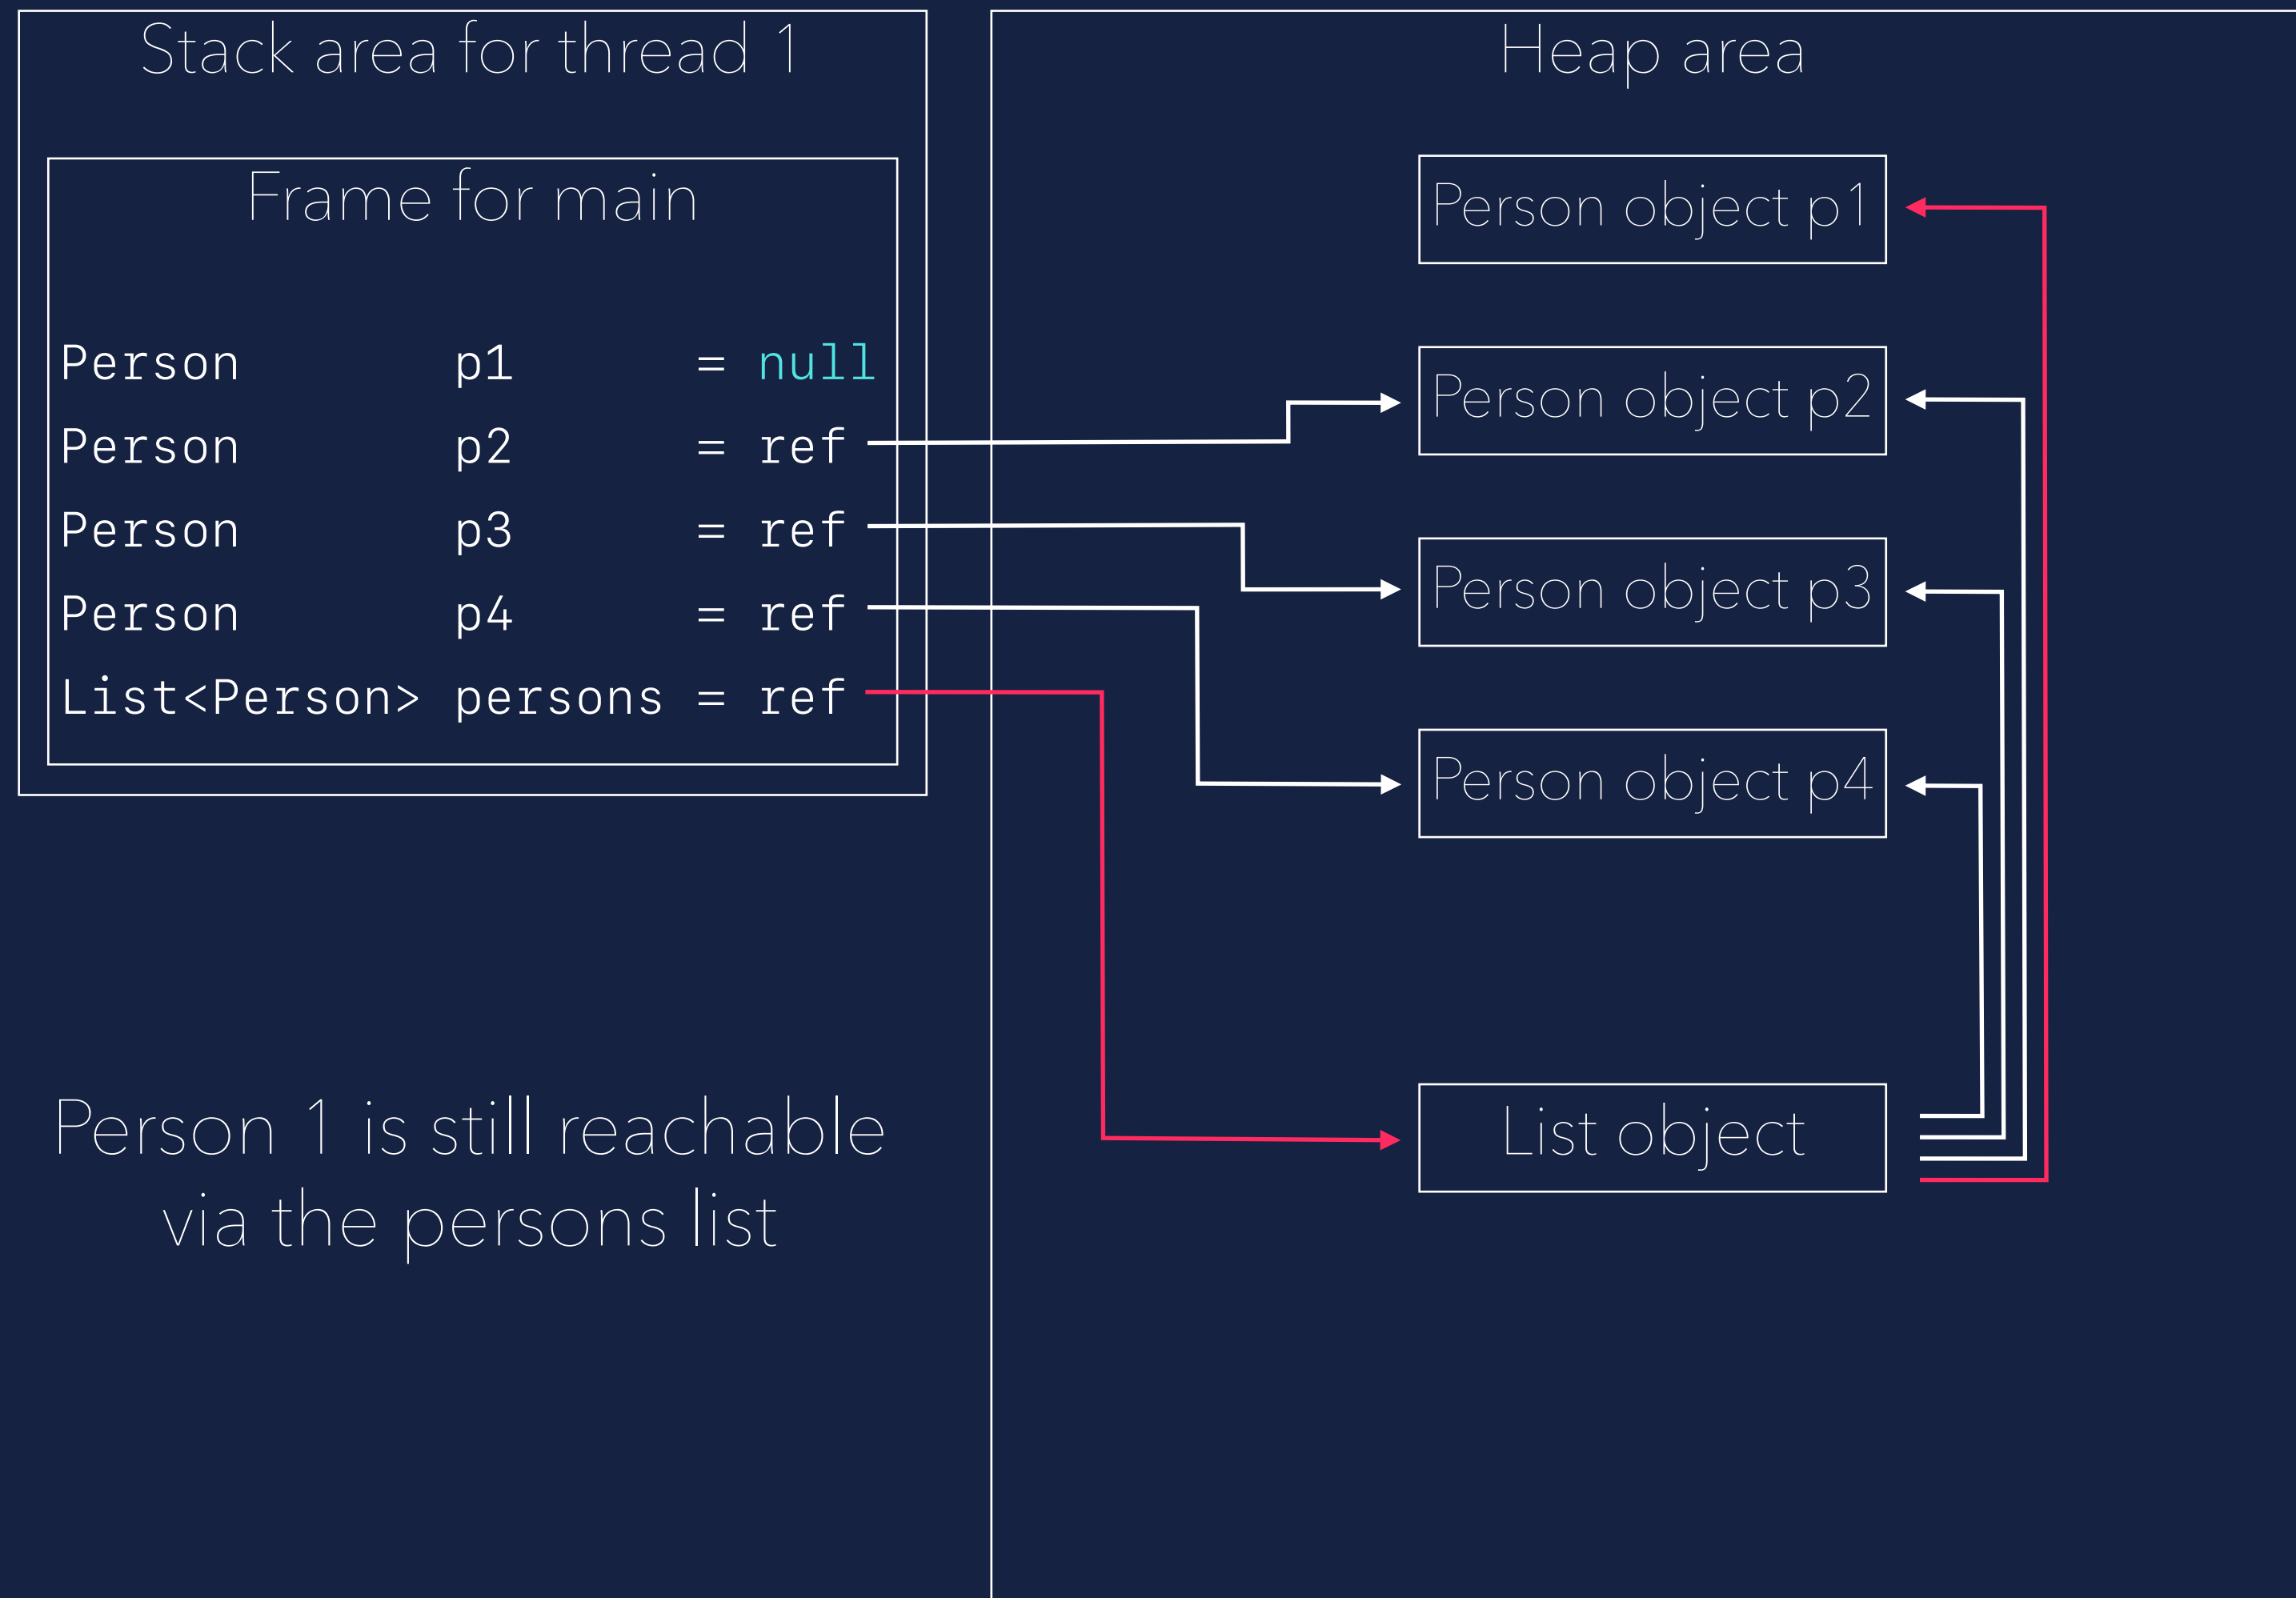
Person object p4

List object

# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
}
```



# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
  
    persons = null;  
}
```

Stack area for thread 1

Frame for main

Person	p1	= null
Person	p2	= ref
Person	p3	= ref
Person	p4	= ref
List<Person>	persons	= null

Setting persons = null

Heap area

Person object p1

Person object p2

Person object p3

Person object p4

List object

# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
  
    persons = null;  
}
```

Stack area for thread 1

Frame for main

Person	p1	= null
Person	p2	= ref
Person	p3	= ref
Person	p4	= ref
List<Person>	persons	= null

Only p2, p3 and p4 are  
reachable

Heap area

Person object p1

Person object p2

Person object p3

Person object p4

List object

# MEMORY MANAGEMENT

In the JVM...

```
public static void main(String[] args) {  
  
    record Person(String name) {  
        @Override public String toString() { return name(); }  
    }  
  
    Person p1 = new Person("Gerrit");  
    Person p2 = new Person("Sandra");  
    Person p3 = new Person("Lilli");  
    Person p4 = new Person("Anton");  
  
    List<Person> persons = Arrays.asList(p1, p2, p3, p4);  
  
    System.out.println(p1); // -> Gerrit  
  
    p1 = null;  
  
    System.out.println(persons.get(0)); // -> Gerrit  
  
    persons = null;  
}
```

Stack area for thread 1

Frame for main

Person	p1	= null
Person	p2	= ref
Person	p3	= ref
Person	p4	= ref
List<Person>	persons	= null

p1 and persons are  
garbage

Heap area

Person object p1

Person object p2

Person object p3

Person object p4

List object

HOW TO GET  
RID OF IT...?

# **GARBAGE COLLECTION**

# GARBAGE COLLECTION

What is it...

 Form of automatic memory management



# GARBAGE COLLECTION

What is it...

- 🗑 Form of automatic memory management
- 🗑 Identifies and reclaims no longer used memory

# GARBAGE COLLECTION

What is it...

- 🗑 Form of automatic memory management
- 🗑 Identifies and reclaims no longer used memory
- 🗑 Ensures efficient memory utilisation

# GARBAGE COLLECTION

What is it...

- 🗑 Form of automatic memory management
- 🗑 Identifies and reclaims no longer used memory
- 🗑 Ensures efficient memory utilisation
- 🗑 Frees user from managing the memory manually

**PHASES**

# GARBAGE COLLECTION

Phases (precise collectors)



## Tracing

Identify live objects on the heap

# GARBAGE COLLECTION

## Phases (precise collectors)



### Tracing

Identify live objects on the heap



### Freeing

Reclaim resources held by dead objects

# GARBAGE COLLECTION

## Phases (precise collectors)

 Tracing  
Identify live objects on the heap

 Freeing  
Reclaim resources held by dead objects


 Compaction  
Periodically relocate live objects

# STOPPING THE WORLD



# STOPPING THE WORLD

Halt of all application threads



# COLLECTORS


# NON MOVING COLLECTOR

Mark & Sweep

# NON MOVING COLLECTOR

## Demo

- 1. Mutator allocates cells in Heap
- 2. Heap is out of memory -> GC
- 3. Mark all live cells
- 4. Free all dead cells
- 5. Unmark all live cells
- 6. Resume Mutator



Heap



- Free Cell
- Referenced Cell
- Dereferenced Cell
- Marked Cell
- Referenced Cell (survived 1 GC)



Fragmentation

# MOVING COLLECTOR

Compacting Collector & Copy Collector


# COMPACTING COLLECTOR

Mark & Compact

# COMPACTING COLLECTOR

## Demo

- 1. Mutator allocates cells in Heap
- 2. Heap is out of memory -> GC
- 3. Mark all live cells
- 4. Free all dead cells
- 5. Unmark all live cells
- 6. Compact all live cells
- 7. Resume Mutator



⚠  
Headroom  
20-50%

# COPY COLLECTOR

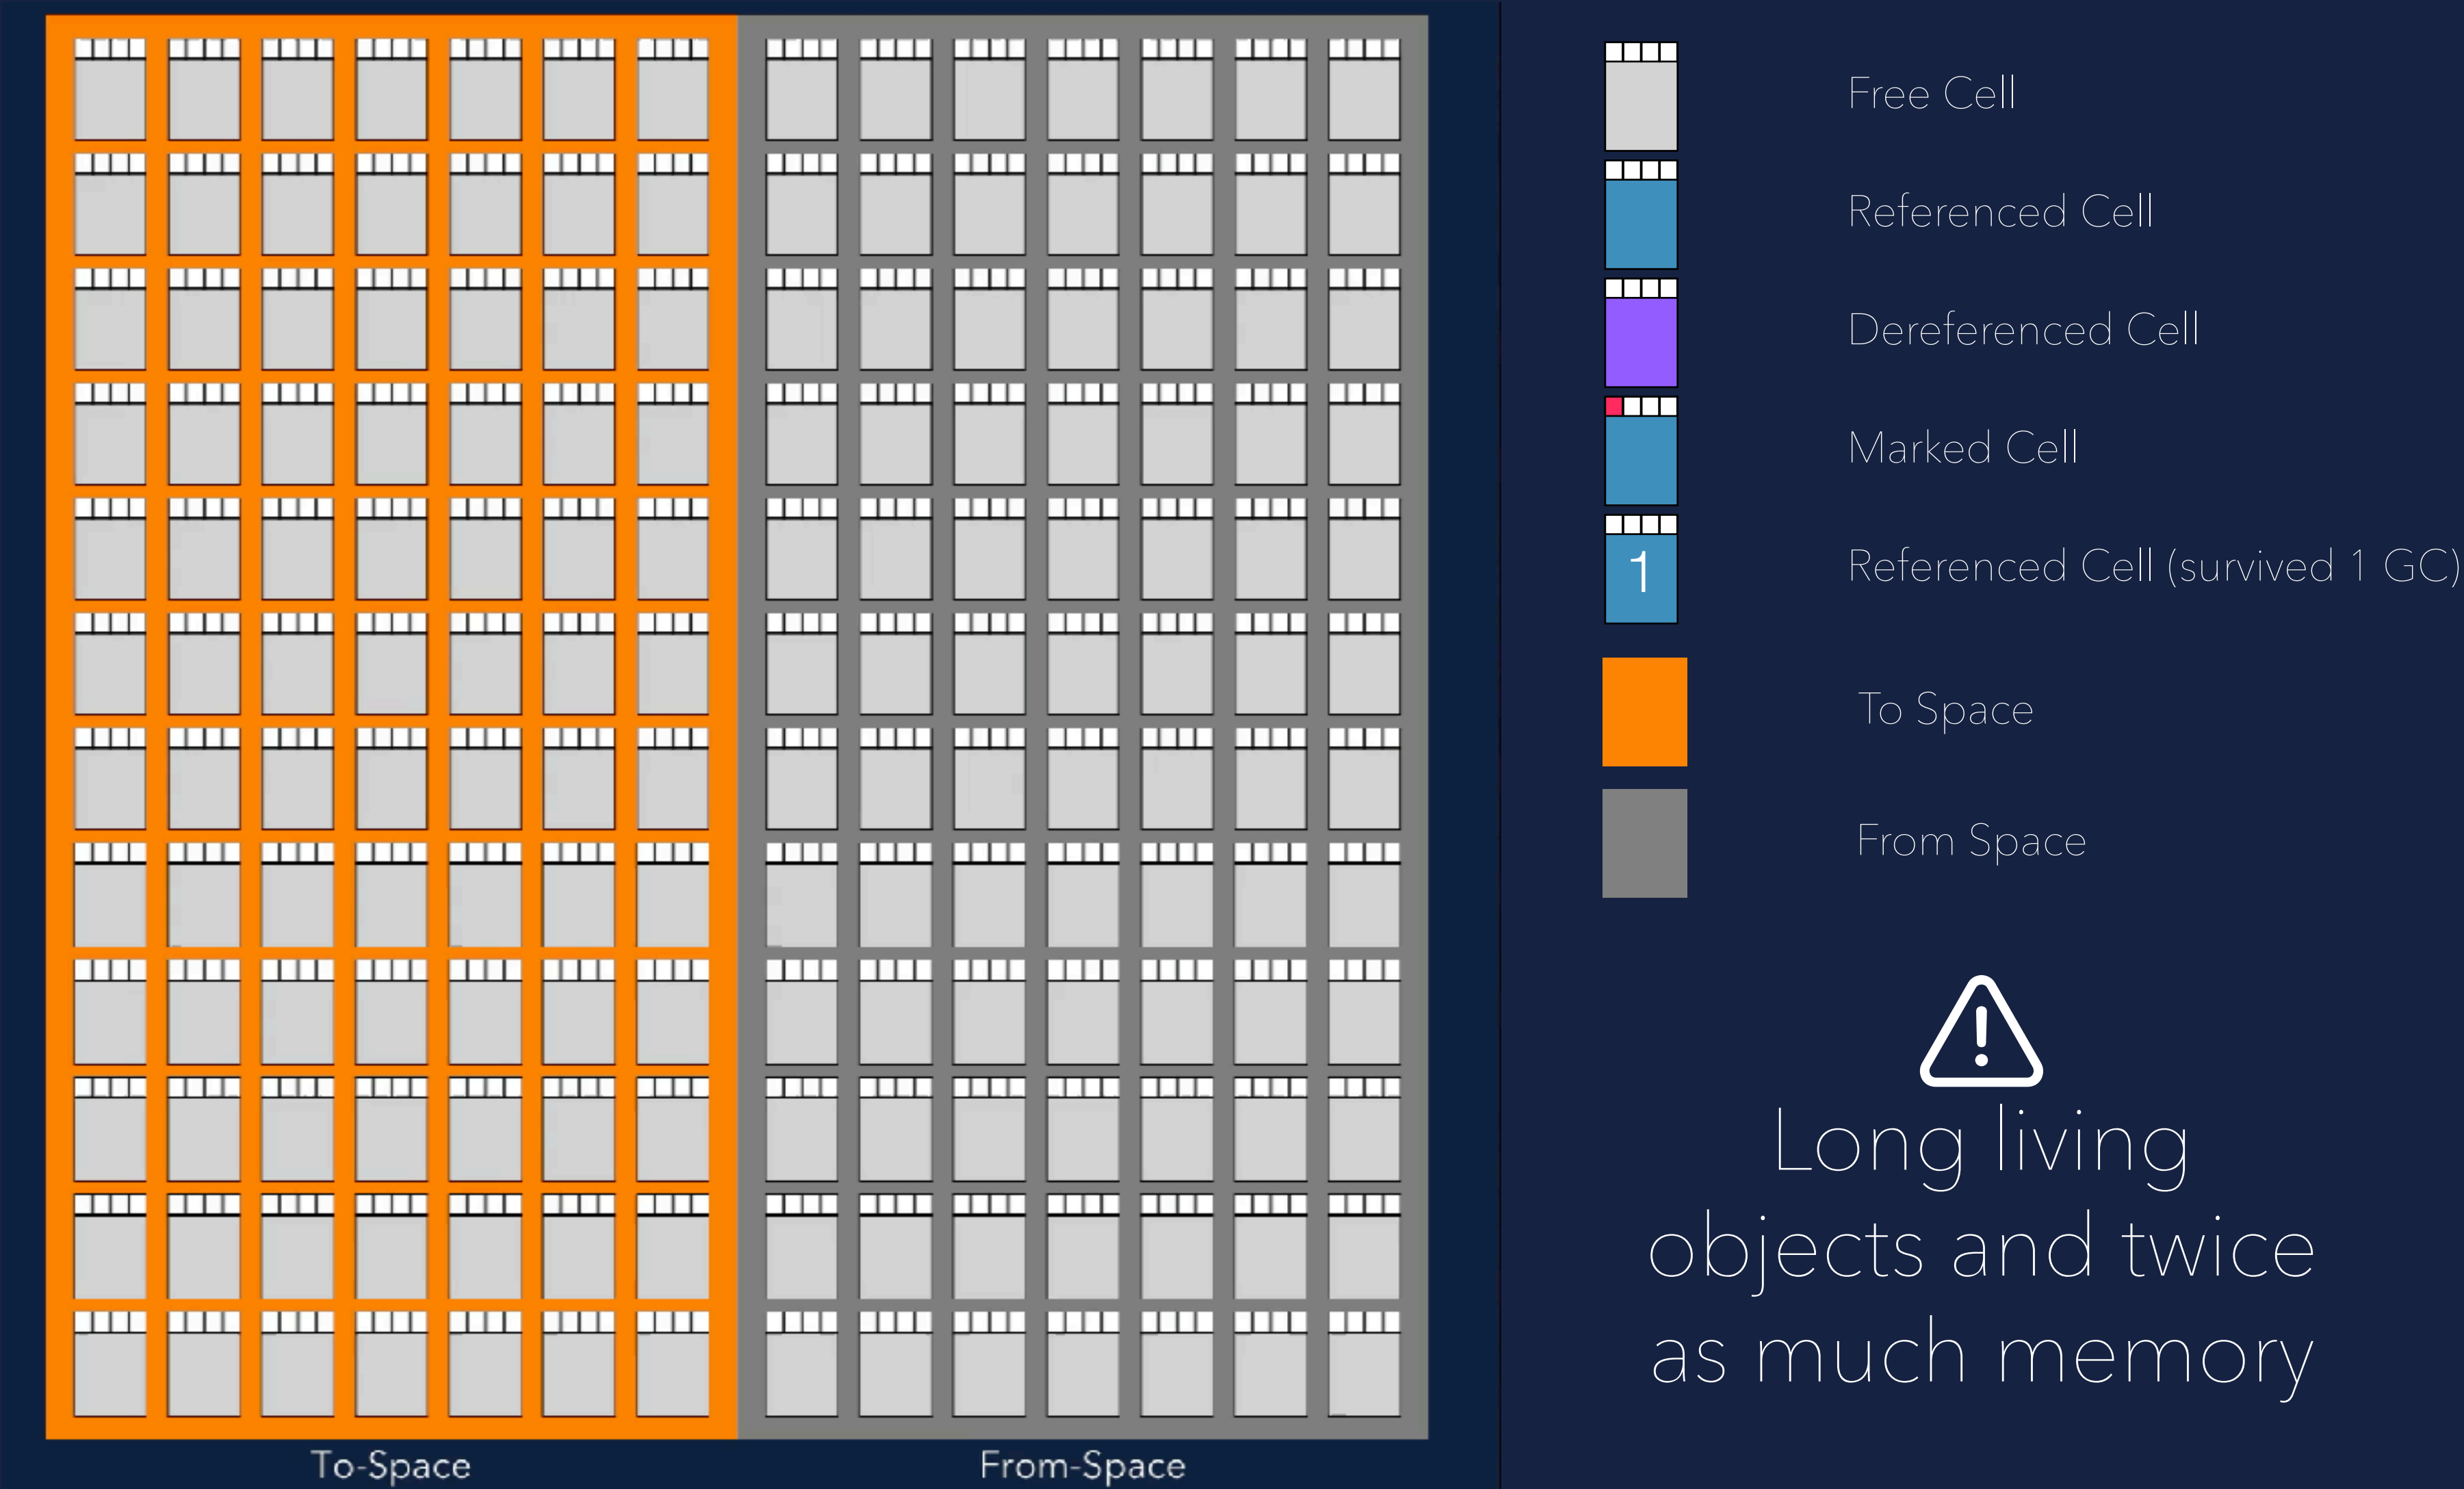
Mark & Copy



# COPY COLLECTOR

## Demo

- 1. Allocating in ToSpace
- 2. ToSpace is out of memory -> GC
- 3. Toggle To- and FromSpace
- 4. Mark live cells in FromSpace
- 5. Copy live cells to ToSpace
- 6. Free all cells in FromSpace
- 7. Resume Mutator




# GENERATIONAL COLLECTOR

Generational Mark & Compact


# GENERATIONAL COLLECTOR

Weak Generational Hypothesis (Most objects die young)



# GENERATIONAL COLLECTOR

Weak Generational Hypothesis (Most objects die young)



# GENERATIONAL COLLECTOR

## Demo

- 1. Mutator allocates cells in Eden
- 2. Eden is out of memory -> GC
- 3. Toggle To- and FromSpace
- 4. Copy all live cells from FromSpace to ToSpace
- 5. Copy all live cells from Eden to ToSpace
- 6. Promote live cells from FromSpace to TenuredSpace
- 7. Free all dead cells
- 8. Resume Mutator





# GENERATIONAL COLLECTOR


How to do a minor collection  
with references from old to  
young generation...?


# GENERATIONAL COLLECTOR


Remembered Set (Card Table)



Free Cell

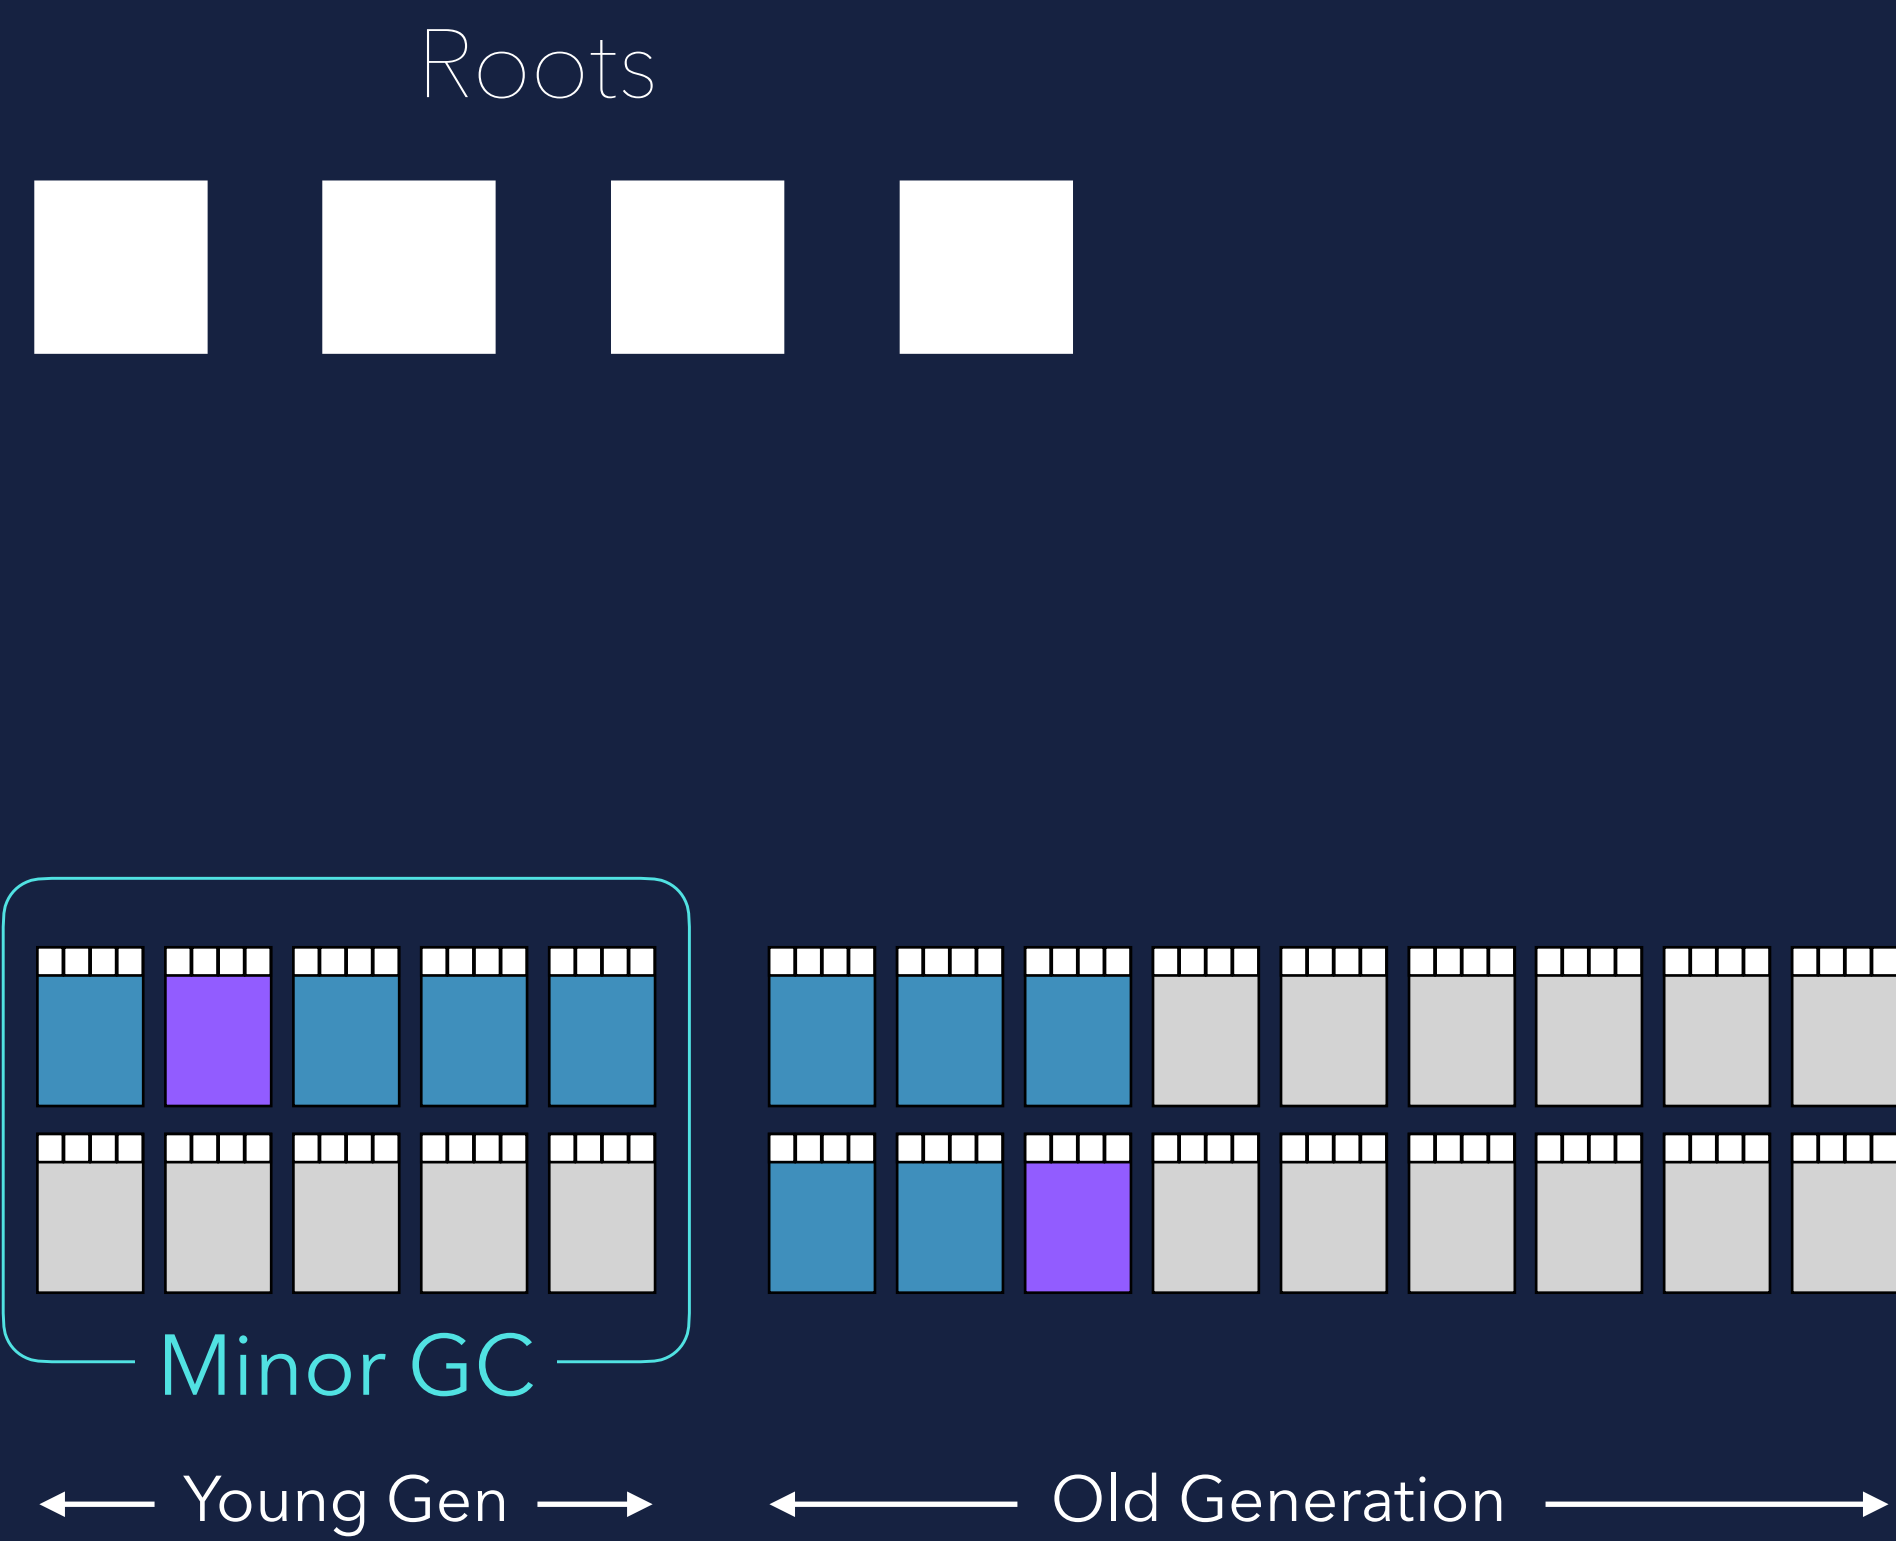
Referenced Cell

Dereferenced Cell

Marked Cell

# GENERATIONAL COLLECTOR

Remembered Set (Card Table)



Free Cell

Referenced Cell


Dereferenced Cell

Marked Cell




# GENERATIONAL COLLECTOR

Remembered Set (Card Table)




# GENERATIONAL COLLECTOR

Remembered Set (Card Table)




# GENERATIONAL COLLECTOR

Remembered Set (Card Table)



# GENERATIONAL COLLECTOR

## Remembered Set (Card Table)



GC looks up Card Table,  
finds the reference and  
marks it as live

# CONCURRENT COLLECTION ?

**CONCURRENCY**

**IS**

**HARD...**

**STOP THE  
WORLD**

# CONCURRENT COLLECTION

Application and GC running concurrently

Application  
Threads




GC Thread(s)



# CONCURRENT MARKING

# CONCURRENCY IS HARD...


## Concurrent Marking






# CONCURRENCY IS HARD...

## Concurrent Marking



# CONCURRENCY IS HARD...


## Concurrent Marking



Mutator removes reference and creates a new one from an already visited cell !

# CONCURRENCY IS HARD...


## Concurrent Marking



Mutator removes reference and creates a new one from an already visited cell !

# CONCURRENCY IS HARD...


## Concurrent Marking



Won't be detected by the Garbage Collector !

# CONCURRENCY IS HARD...

# Concurrent Marking



Won't be detected by the Garbage Collector !



???

# BARRIERS TO THE RESCUE

# BARRIERS

## Read / Write Barriers

- 🗑 Mechanisms to execute memory management code when a read/write on some object takes place

# BARRIERS

## Read / Write Barriers

- 🗑 Mechanisms to execute memory management code when a read/write on some object takes place
- 🗑 Used to keep track of inter-generational references.  
(references from old generation to young generation, the so called Remembered Set)





# BARRIERS

## Read / Write Barriers

- 🗑 Mechanisms to execute memory management code when a read/write on some object takes place
- 🗑 Used to keep track of inter-generational references.  
(references from old generation to young generation, the so called Remembered Set)
- 🗑 Used to synchronize action between mutator and collector  
(allocation concurrent to collection)


# BARRIERS

## Read / Write Barriers

-  Mechanisms to execute memory management code when a read/write on some object takes place
-  Used to keep track of inter-generational references.  
(references from old generation to young generation, the so called Remembered Set)
-  Used to synchronize action between mutator and collector  
(allocation concurrent to collection)
-  Read Barriers are usually more expensive  
(reads 75% to writes 25% -> Read Barriers must very efficient)


# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers



# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers




## Collector starts marking objects




# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers



# CONCURRENCY IS HARD...


## Concurrent Marking using Write Barriers



Mutator hits write barrier and removes reference and adds a new one

# CONCURRENCY IS HARD...


## Concurrent Marking using Write Barriers



Mutator hits write barrier and removes reference and adds a new one


# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers




# CONCURRENCY IS HARD...

## Concurrent Marking using Write Barriers




# CONCURRENCY IS HARD...

Concurrent Marking using Write Barriers




# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers



# CONCURRENCY IS HARD...


# Concurrent Marking using Write Barriers






# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers



# CONCURRENCY IS HARD...


## Concurrent Marking using Write Barriers



In Re-Mark phase, in between marked references will be marked as live


# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers



# CONCURRENCY IS HARD...

# Concurrent Marking using Write Barriers




"Snapshot at the beginning"

# CONCURRENT COPYING

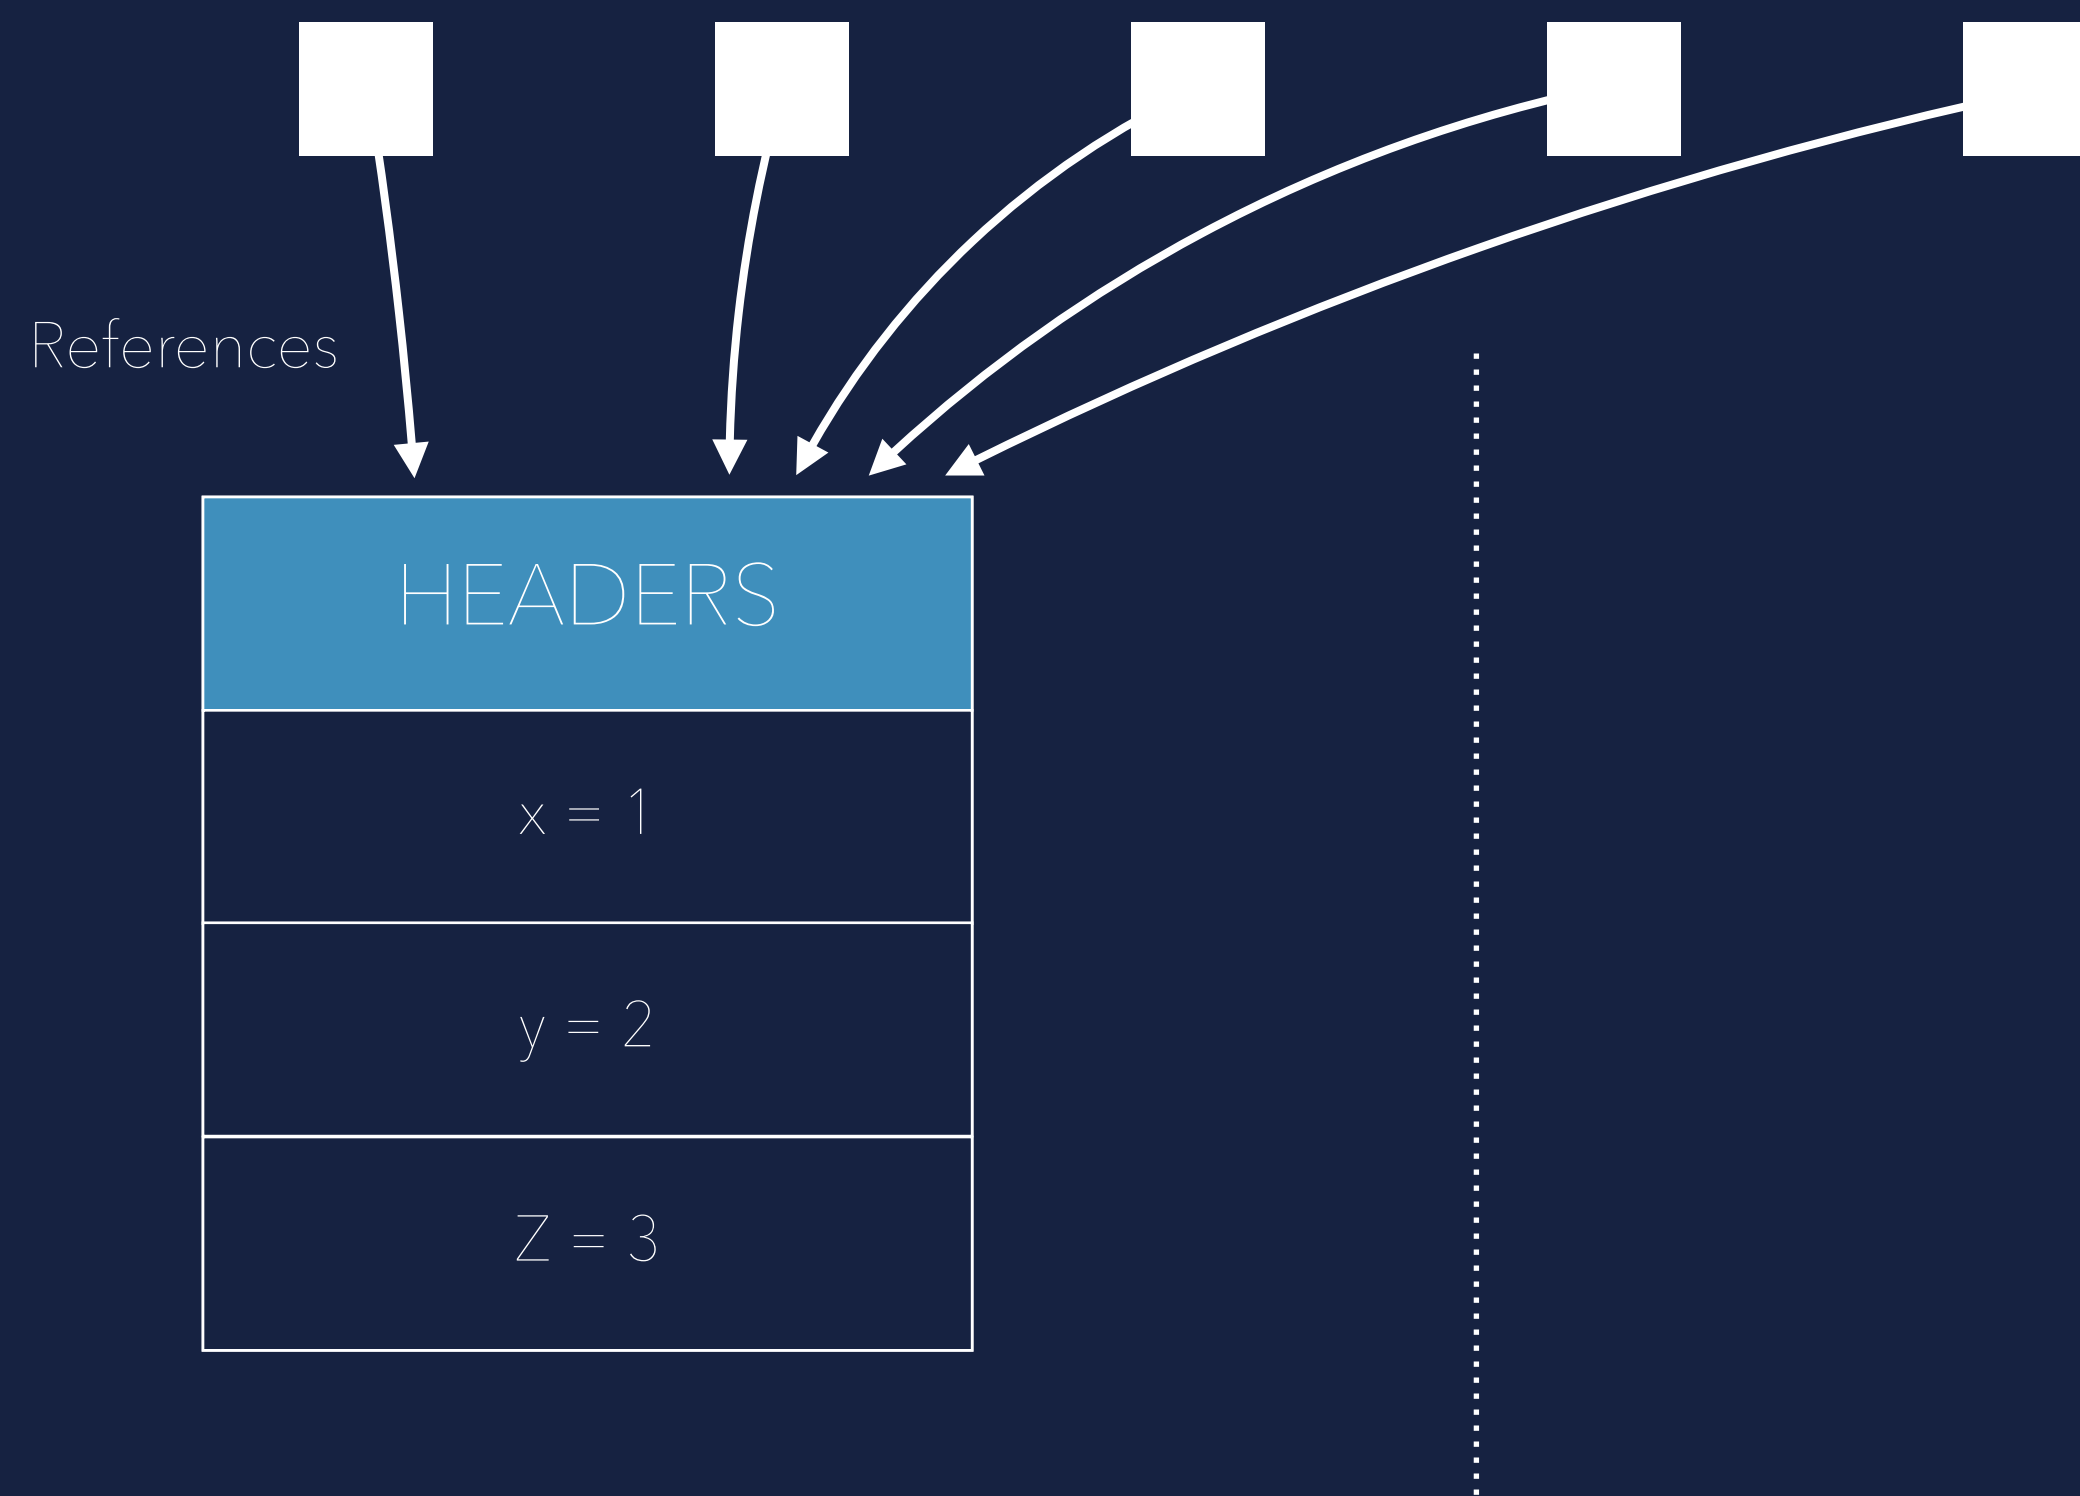
# CONCURRENCY IS HARD...

Stop the world copying



# CONCURRENCY IS HARD...

Stop the world copying



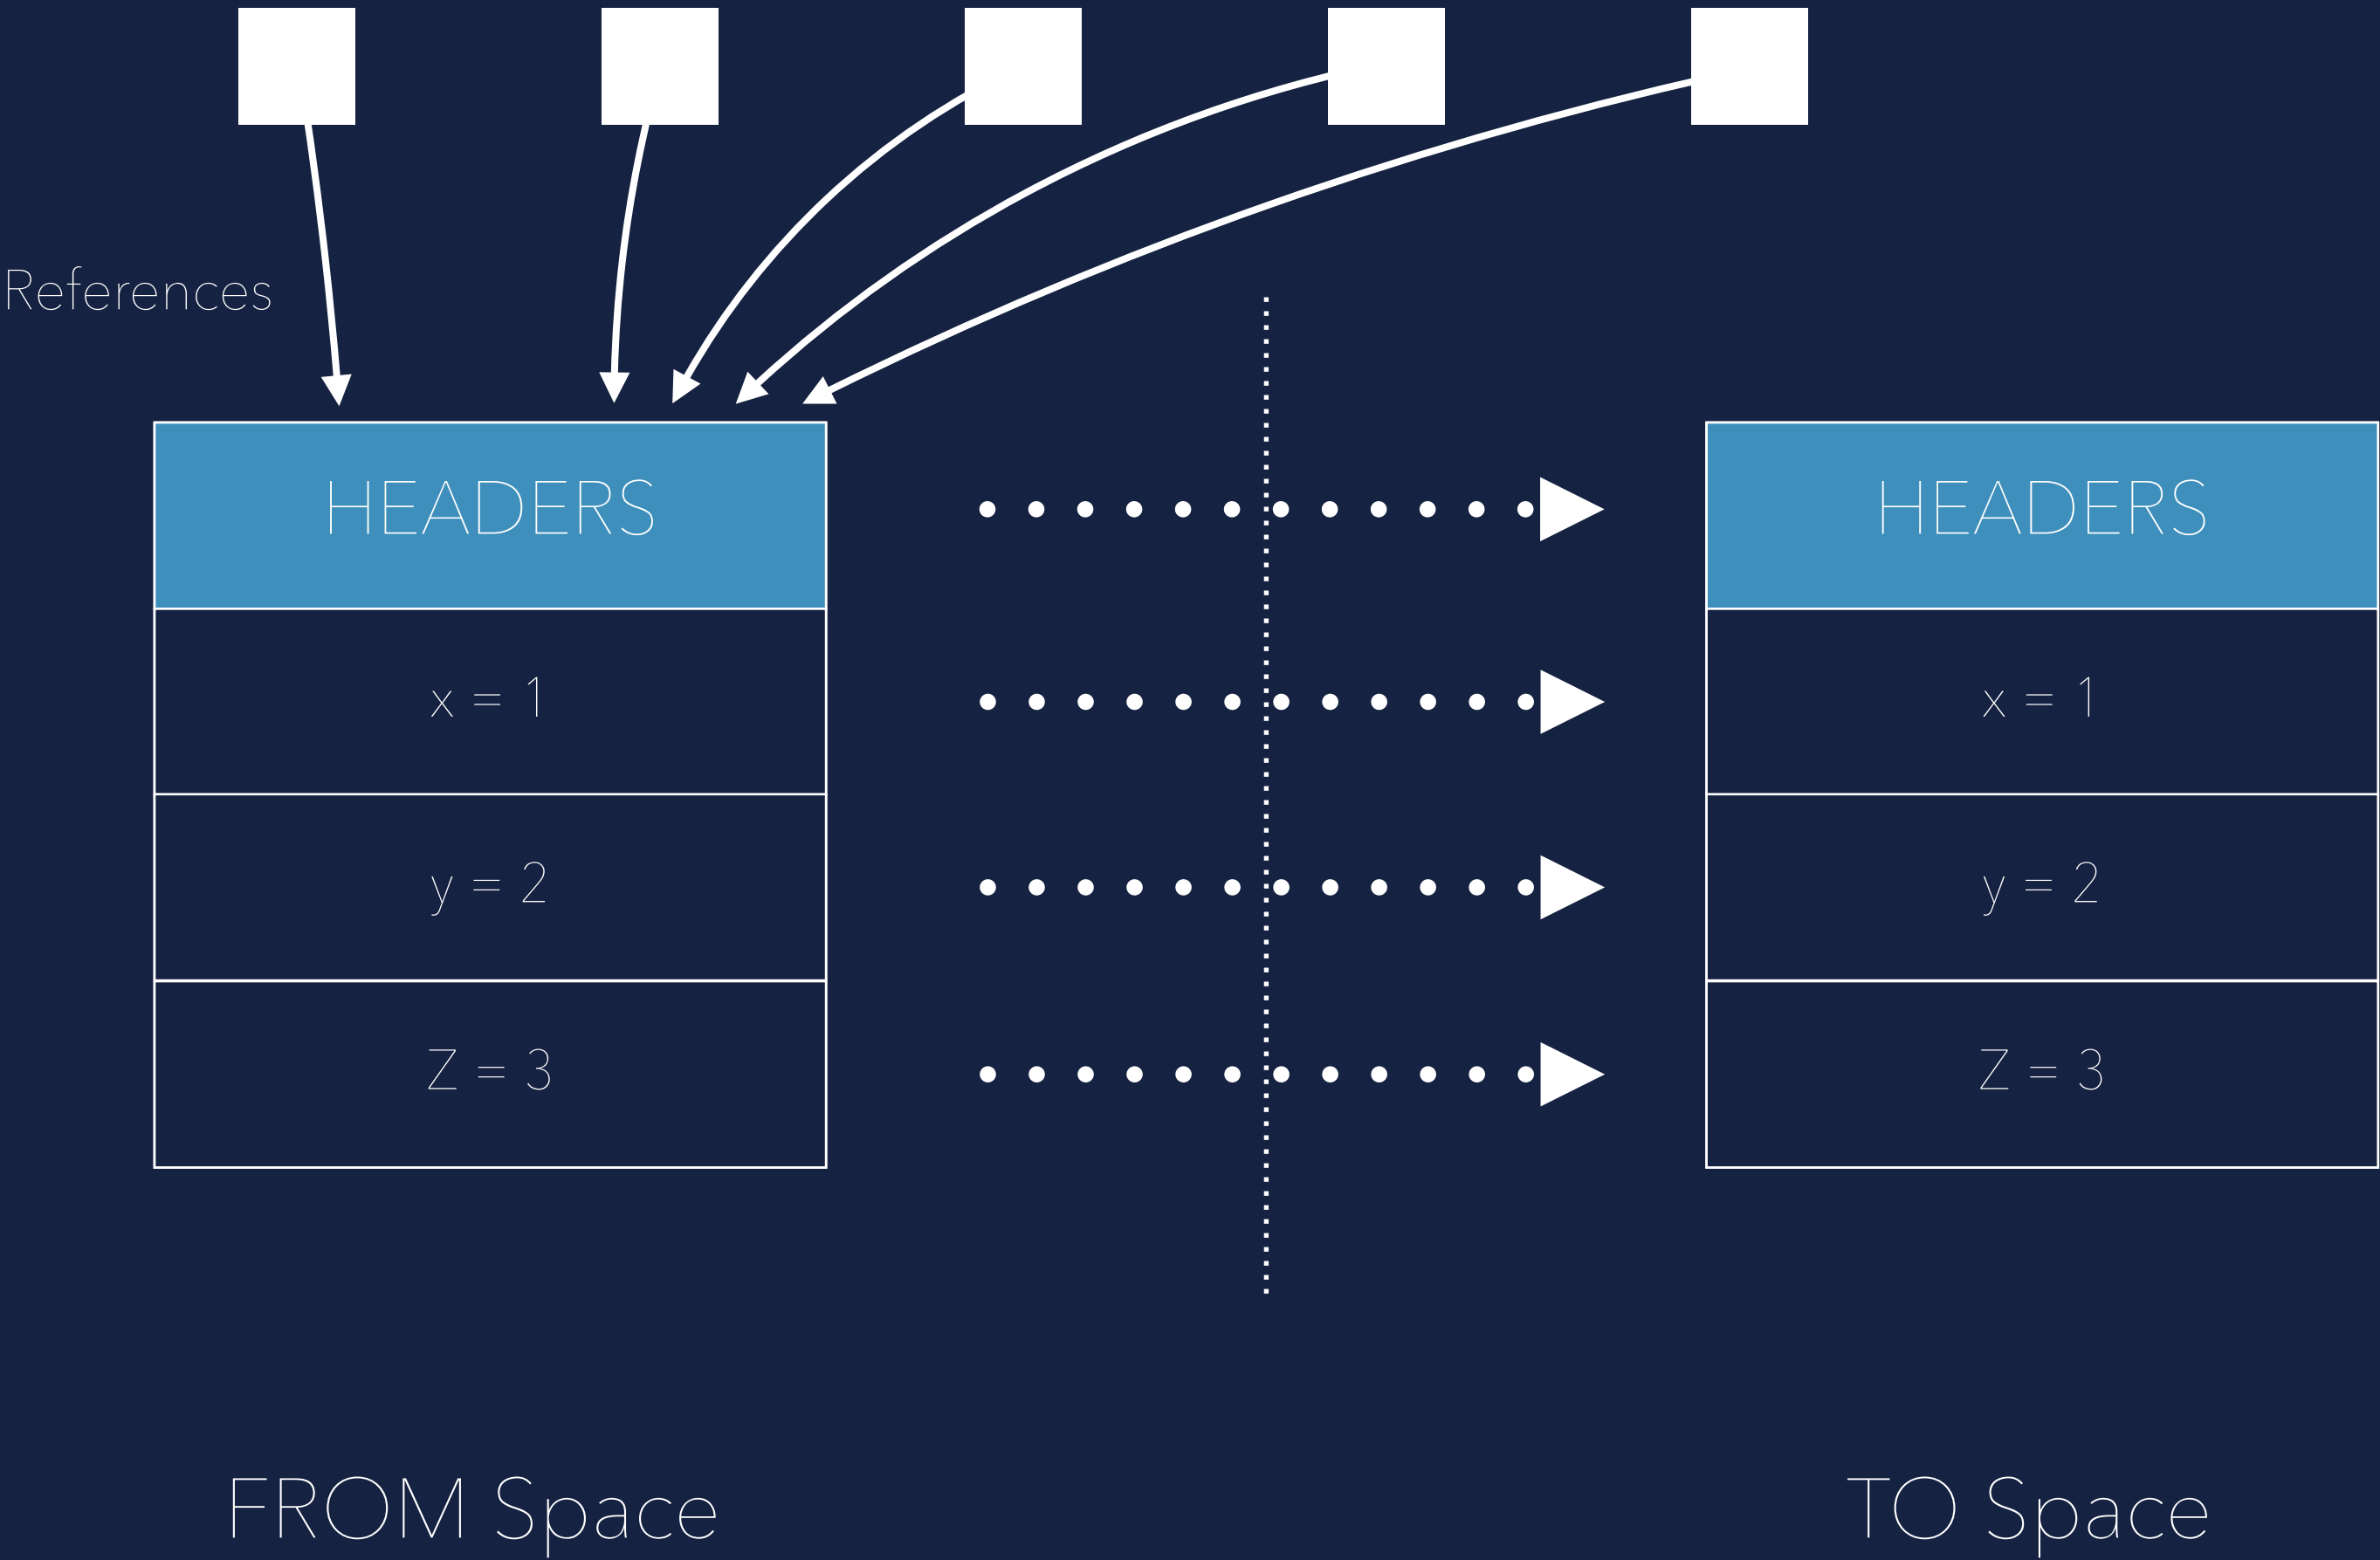
FROM Space

TO Space

Stop the World  
(the Mutator)

# CONCURRENCY IS HARD...

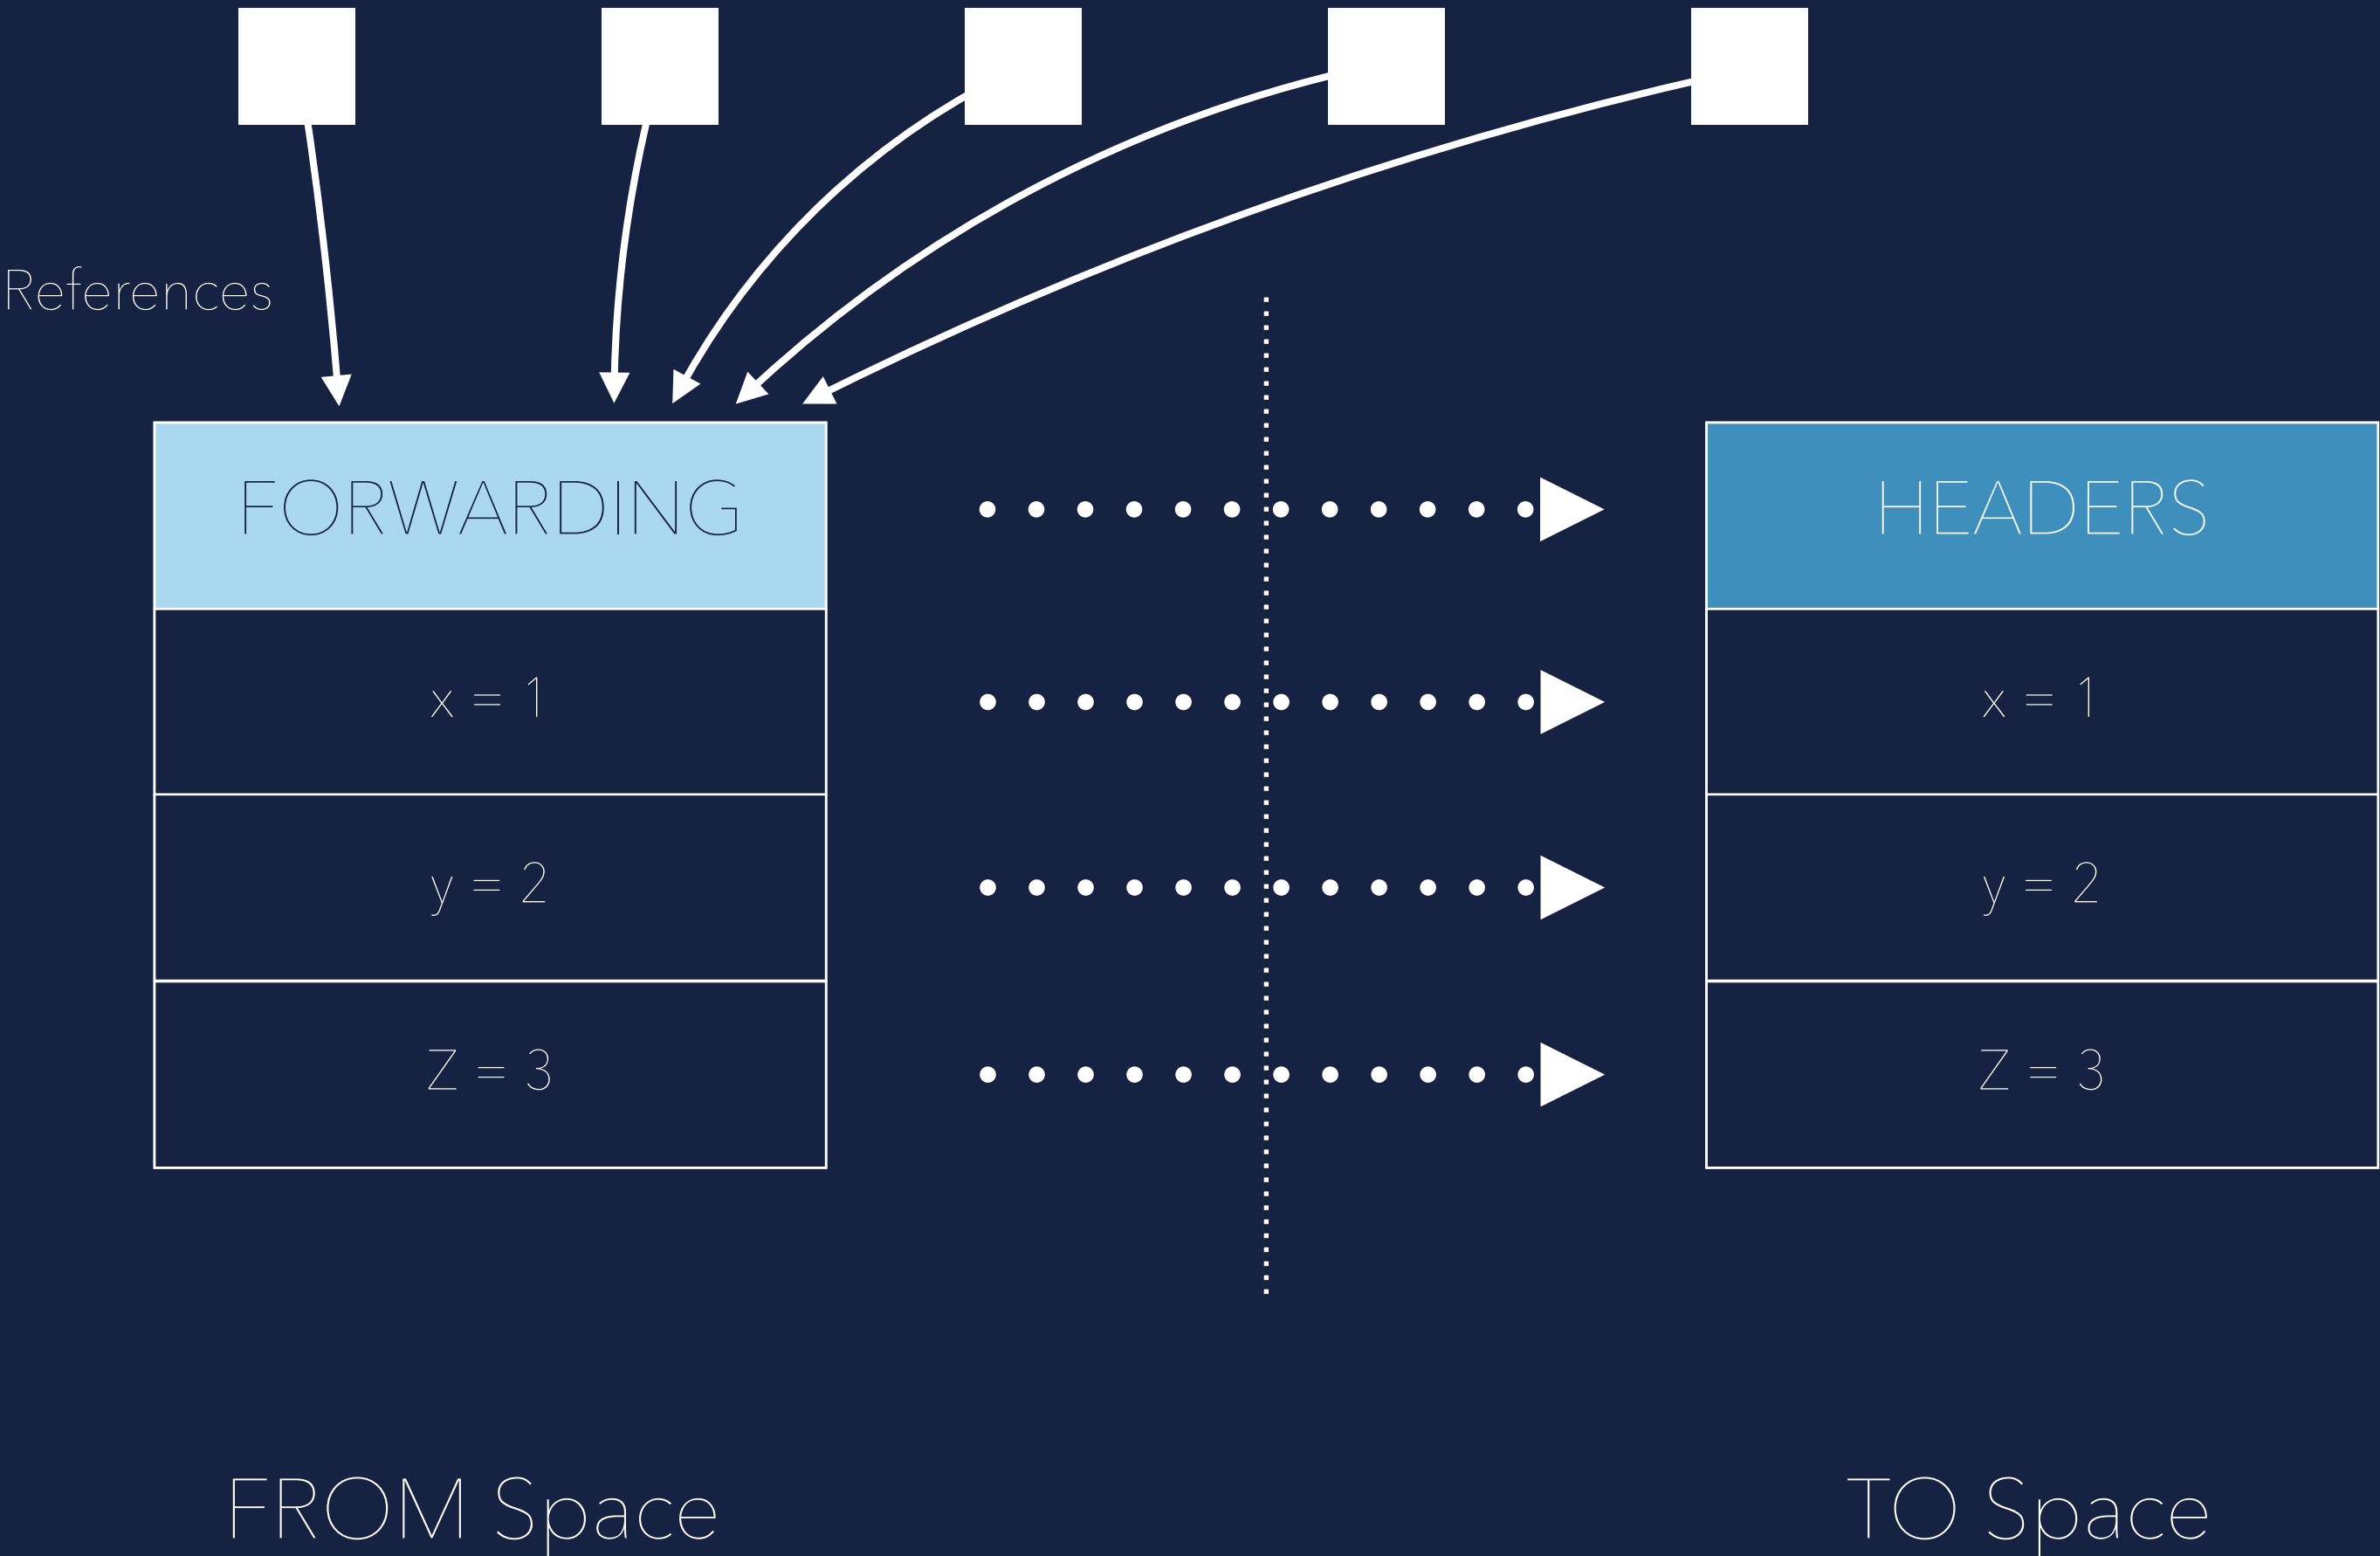
Stop the world copying





# CONCURRENCY IS HARD...

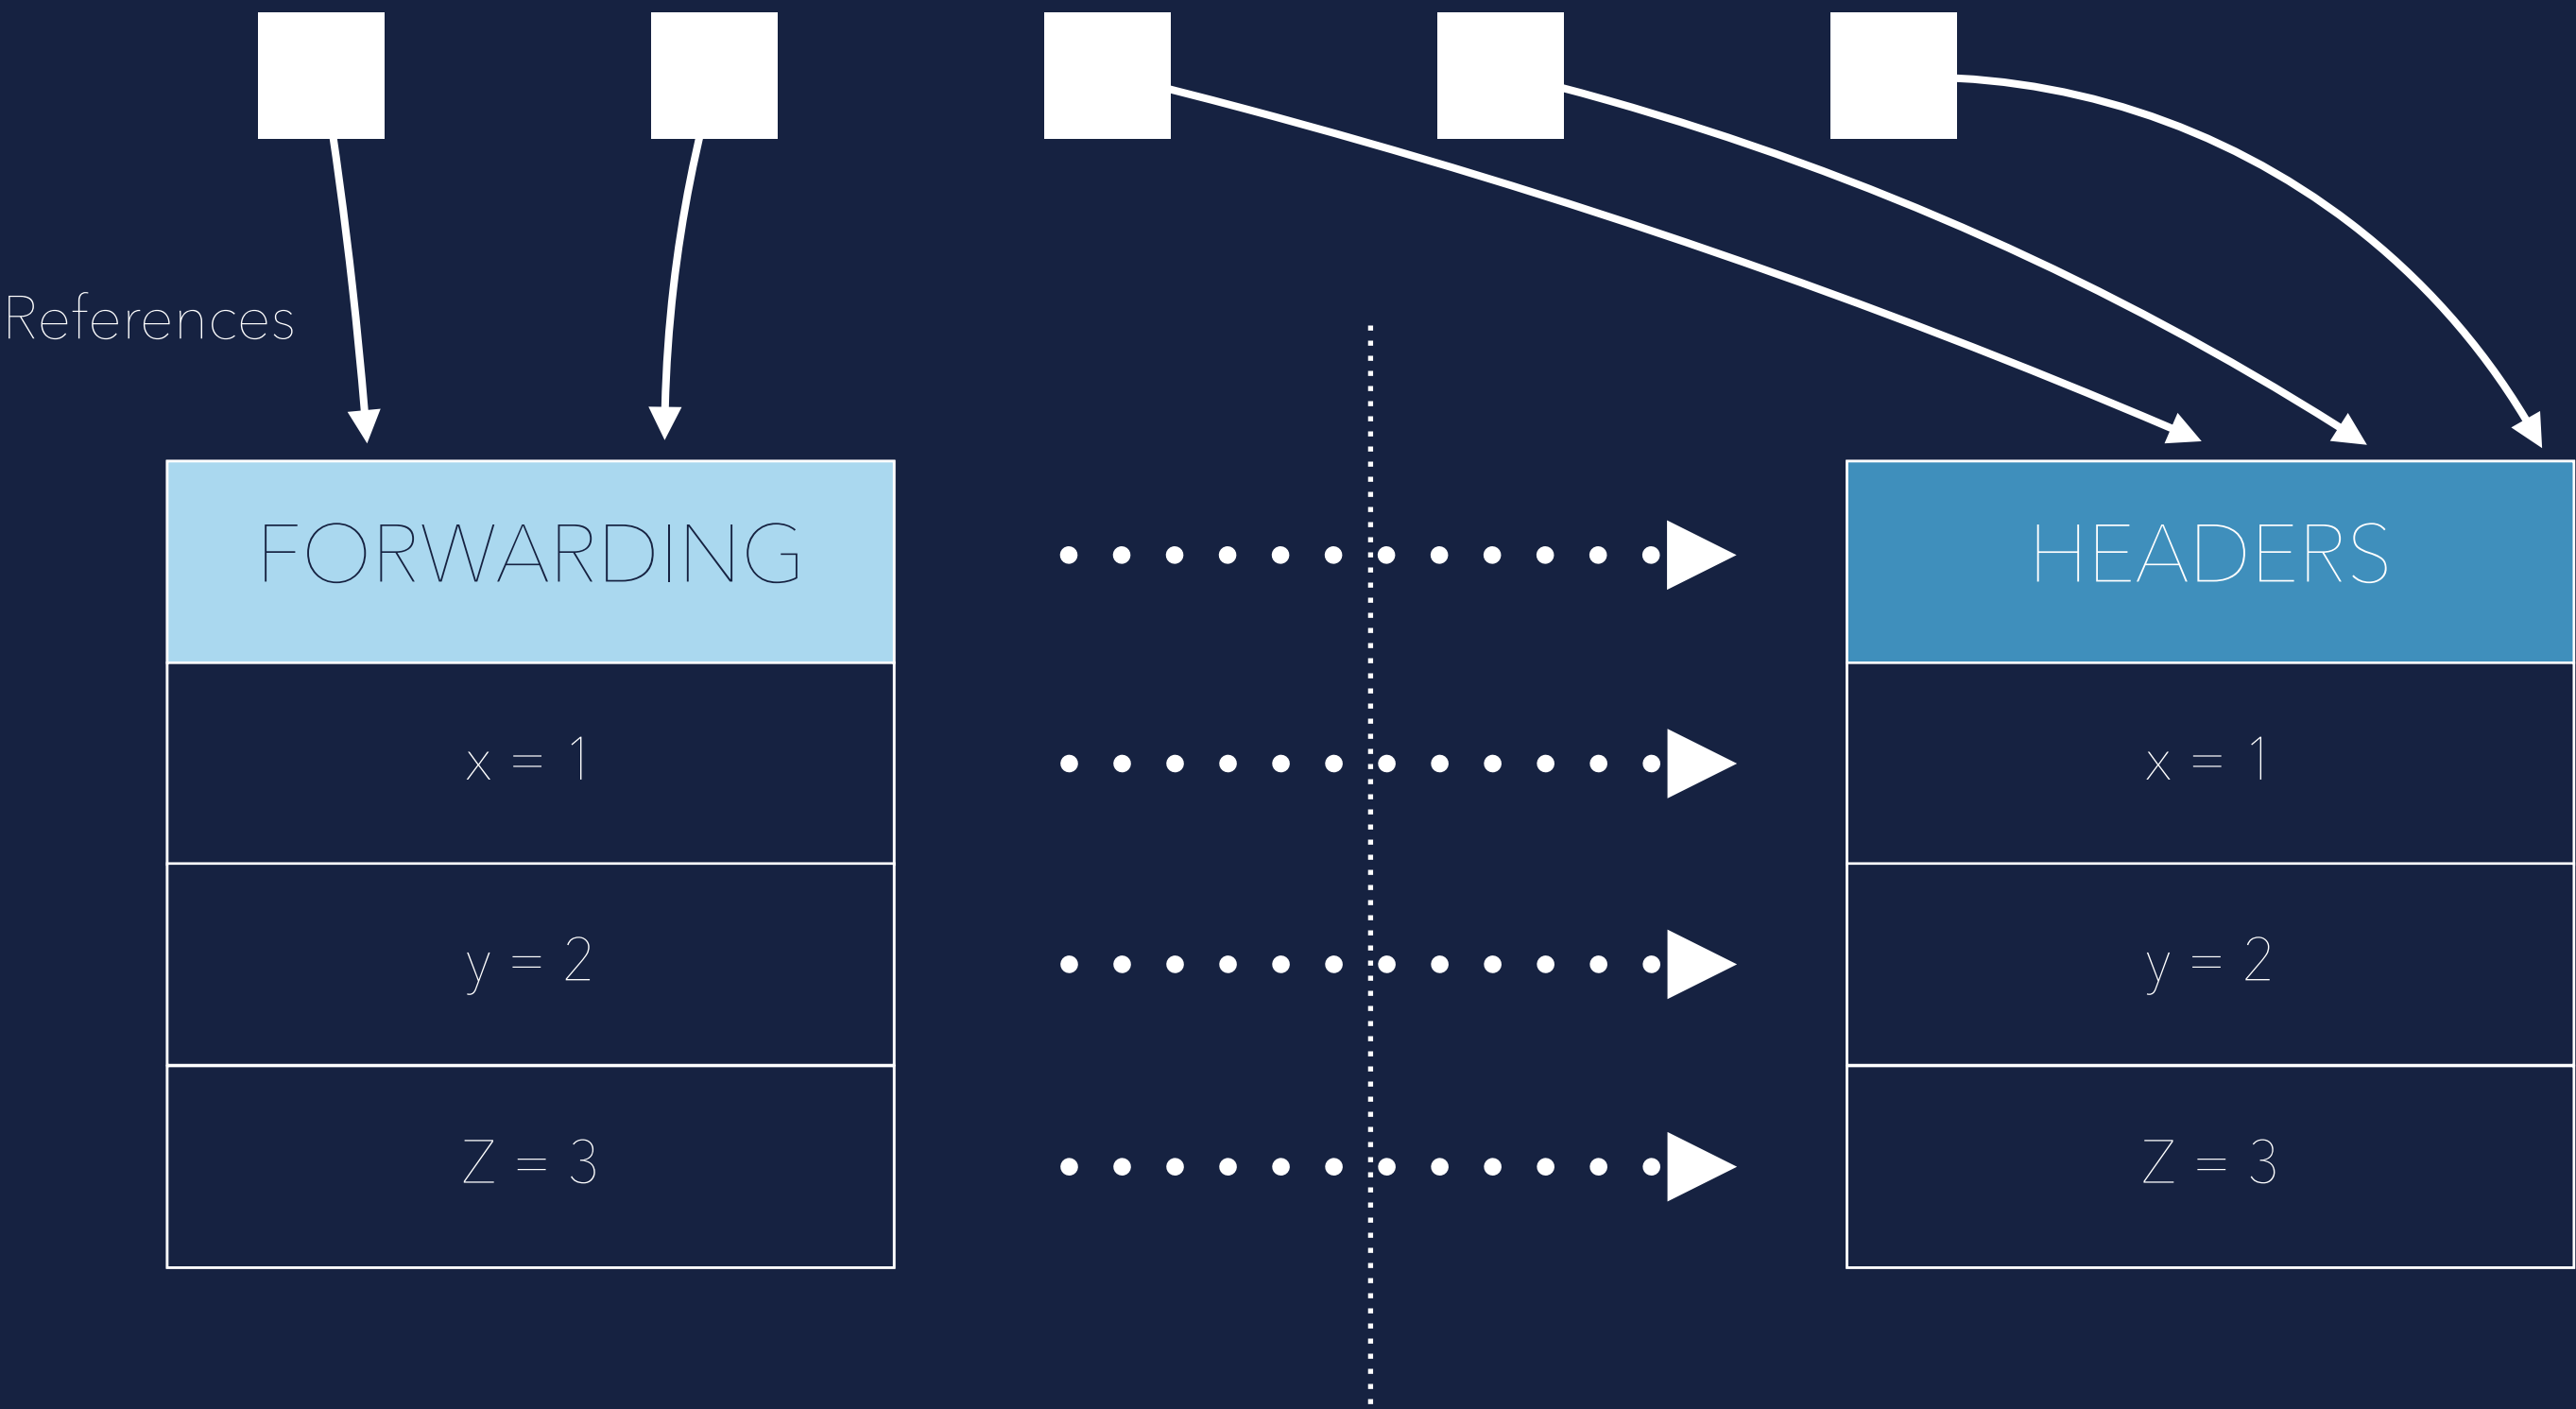
Stop the world copying



Update all references  
(Save the pointer that forwards the copy)

# CONCURRENCY IS HARD...

Stop the world copying




Update all references  
(Walk the heap and replace all references  
with forwarding pointer to new location)

FROM Space

TO Space

# CONCURRENCY IS HARD...


Stop the world copying



Update all references  
(Walk the heap and replace all references with forwarding pointer to new location)

# CONCURRENCY IS HARD...


Stop the world copying



Remove old objects and  
continue running the Mutator

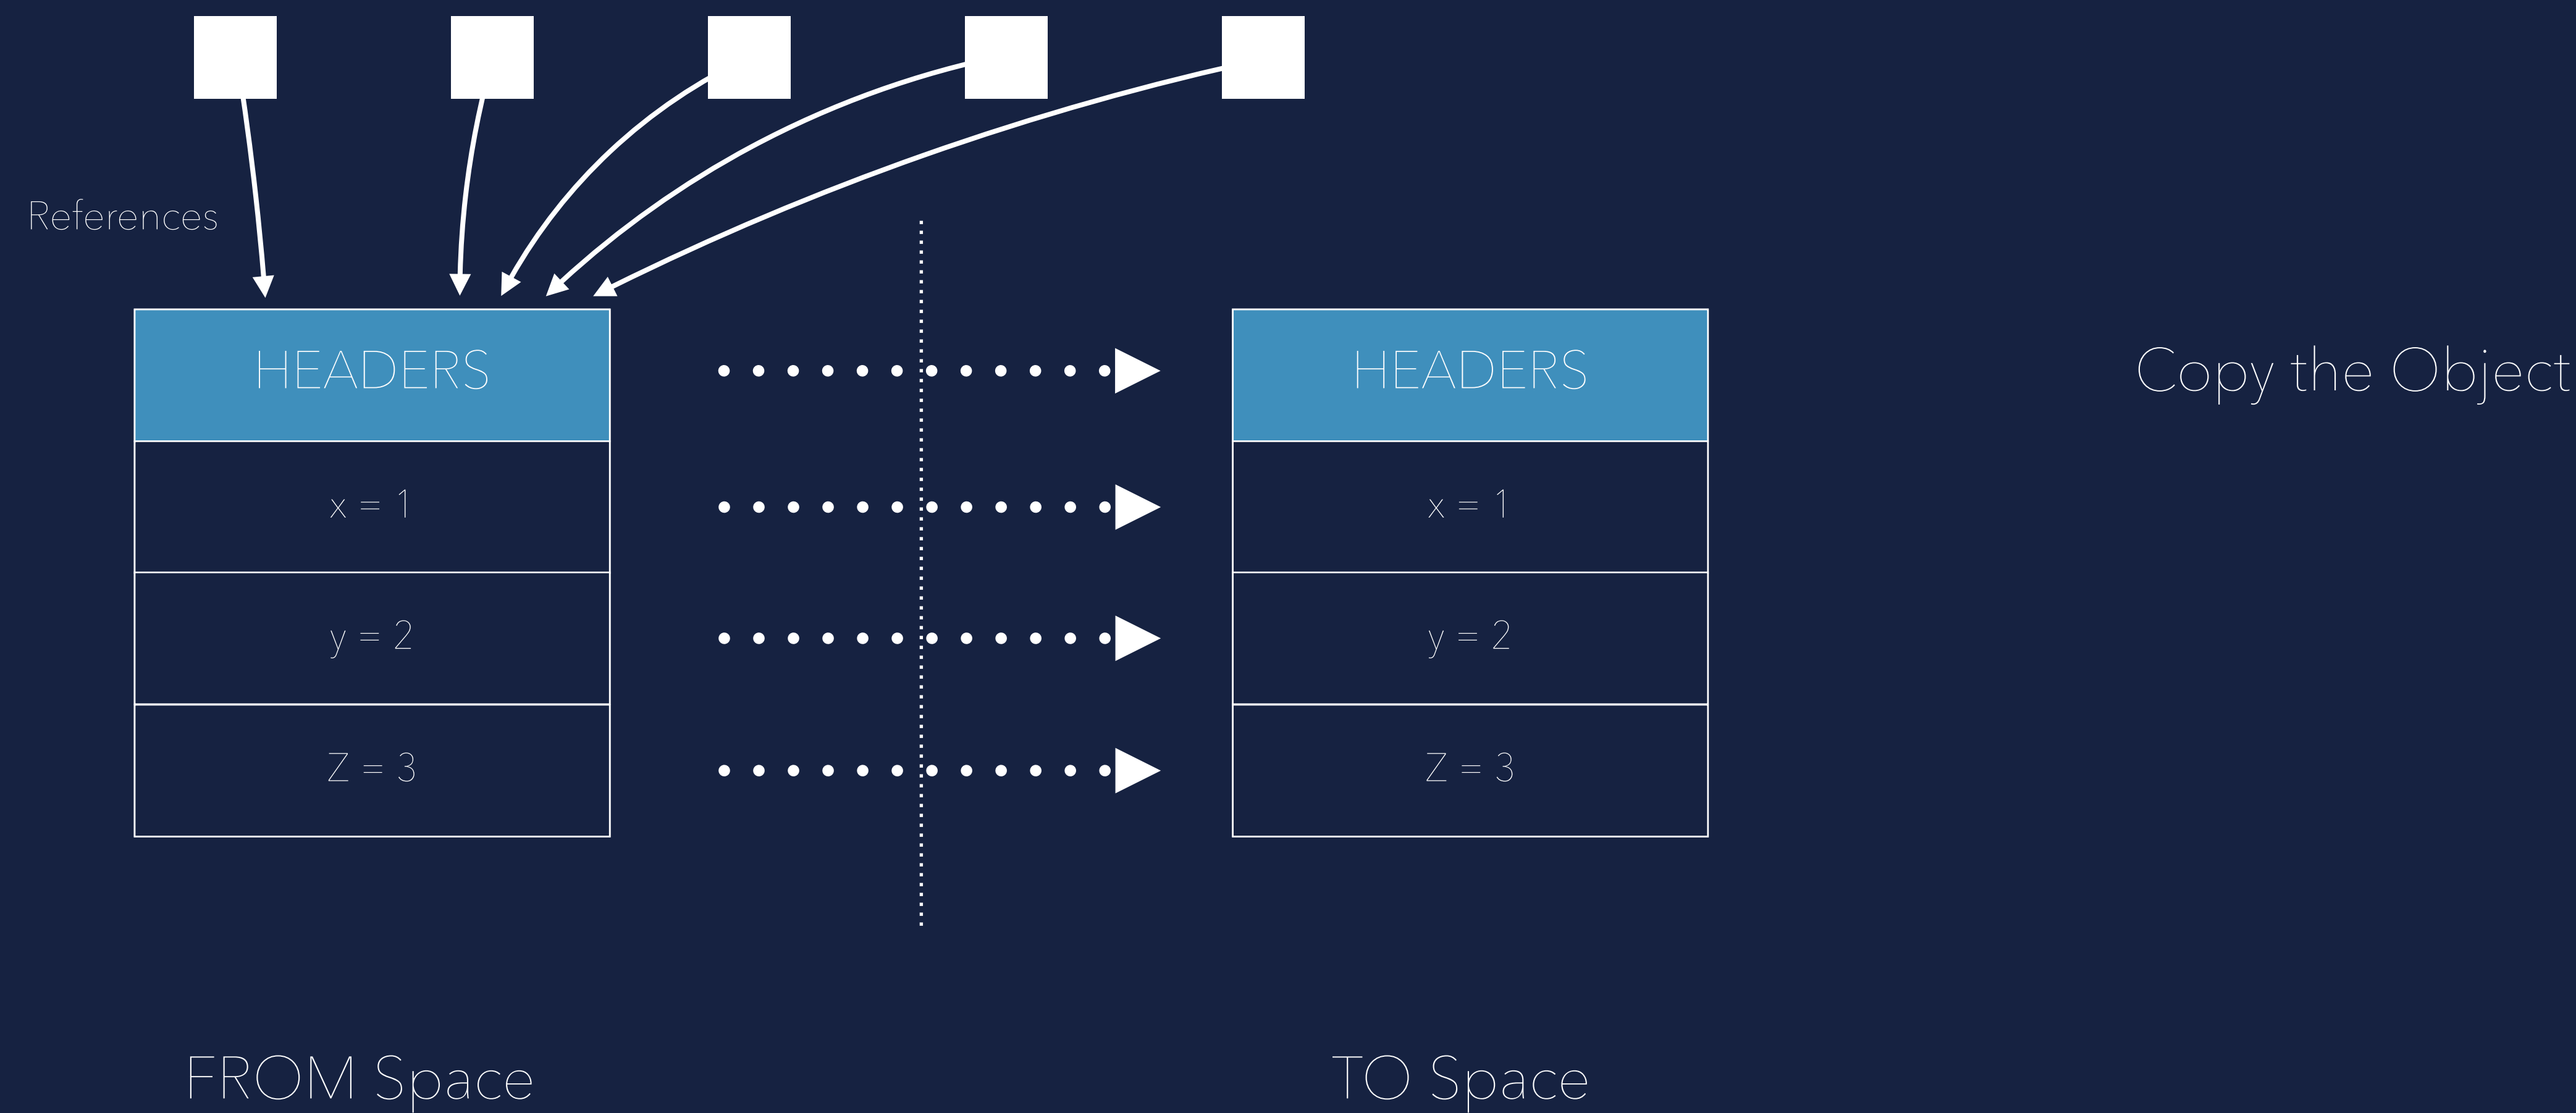
# CONCURRENCY IS HARD...

## Concurrent copying



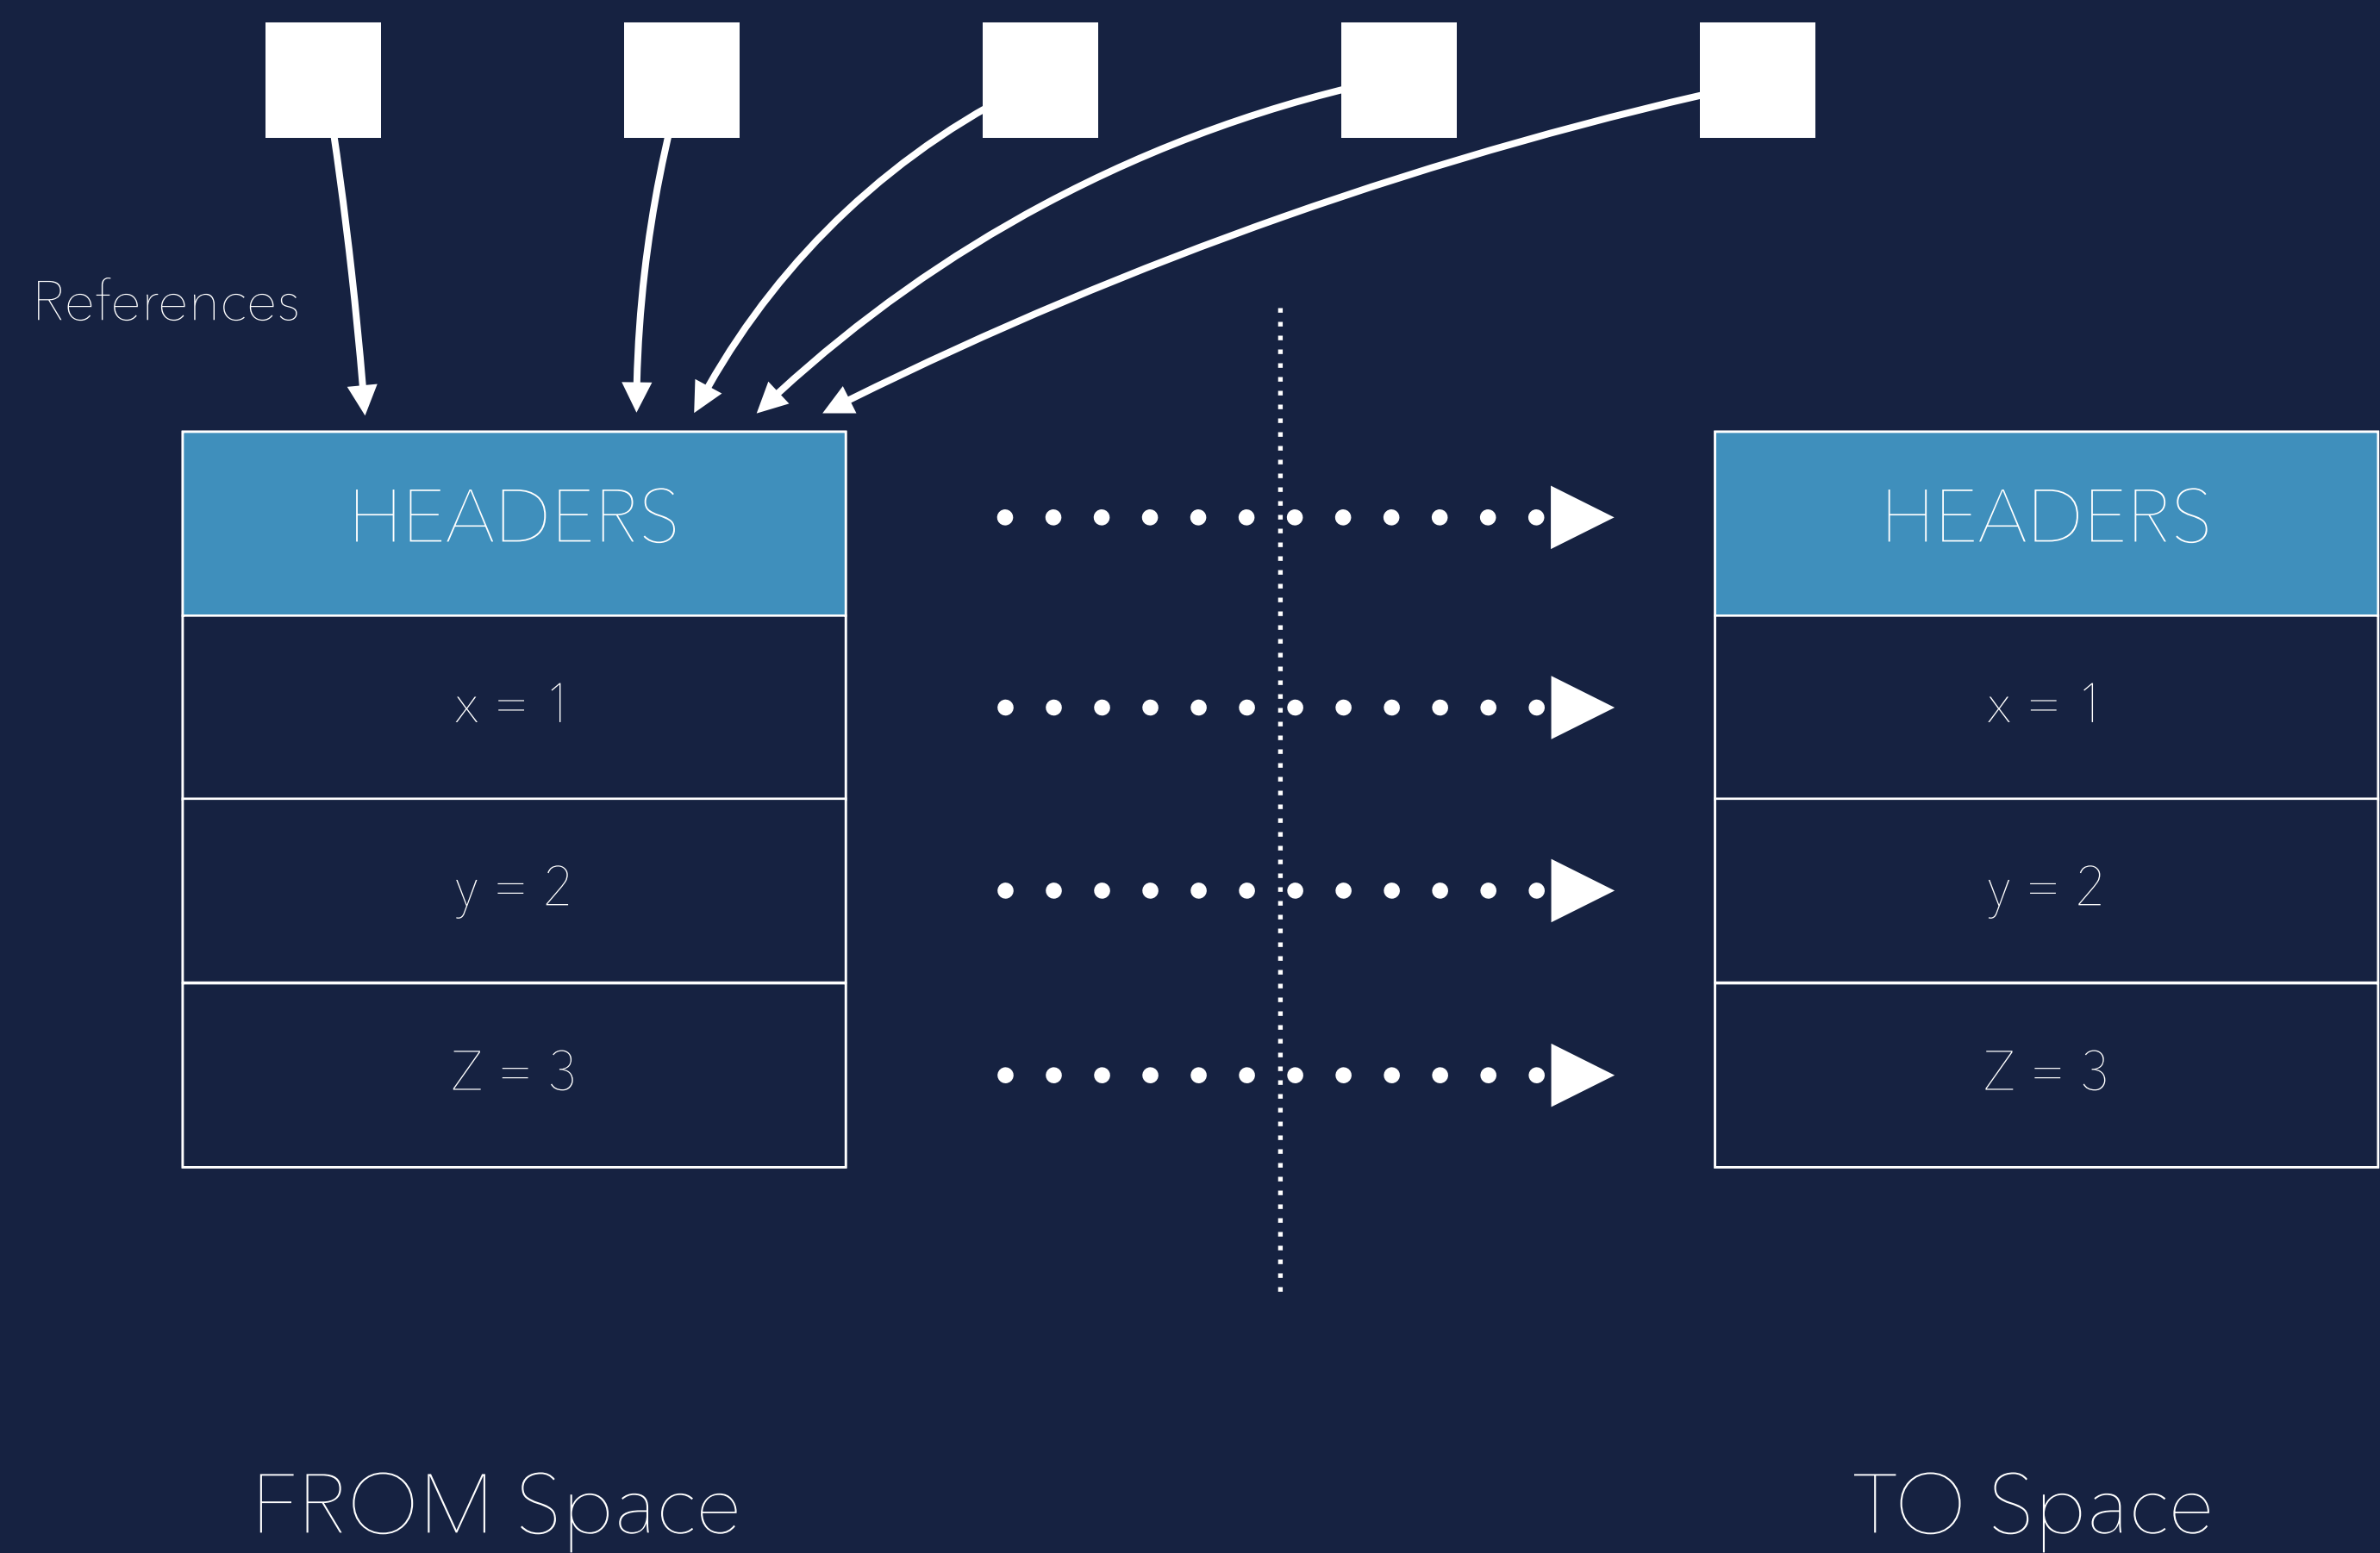
# CONCURRENCY IS HARD...

## Concurrent copying



# CONCURRENCY IS HARD...

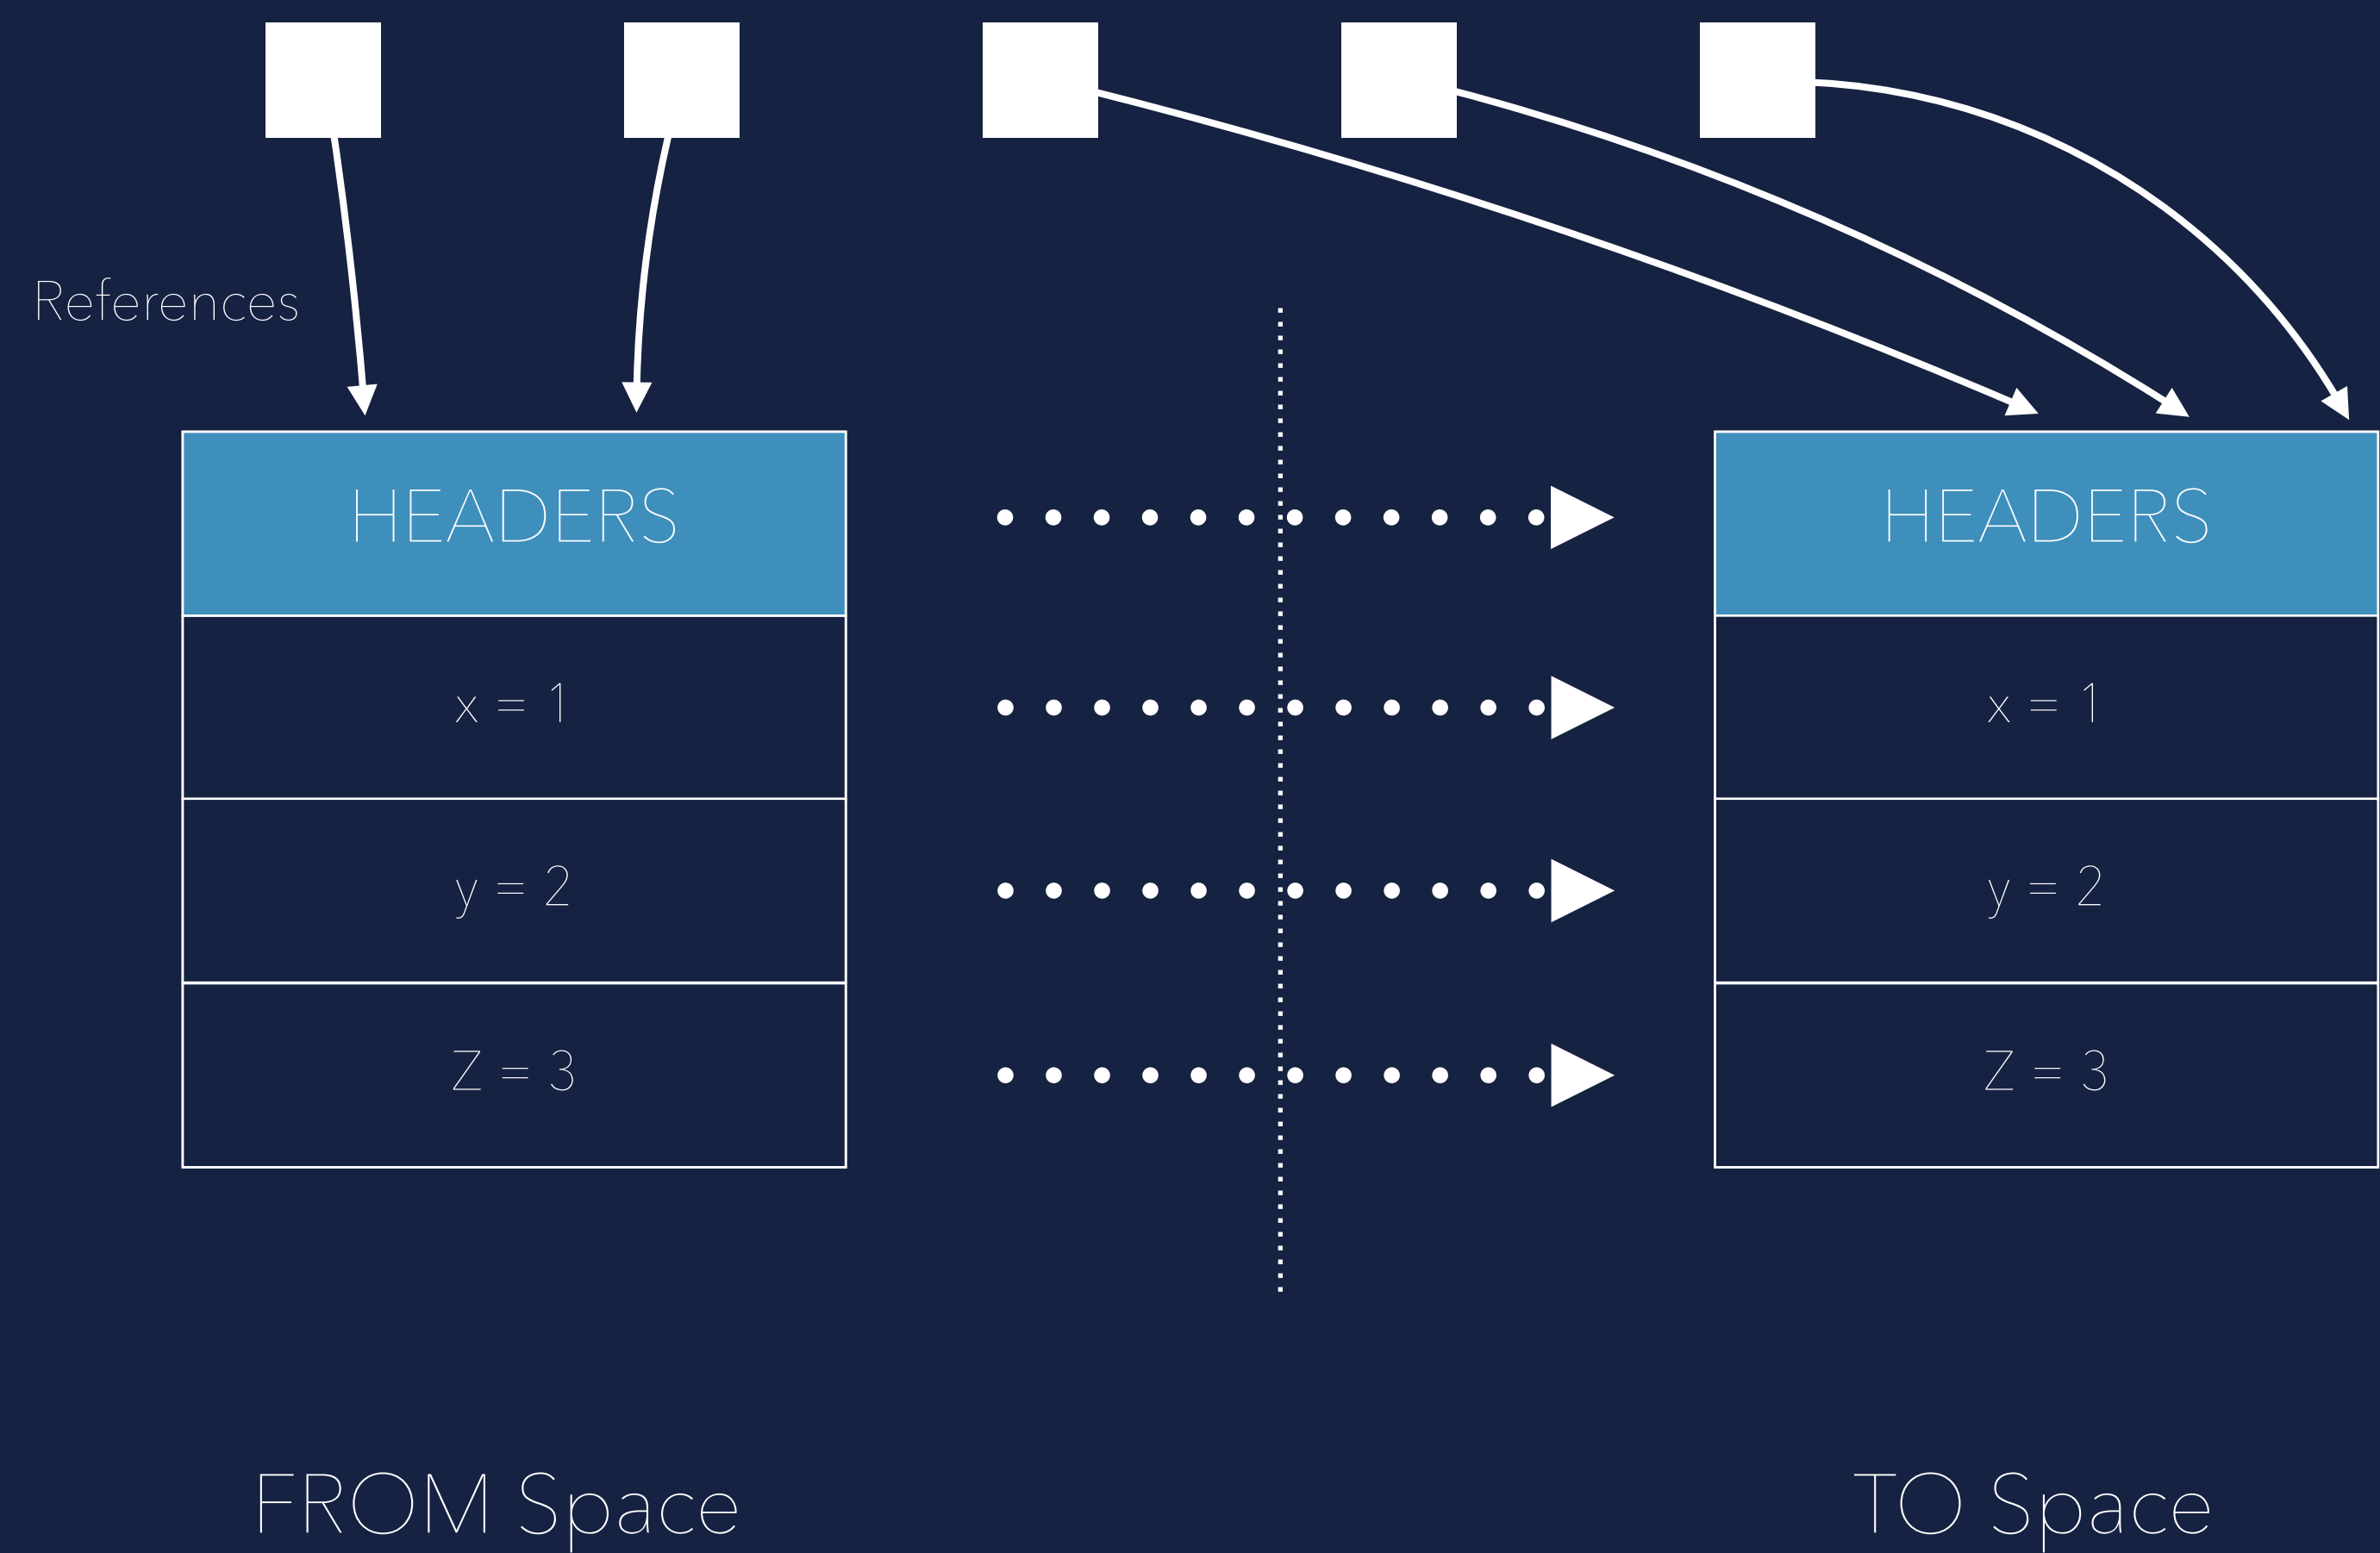
## Concurrent copying



Now both Objects are reachable !

# CONCURRENCY IS HARD...

## Concurrent copying




Now both Objects are reachable !  
And can be accessed in parallel by different Threads.



# CONCURRENCY IS HARD...

## Concurrent copying

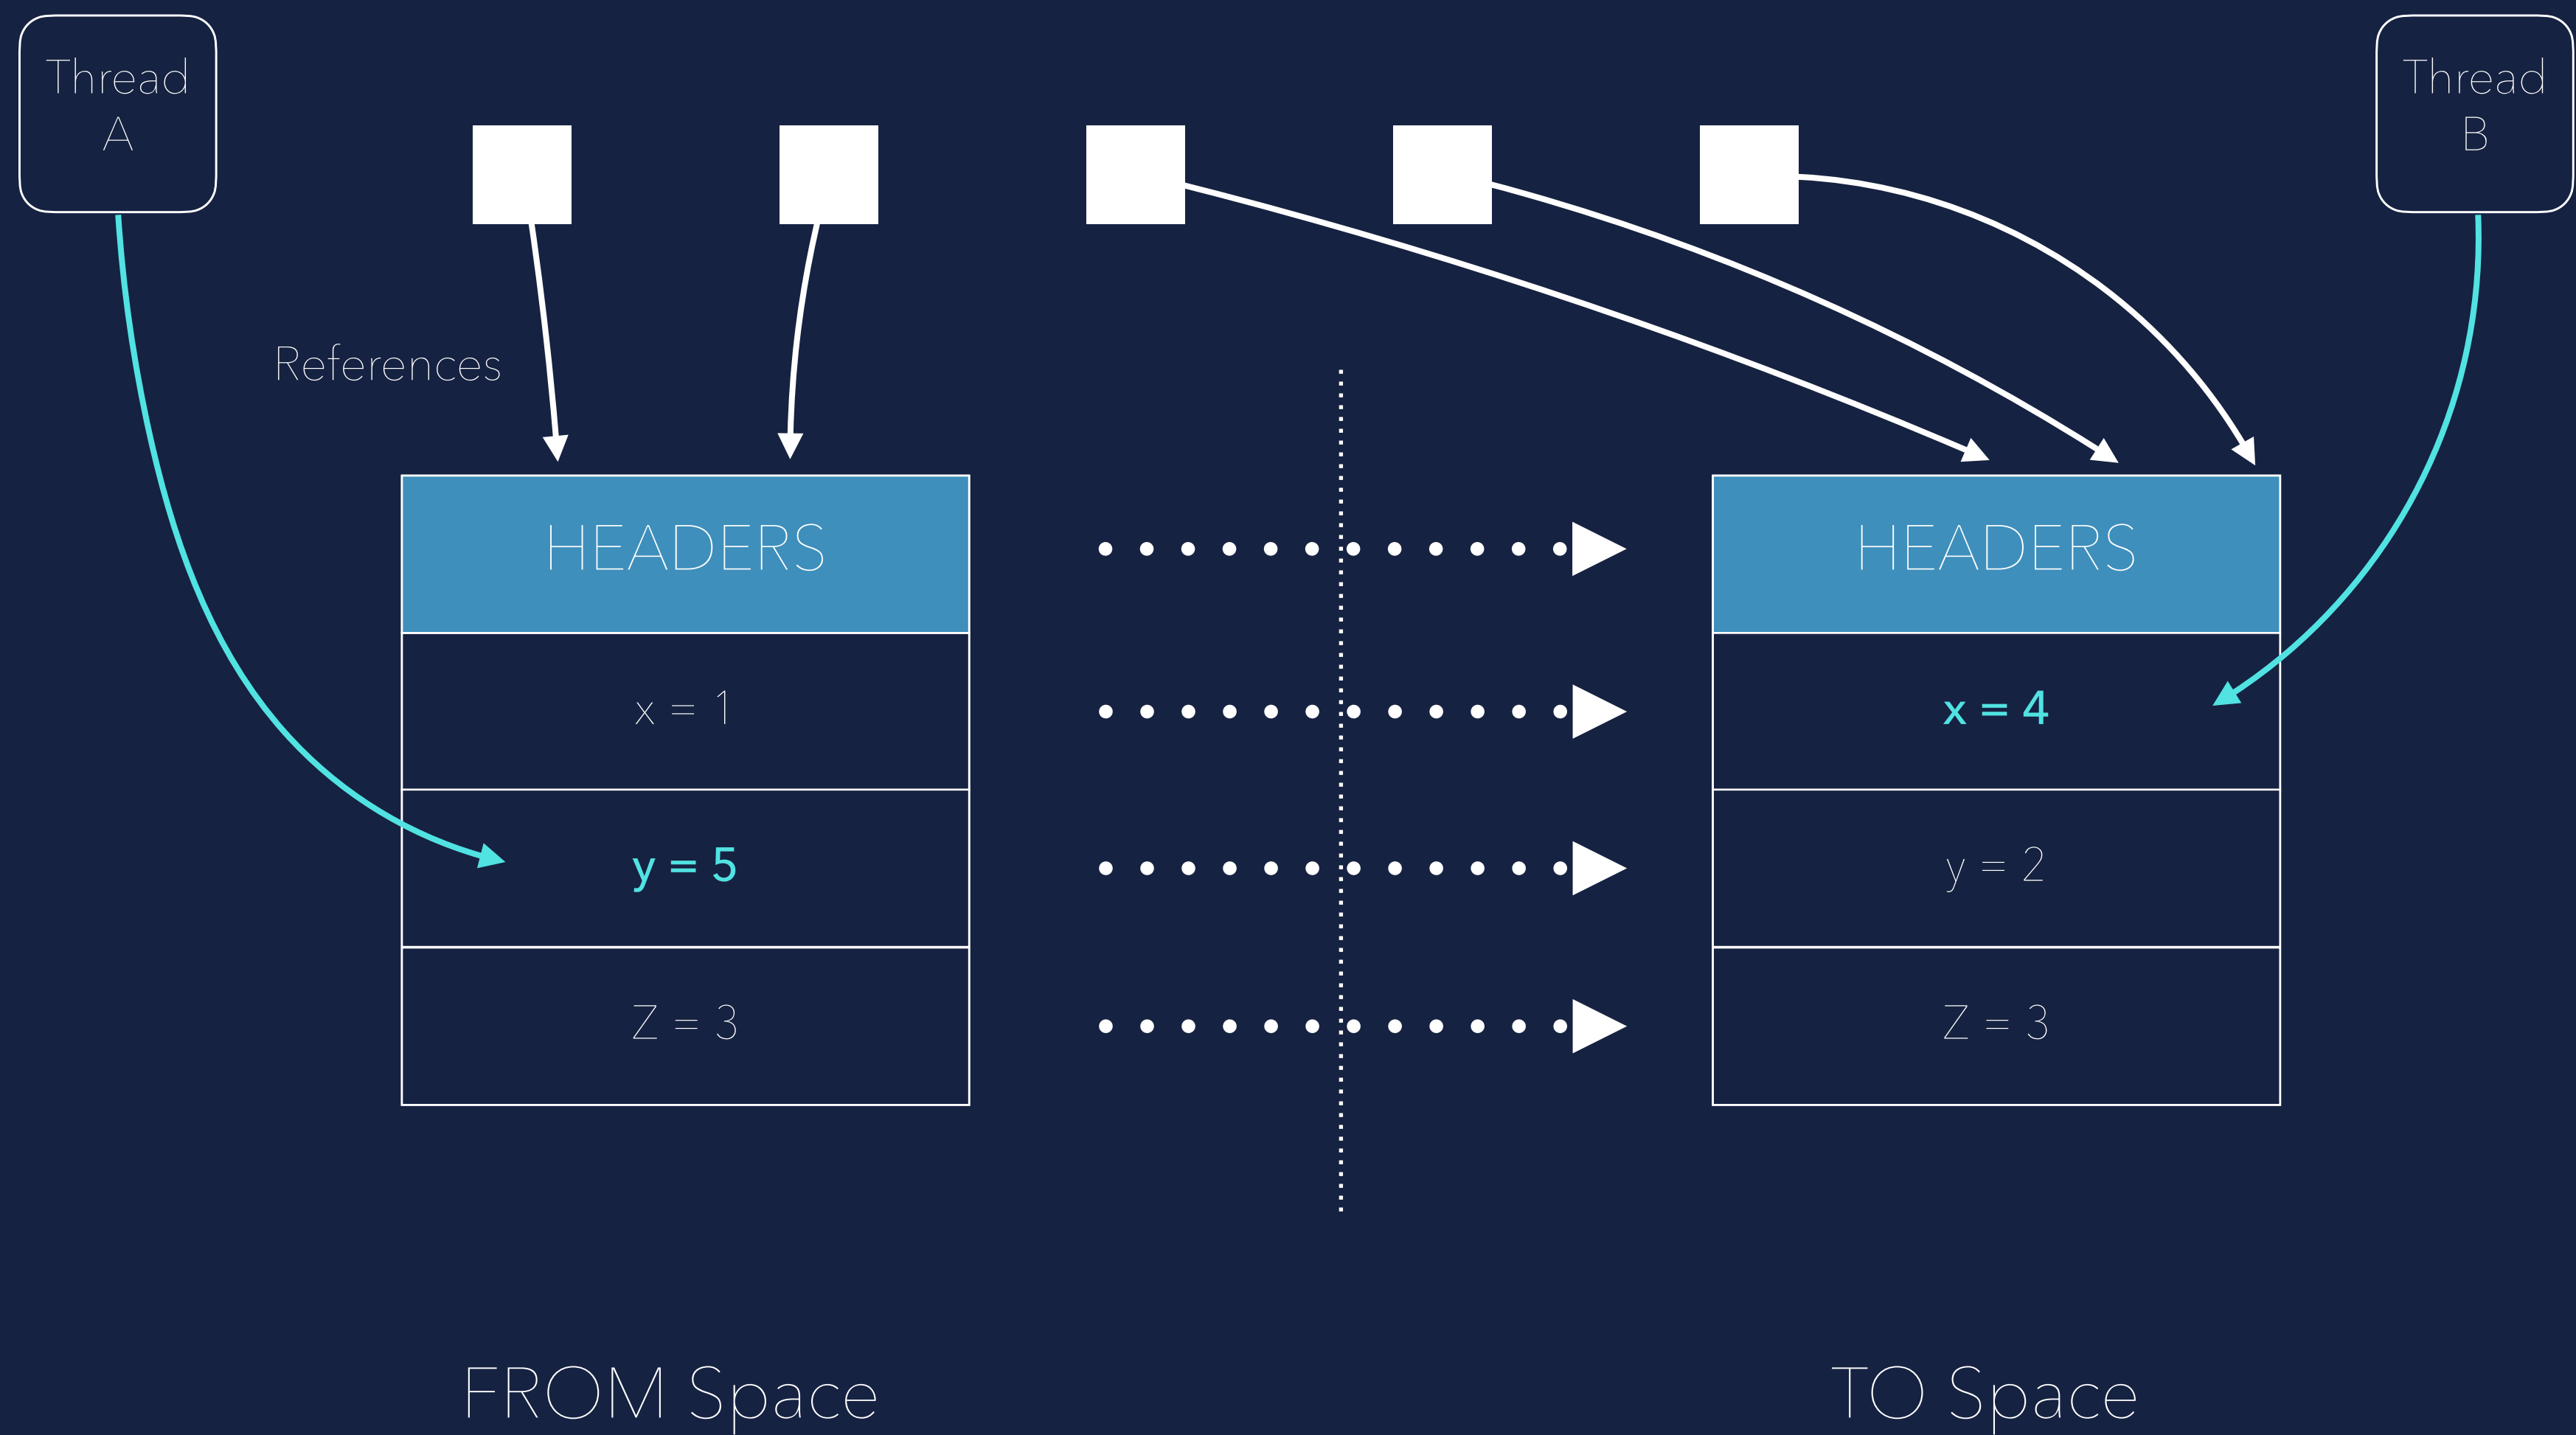


Threads can write to both Objects !



# CONCURRENCY IS HARD...

## Concurrent copying

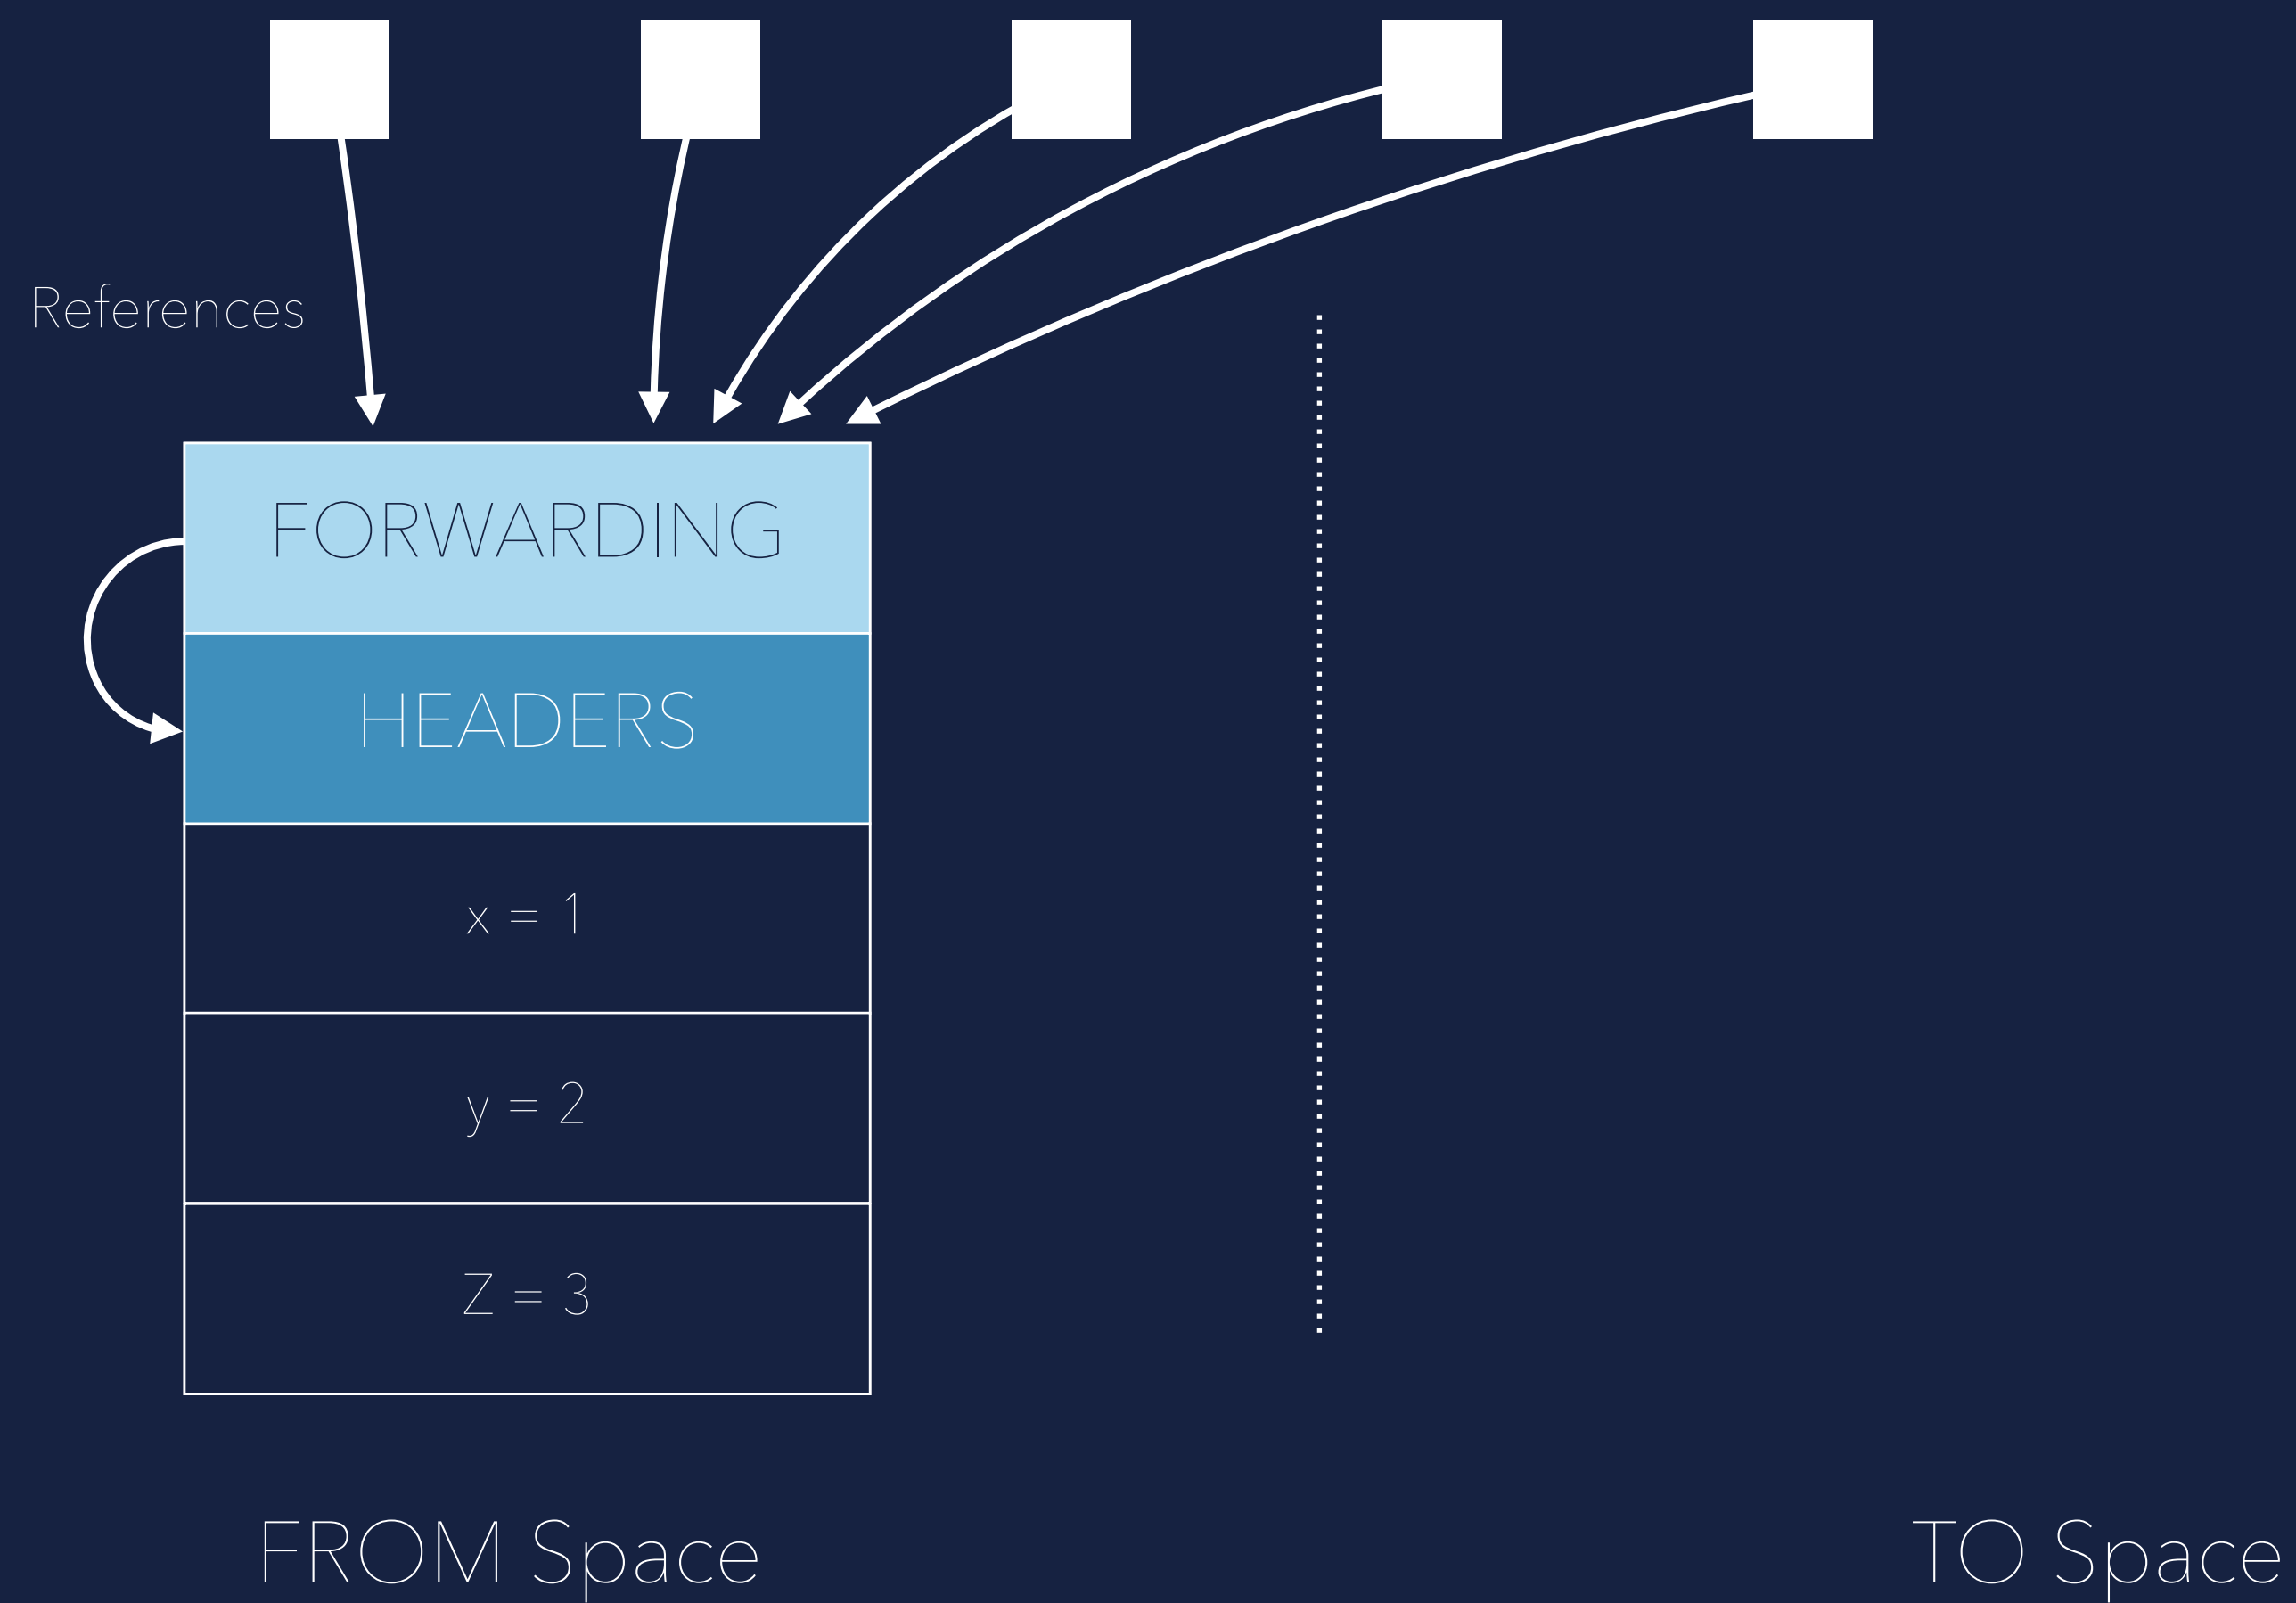


Threads can write to both Objects !  
Which copy is correct ?



# CONCURRENCY IS HARD...

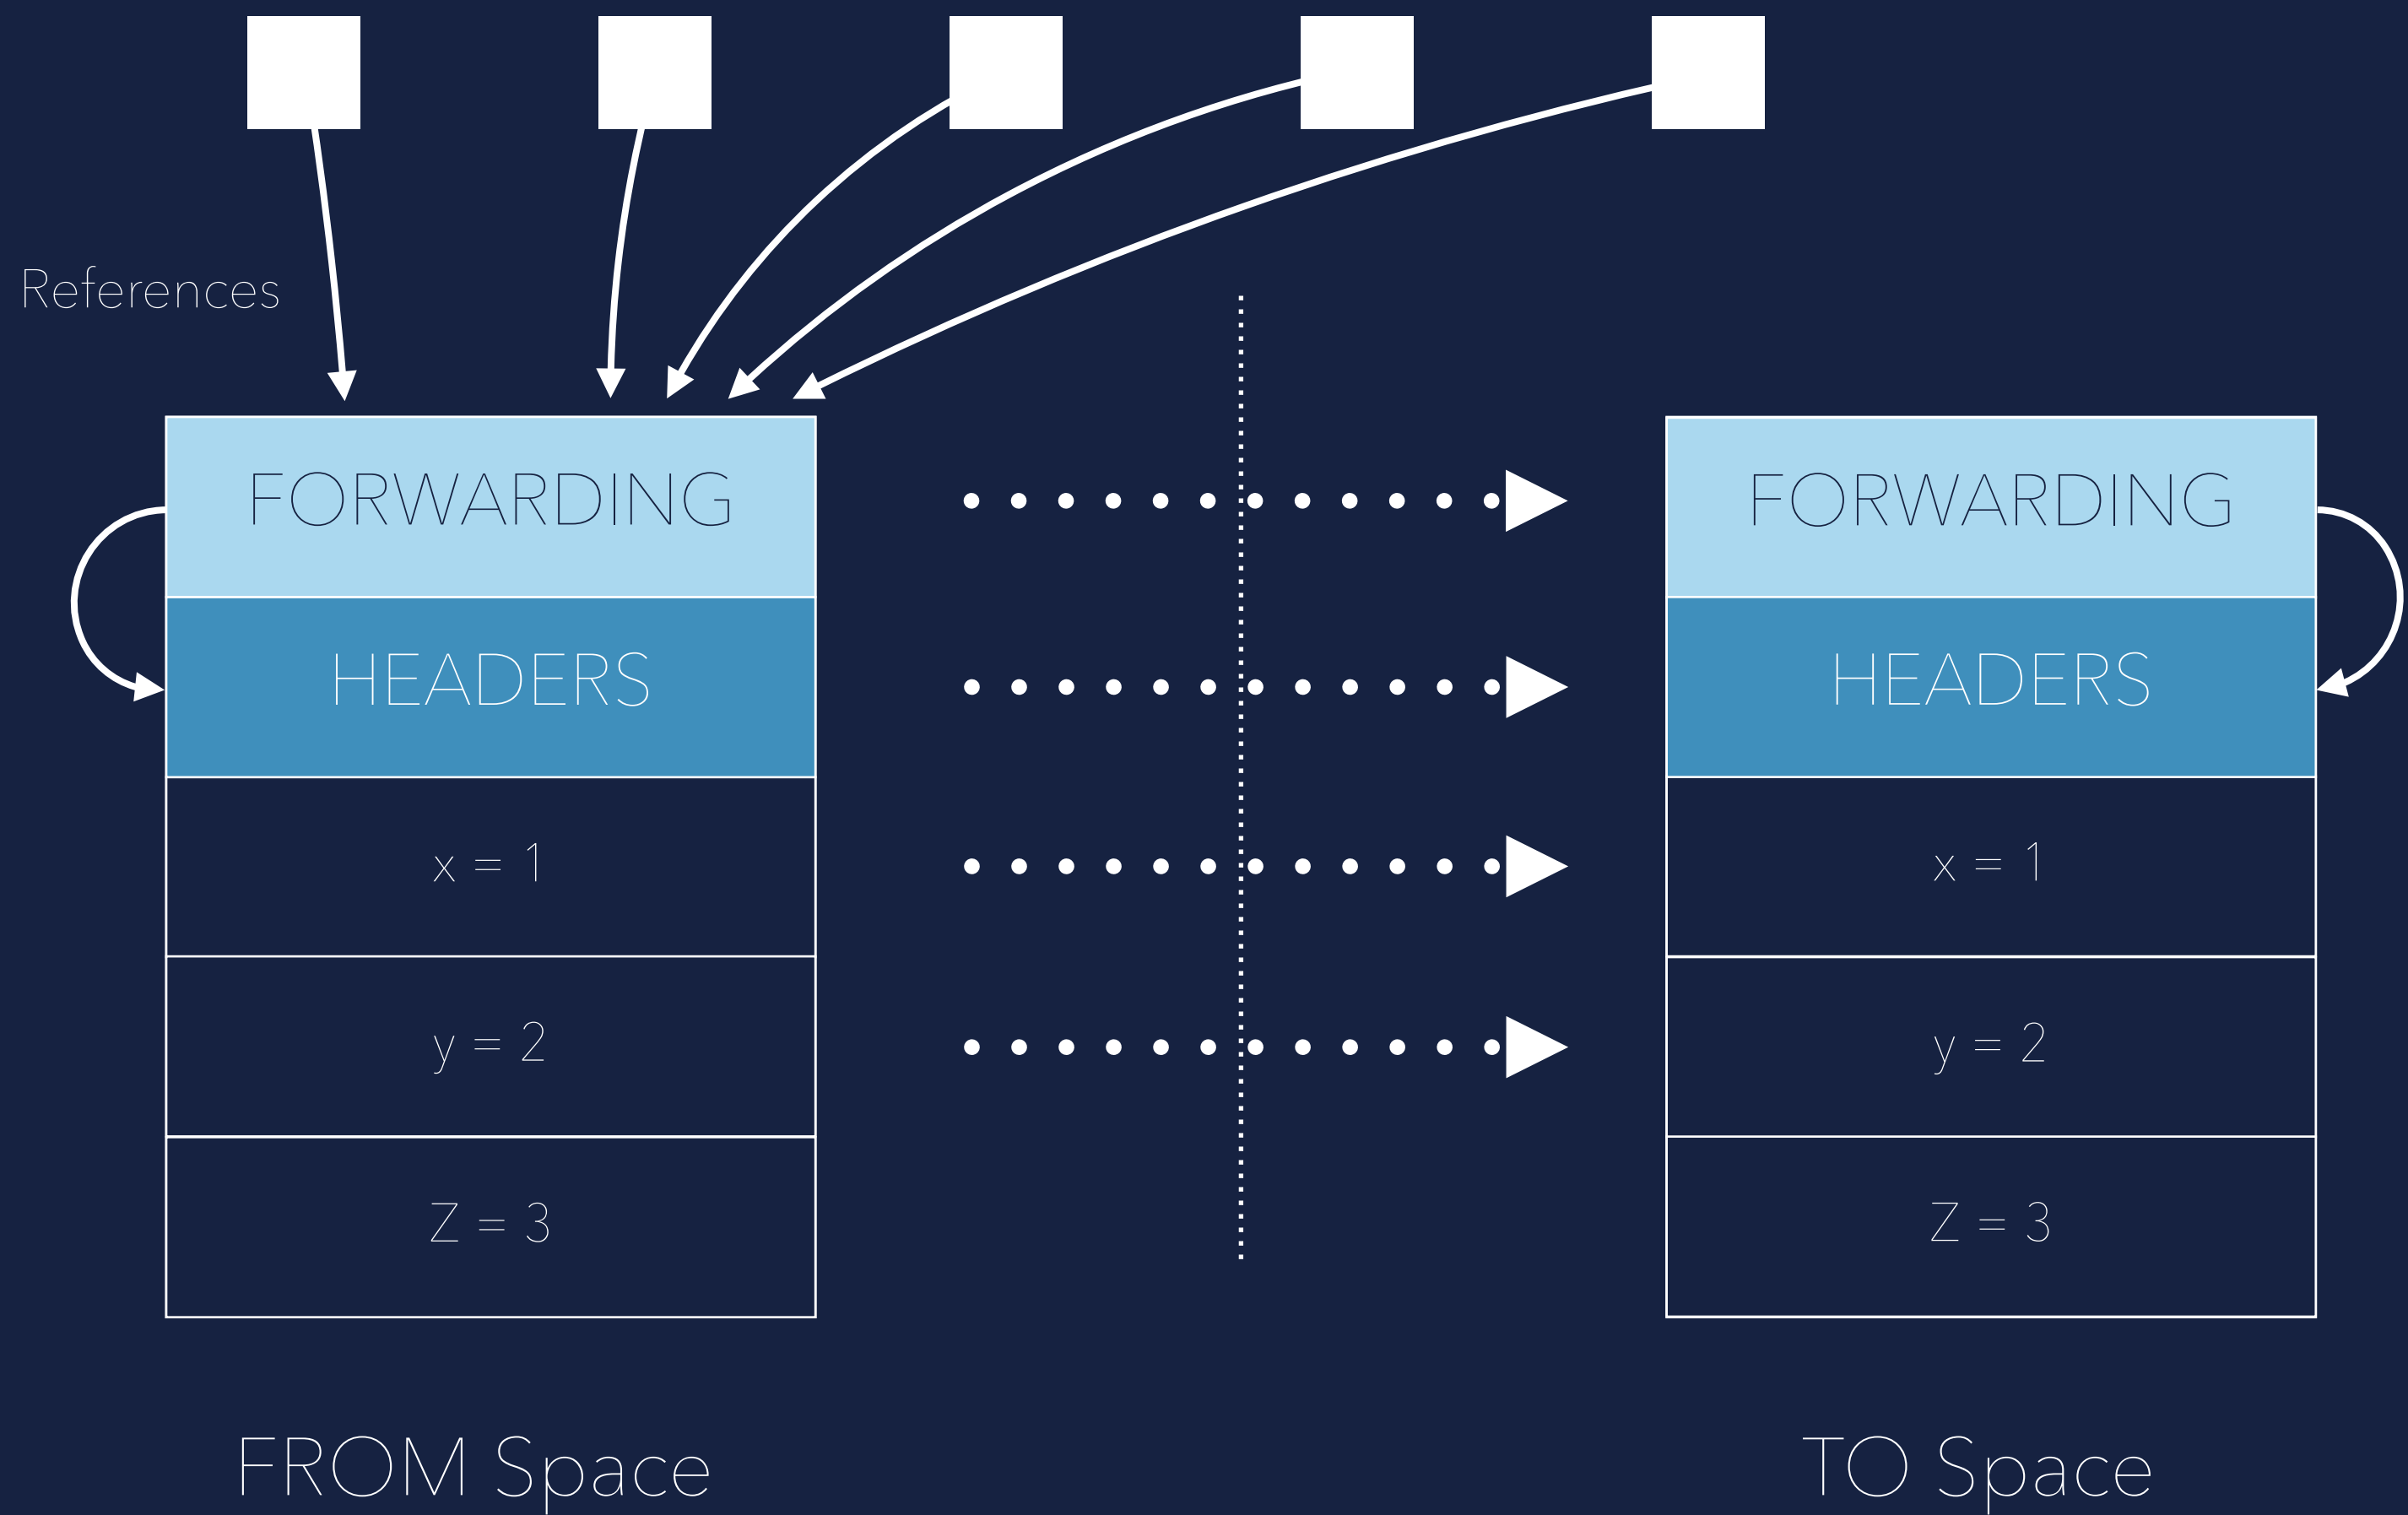
## Concurrent copying



Solution could be a  
Brooks Pointer  
(Initially points to the Object itself)

# CONCURRENCY IS HARD...

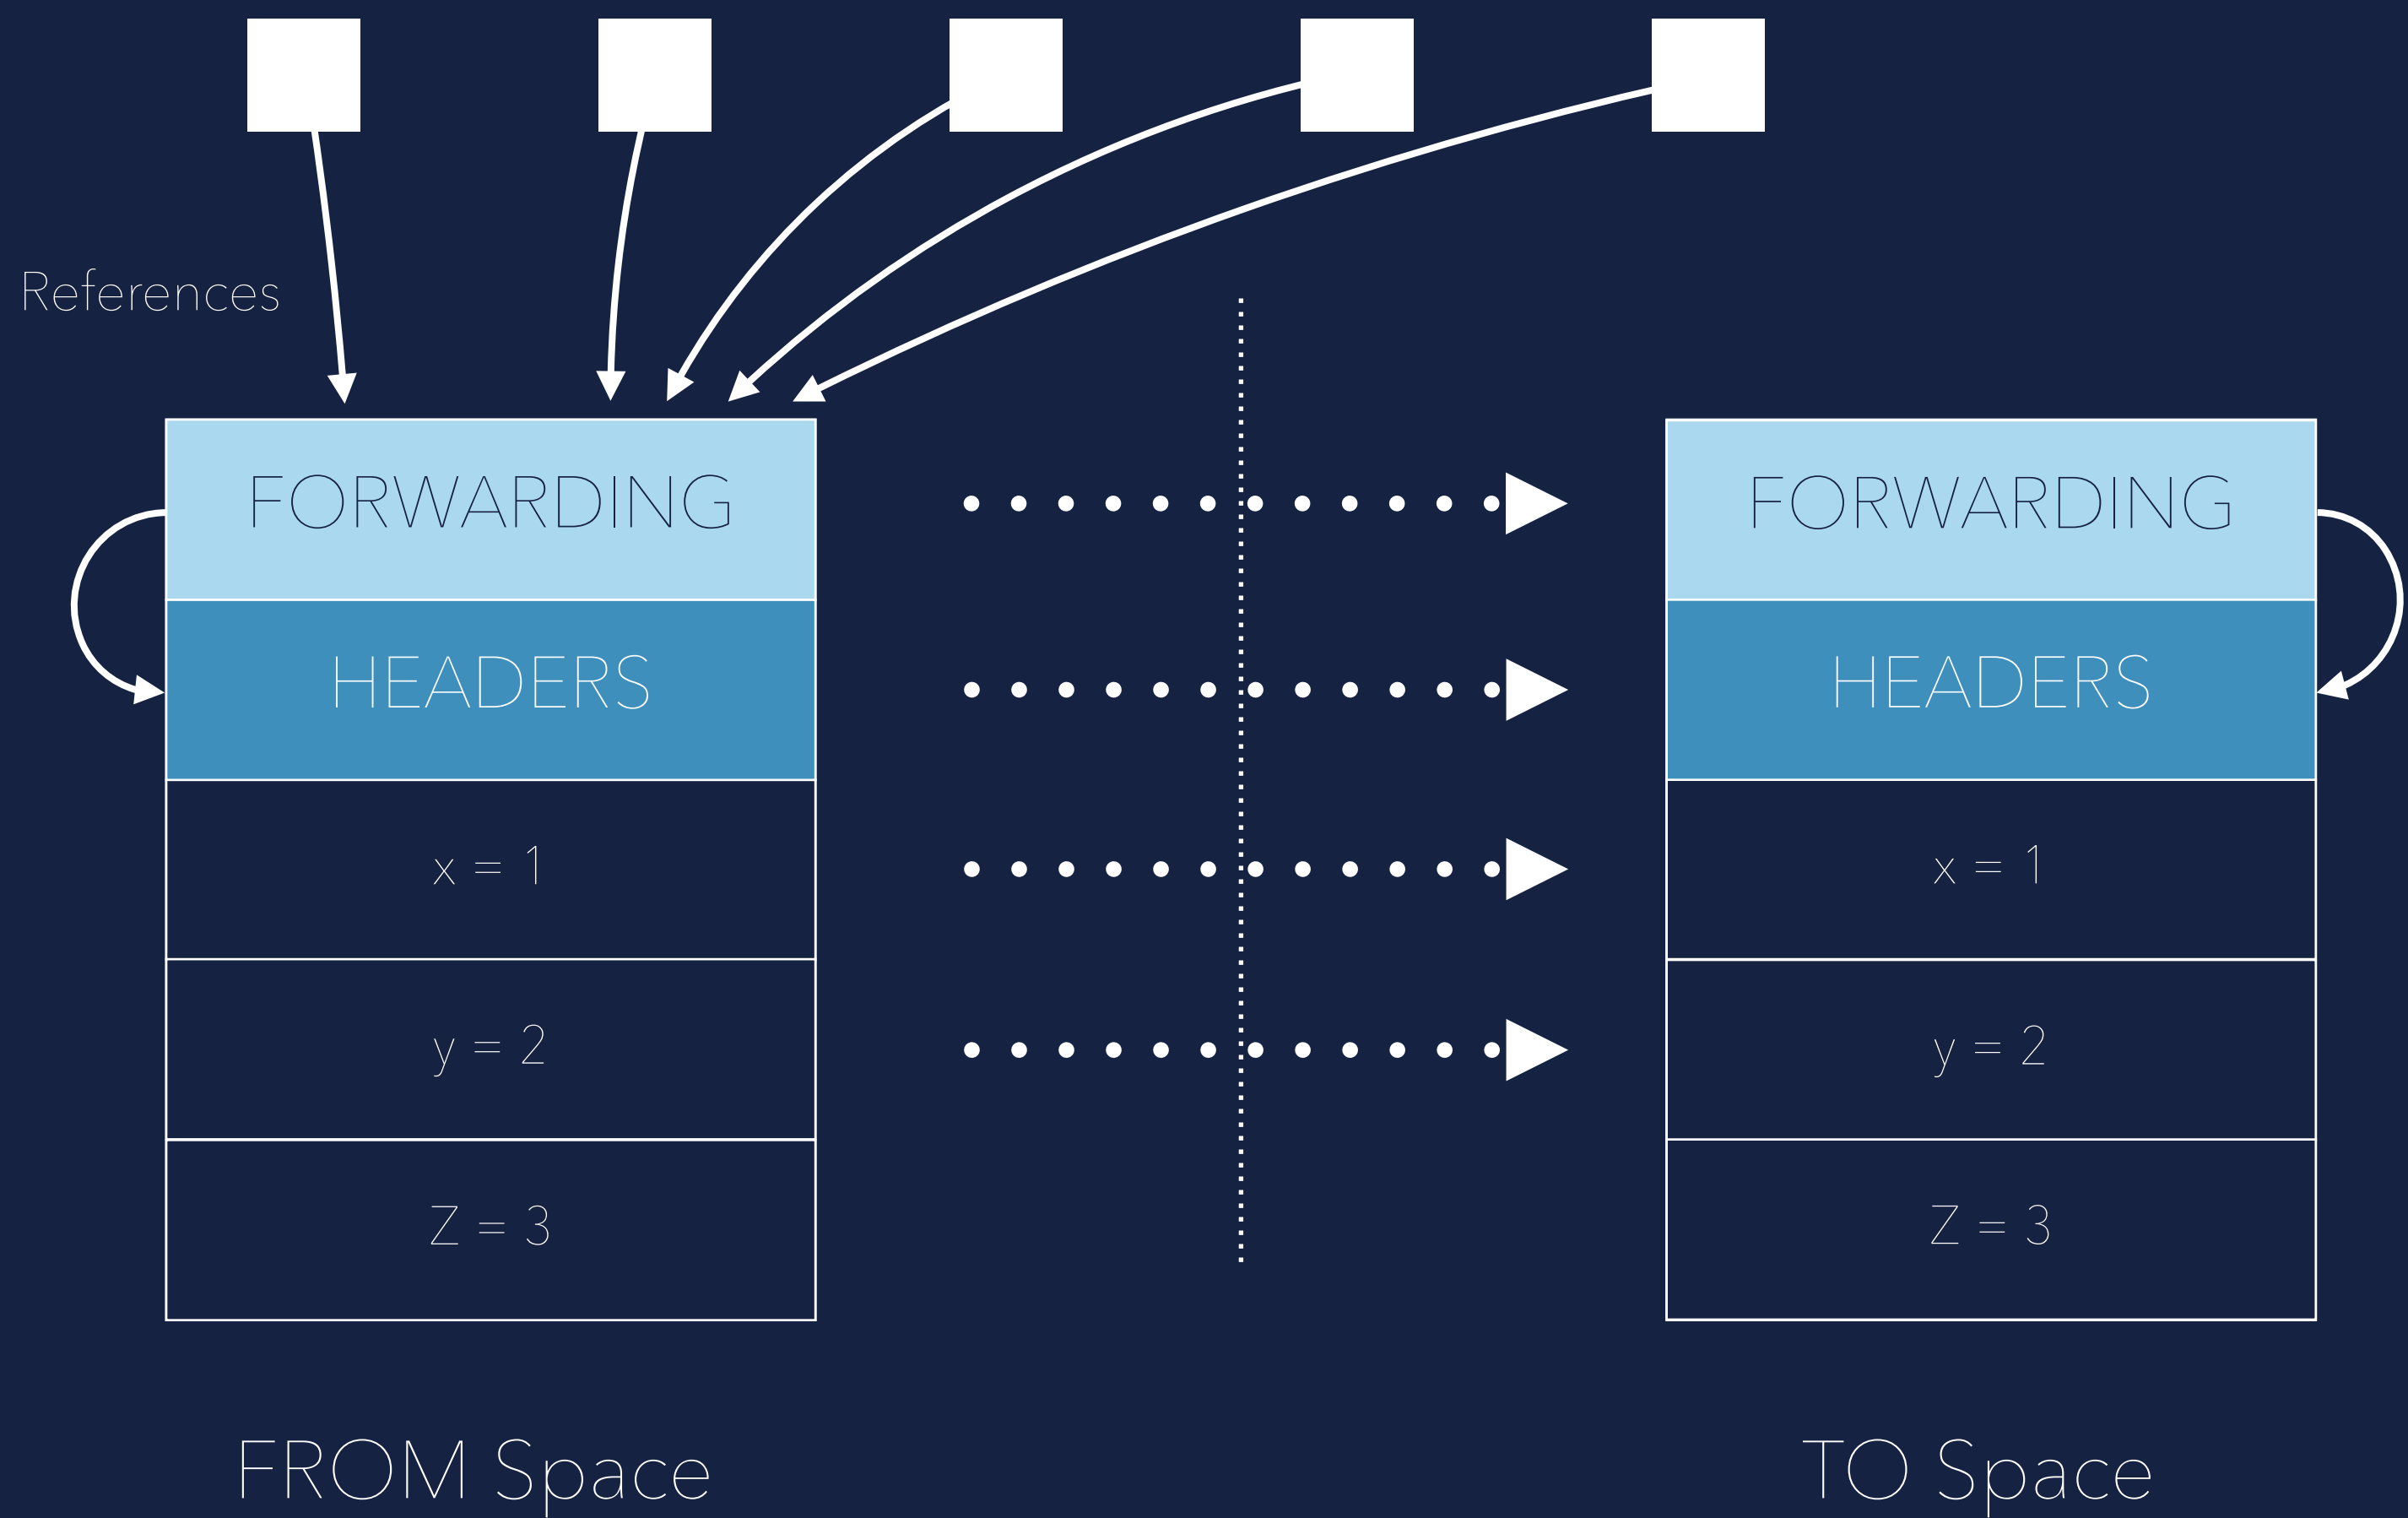
## Concurrent copying



Copy the Object  
(Init forwarding pointer to itself)

# CONCURRENCY IS HARD...

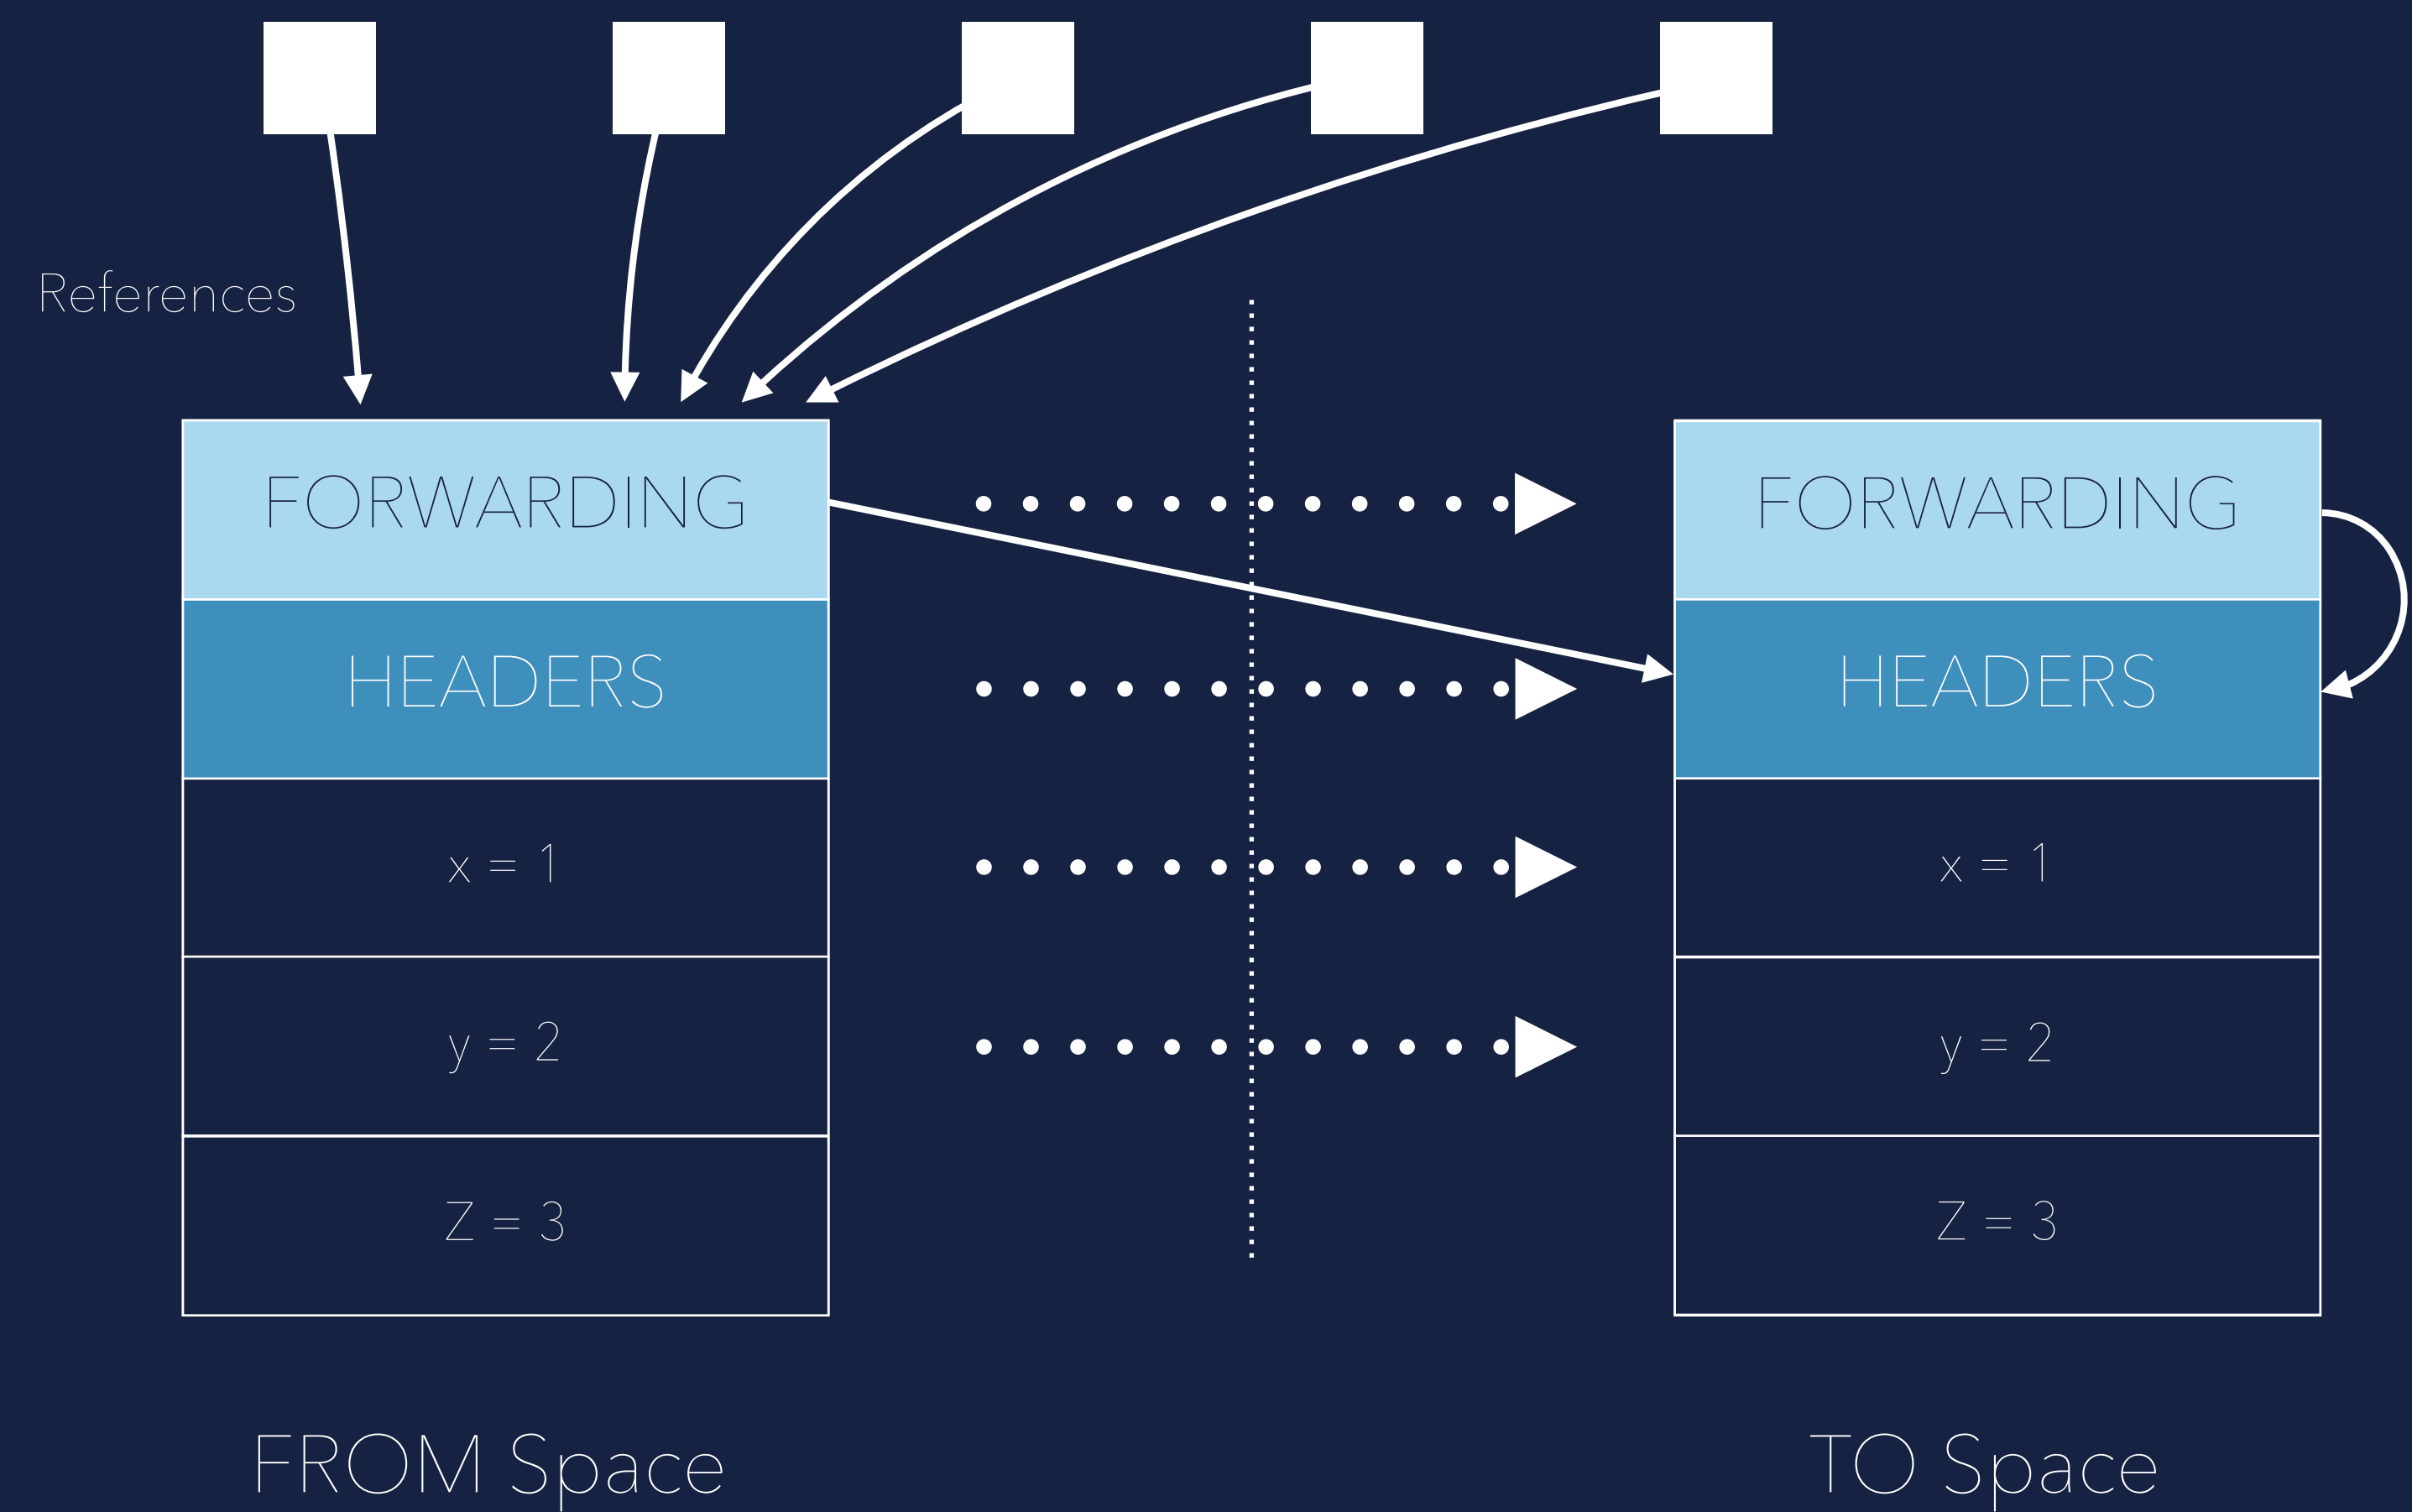
## Concurrent copying



Nobody knows about copy

# CONCURRENCY IS HARD...

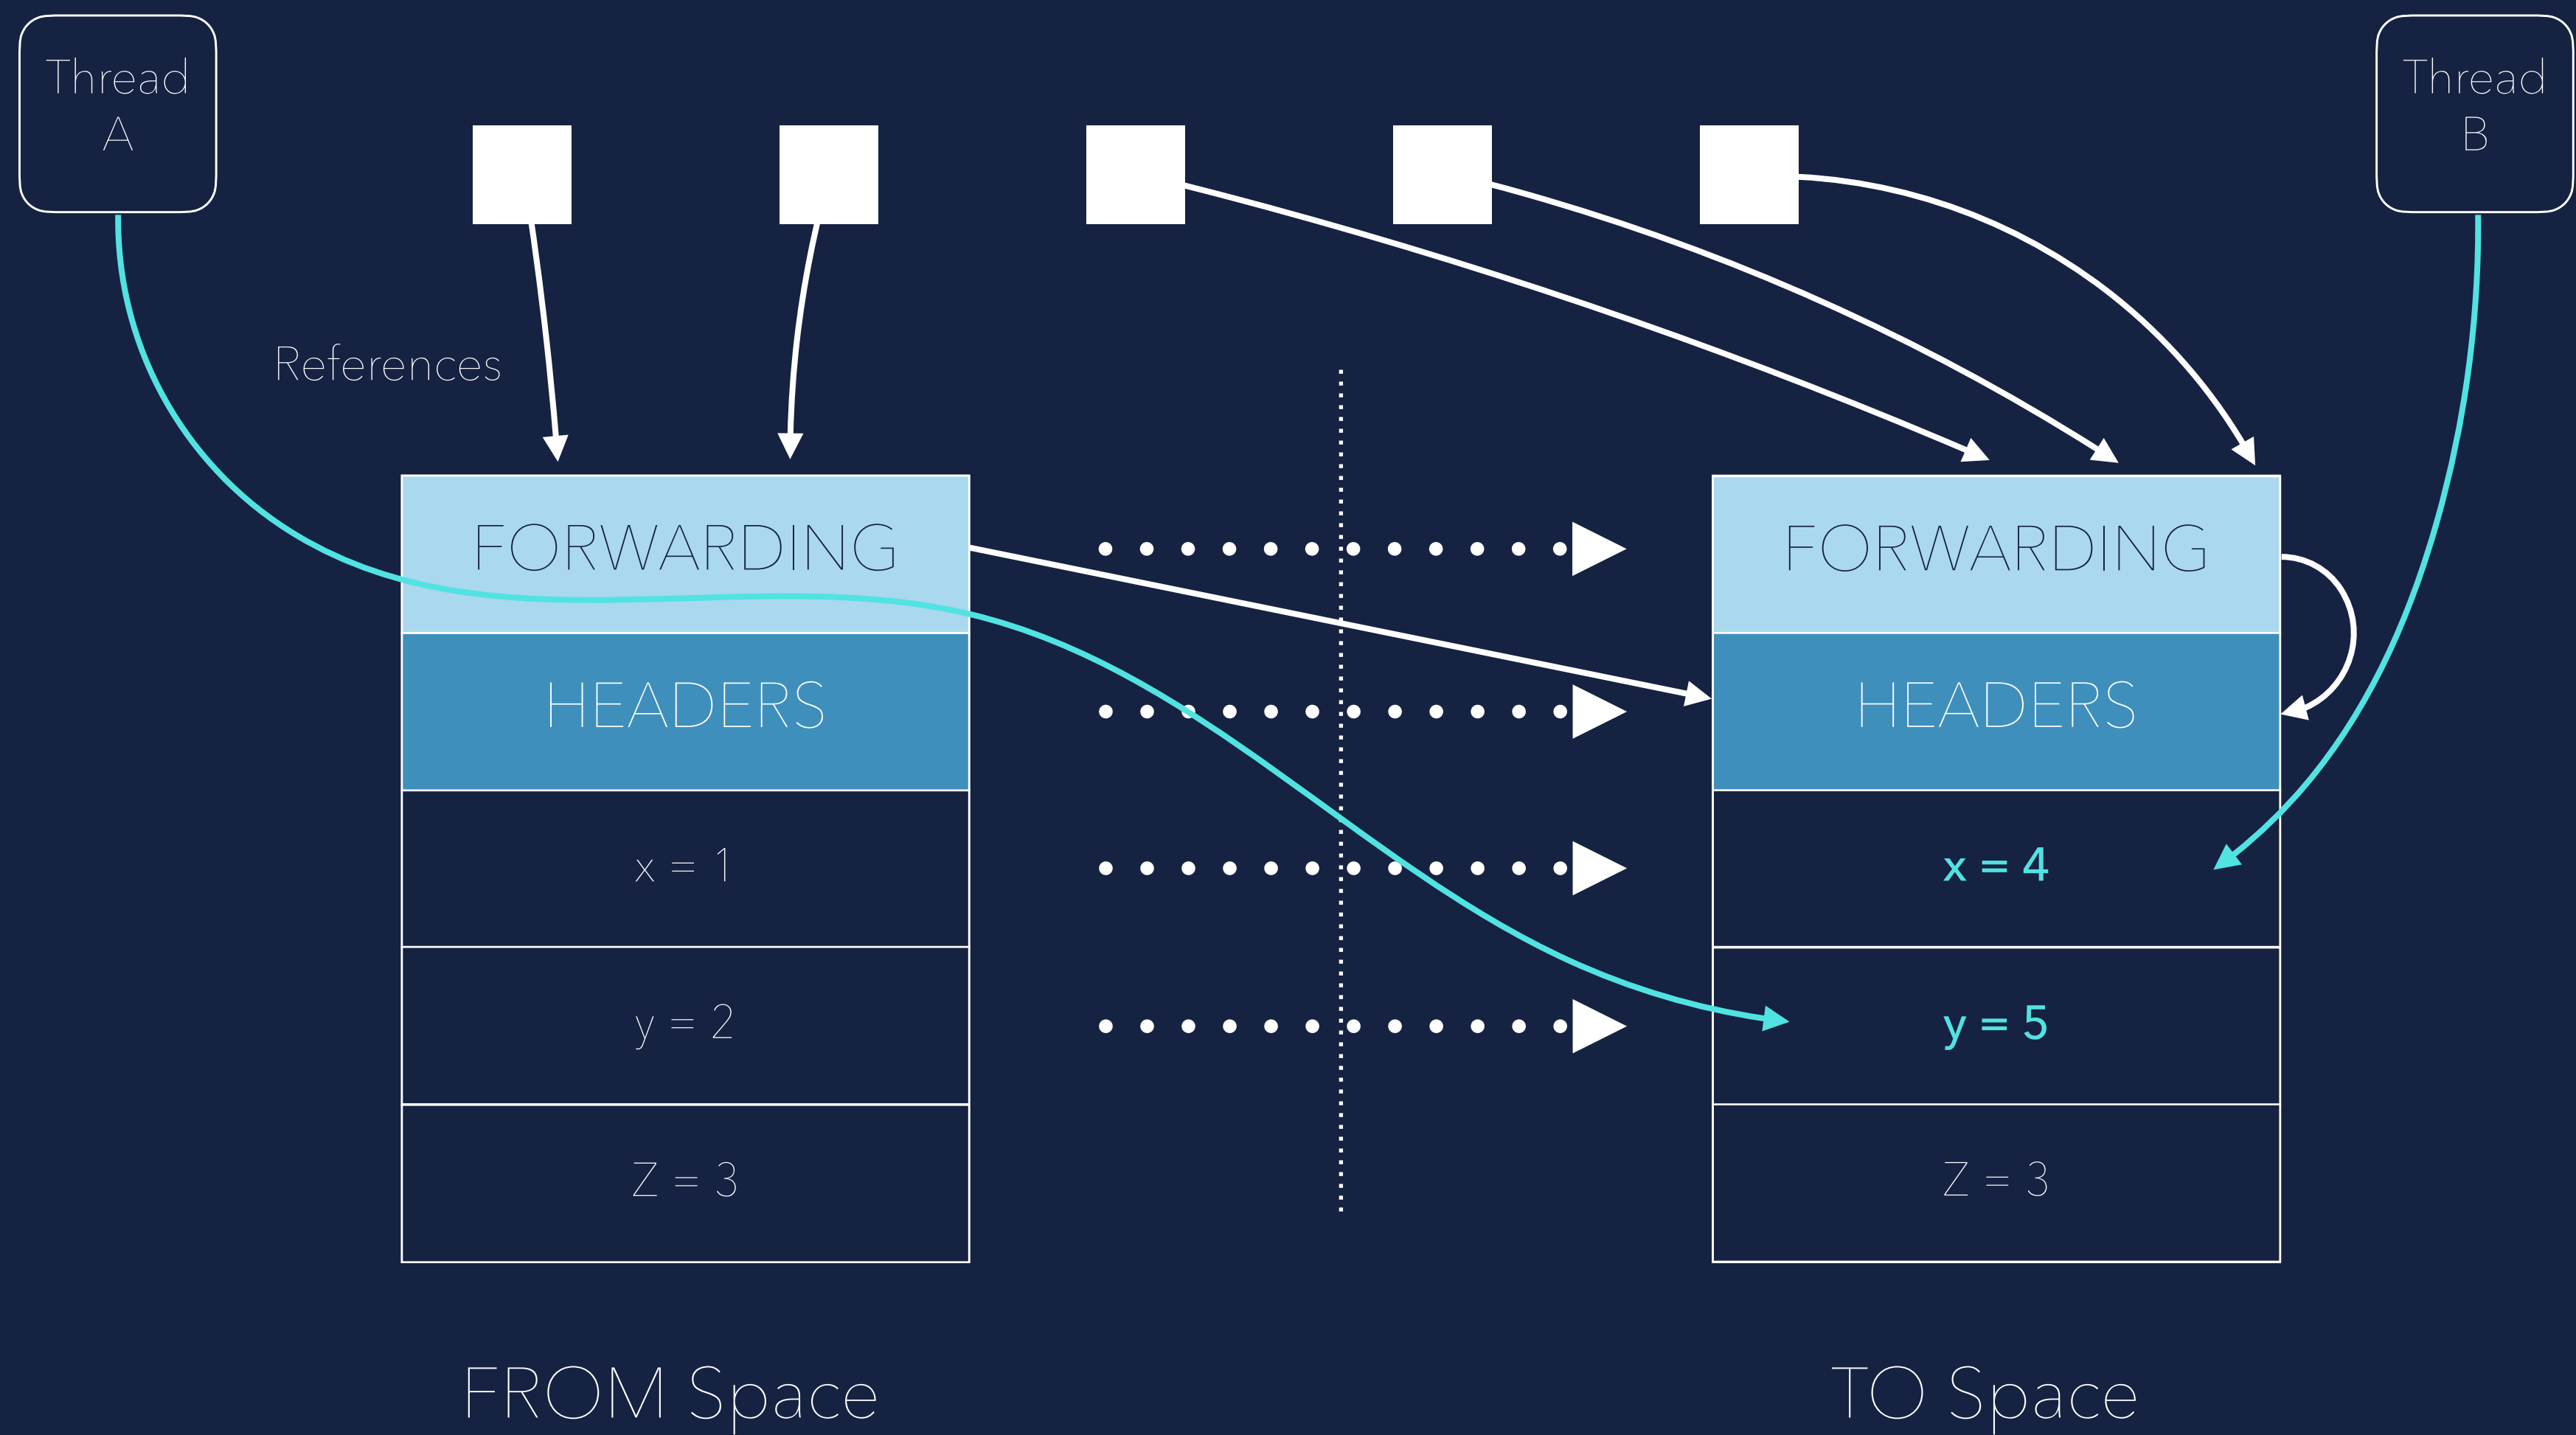
## Concurrent copying



Install forwarding pointer  
of original object  
to new copy

# CONCURRENCY IS HARD...

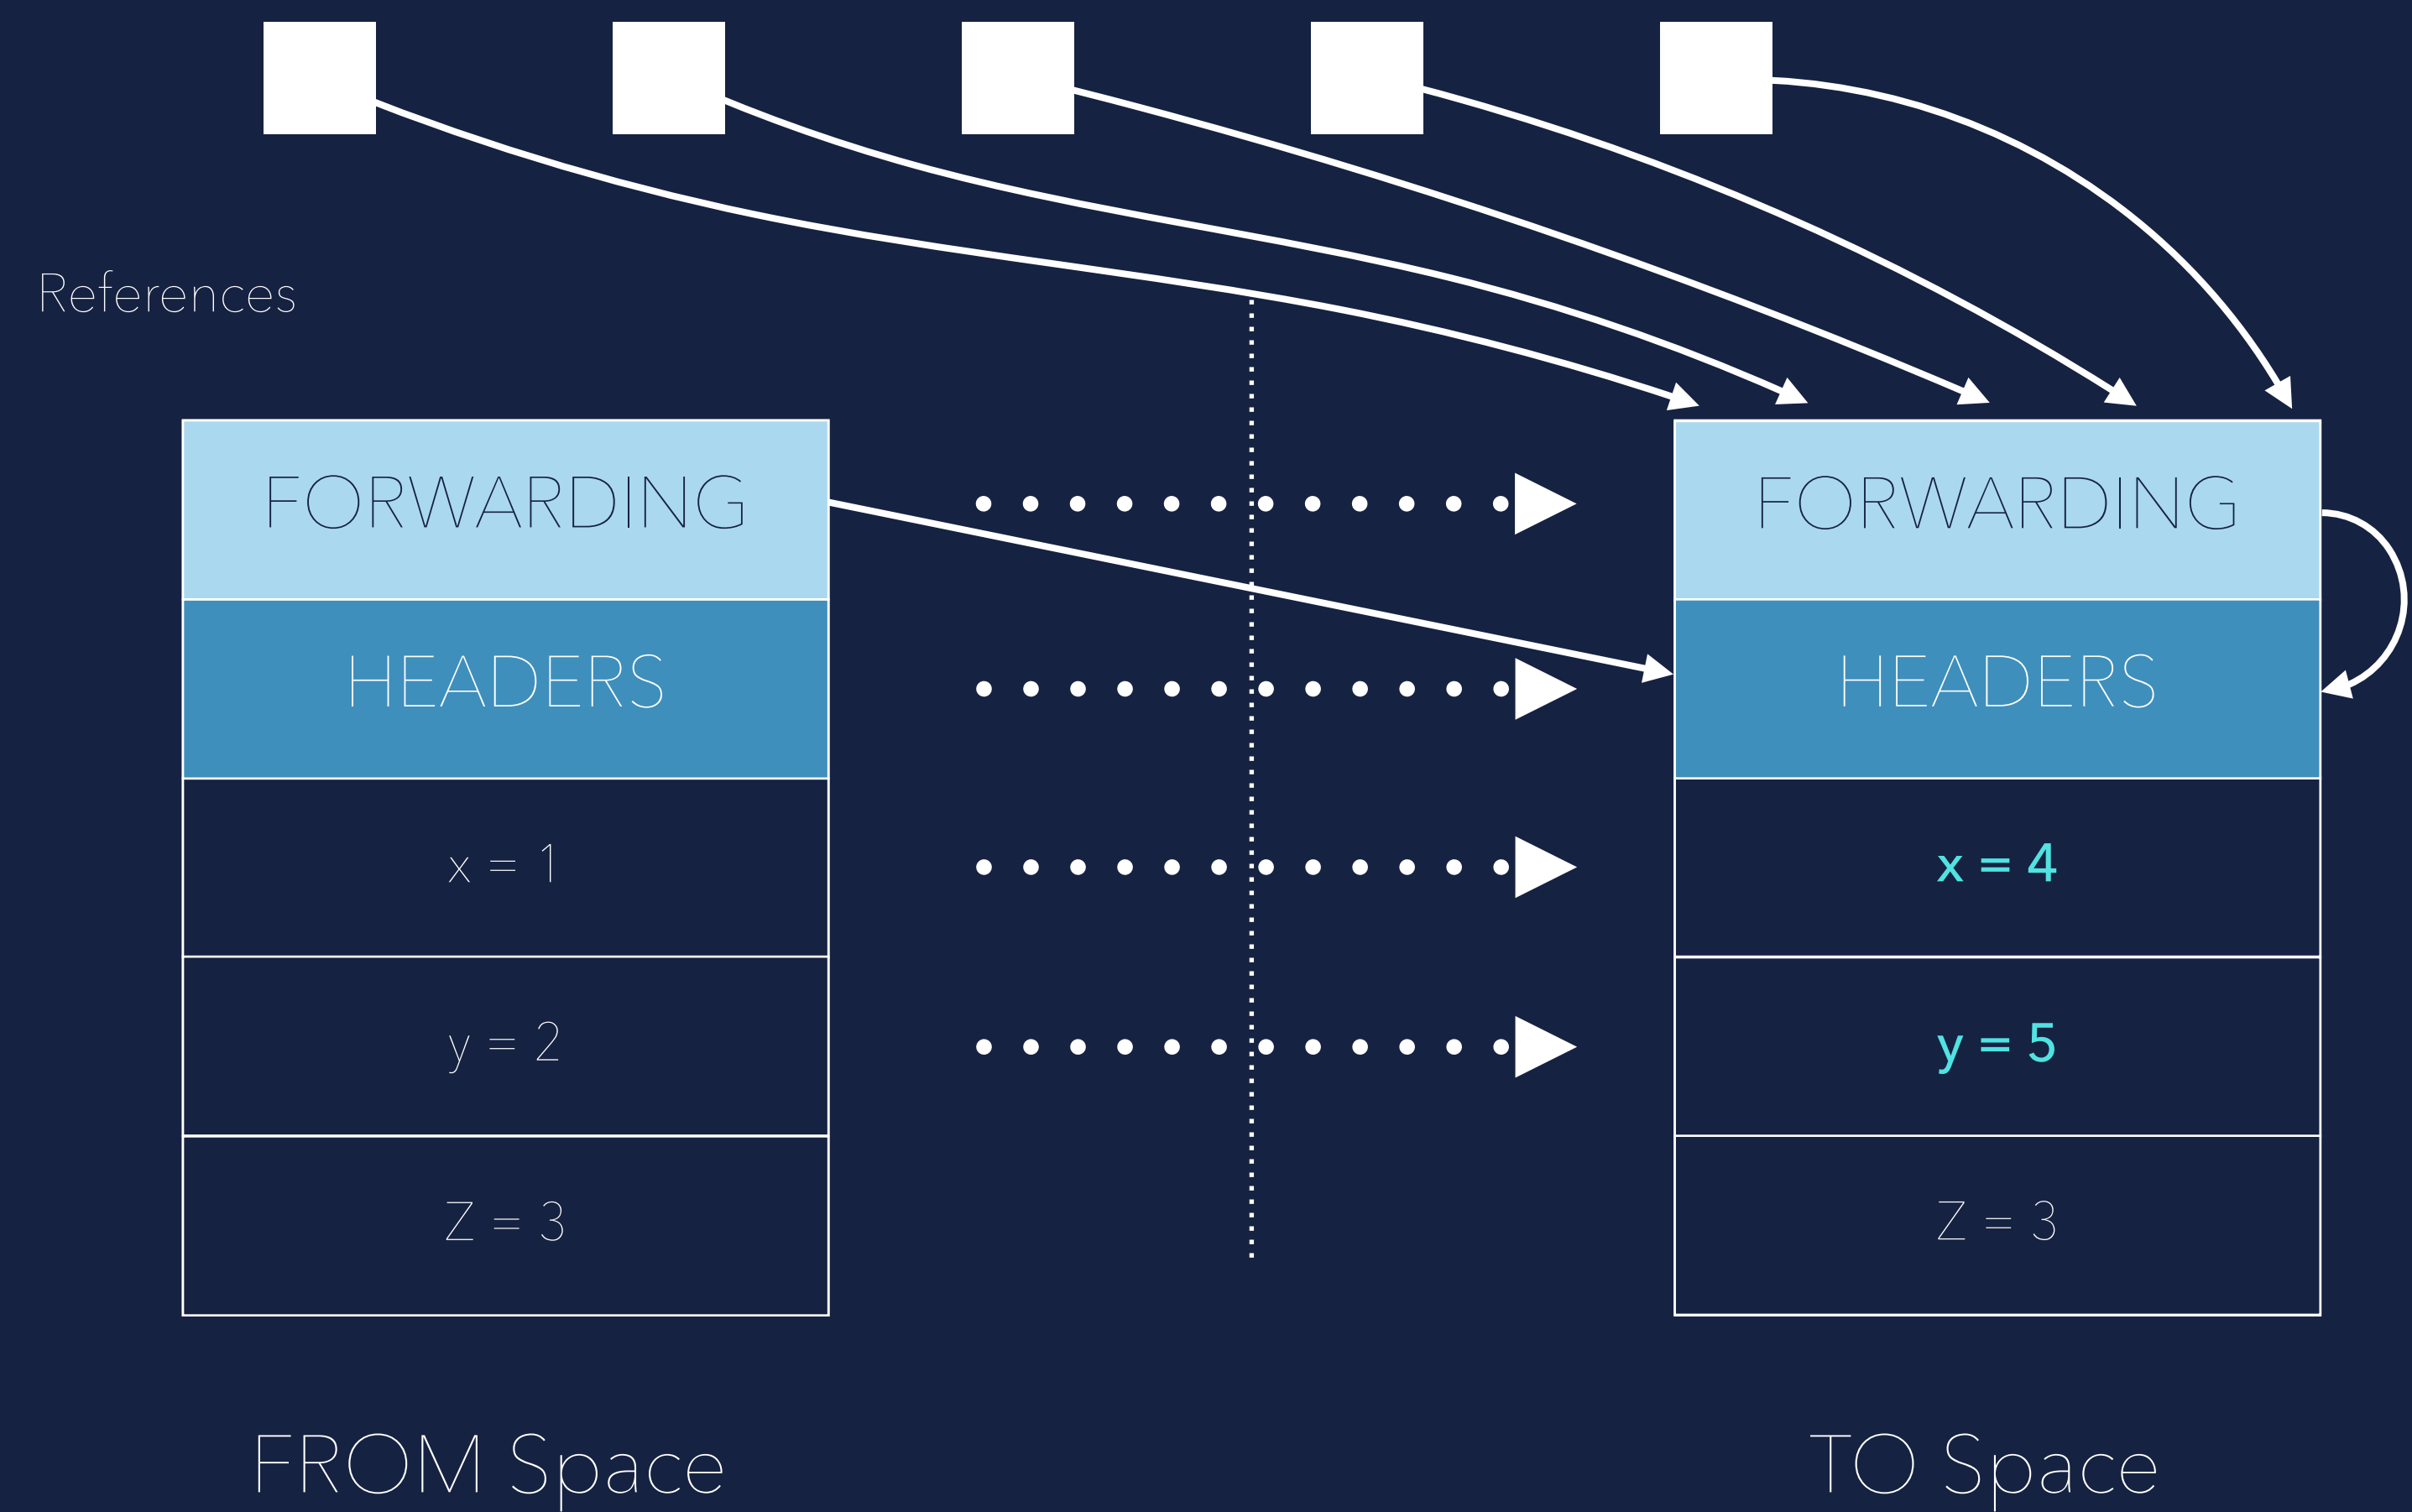
## Concurrent copying



Threads now will always find the right object

# CONCURRENCY IS HARD...

## Concurrent copying

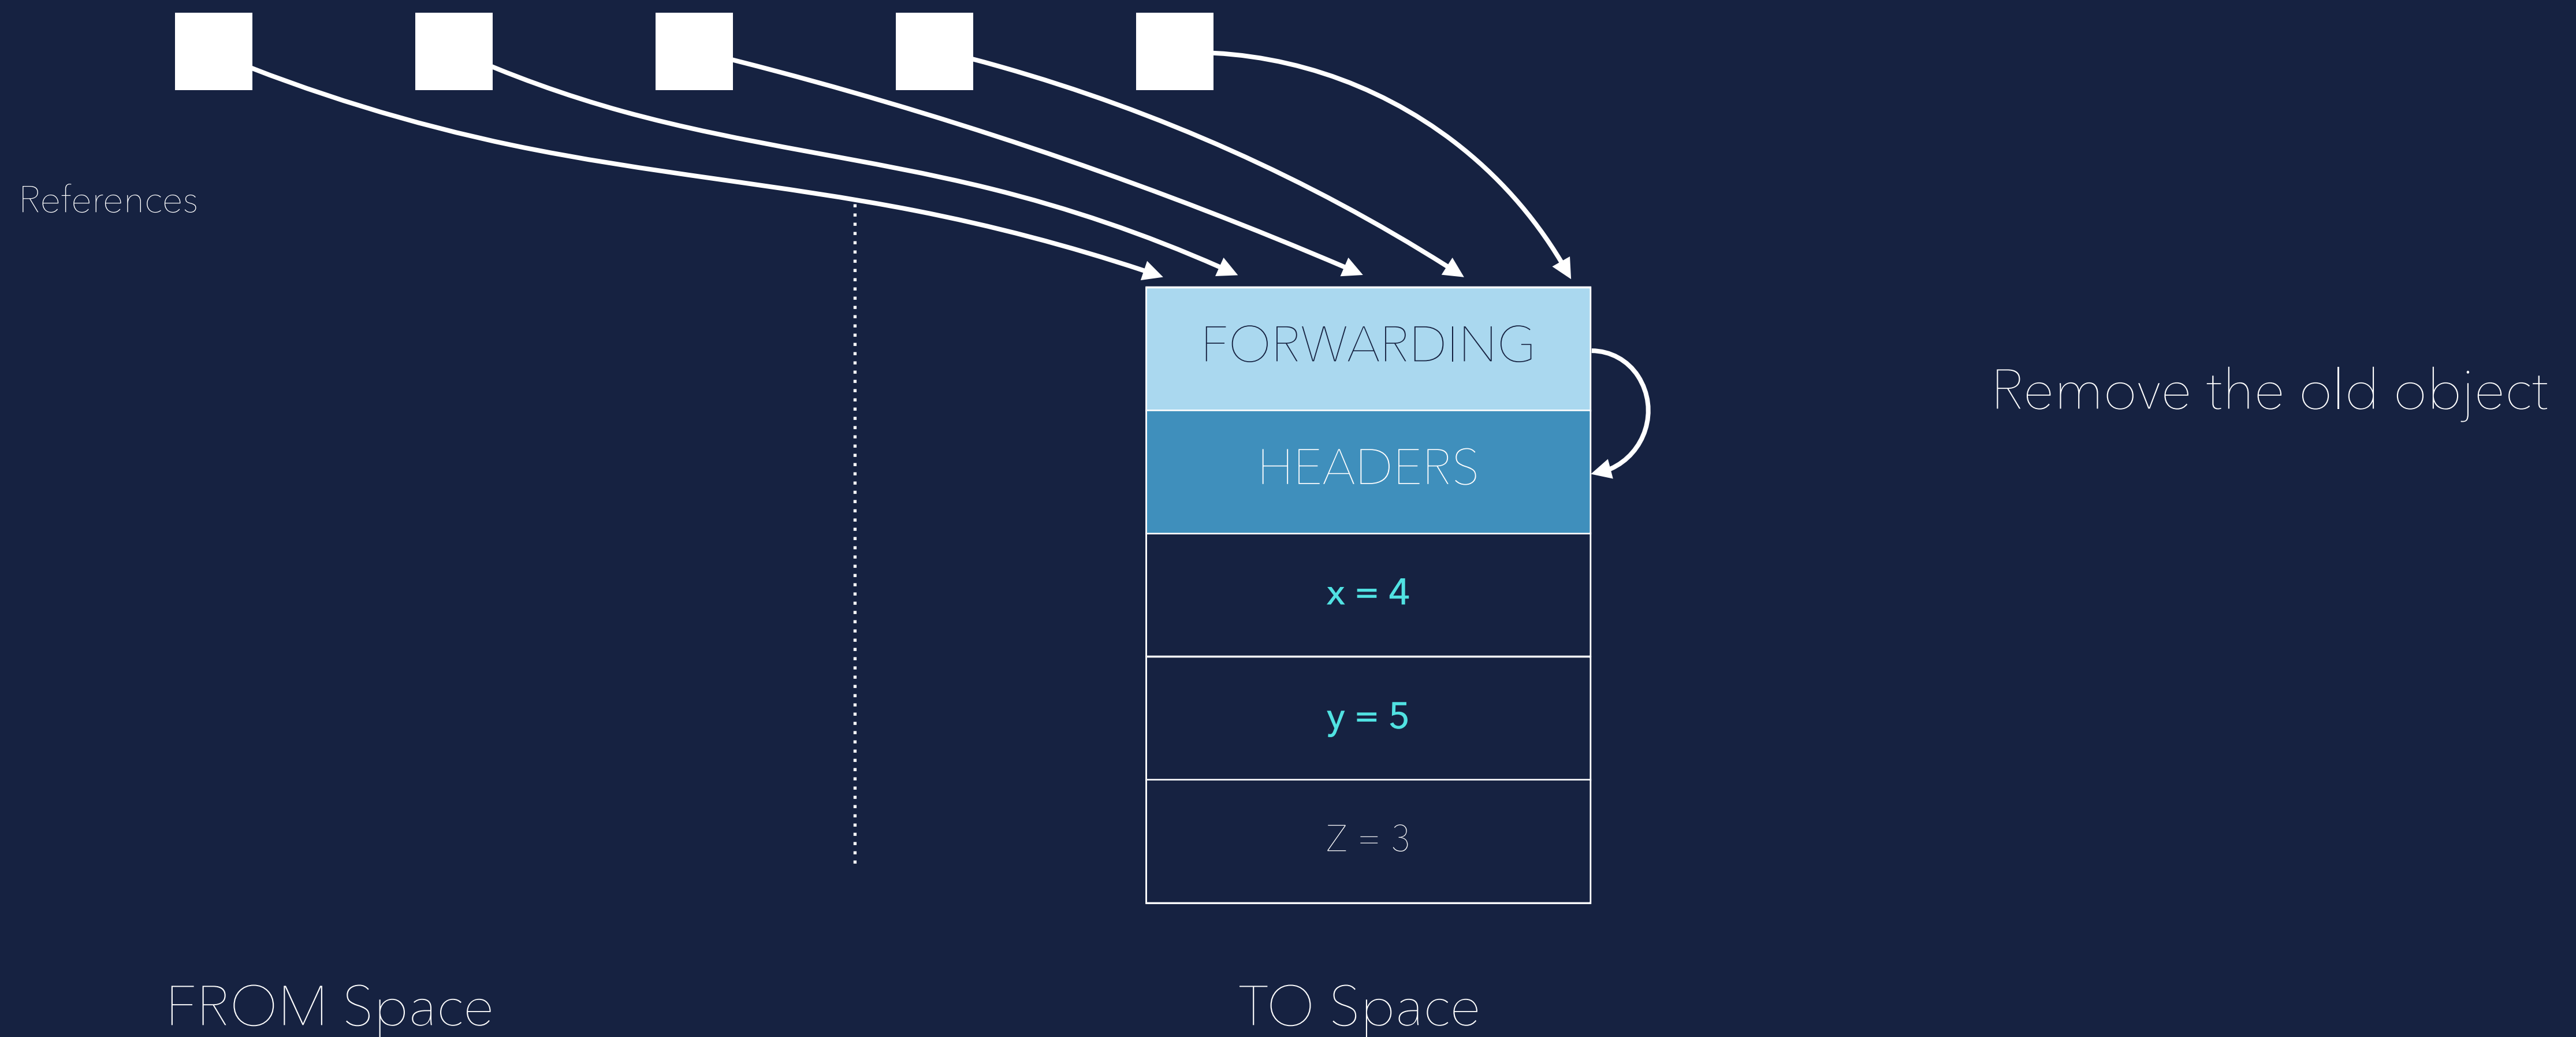


All references are updated



# CONCURRENCY IS HARD...

## Concurrent copying



# COLLECTORS IN THE JVM



# SERIAL

# SERIAL



AVAILABILITY

ALL JDK'S

PARALLEL

NO

CONCURRENT

NO

GENERATIONAL

YES

HEAP SIZE

SMALL - MEDIUM

PAUSE TIMES

LONGER

THROUGHPUT

LOW

LATENCY

HIGHER

CPU OVERHEAD

LOW (1-5%)

## CHOOSE WHEN



Single core systems with small heap (<4GB)



No pause time requirements

## BEST SUITED FOR



Single threaded applications



Development environments



Microservices on small nodes

OS SUPPORT







JVM SWITCH

```
> java -XX:+UseSerialGC
```

# SERIAL



## NOTES

-  Automatically selected if only a single processor is available
-  Automatically selected if the avail. memory less than 1792 MB
-  Young Generation algorithm: Copy Collector
-  Old Generation algorithm: Mark Sweep Compact



# PARALLEL

# PARALLEL



AVAILABILITY

ALL JDK'S

PARALLEL

YES

CONCURRENT

NO

GENERATIONAL

YES

HEAP SIZE

MEDIUM - LARGE

PAUSE TIMES

MODERATE

THROUGHPUT

HIGH

LATENCY

LOWER

CPU OVERHEAD

MODERATE (5-10%)

## CHOOSE WHEN



Multi-core systems with small heap (<4GB)



Peak performance is needed without pause time requirements

## BEST SUITED FOR



Batch processing

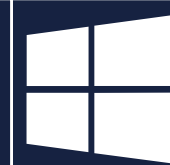


Scientific computing



Data analysis

OS SUPPORT






JVM SWITCH

```
> java -XX:+UseParallelOldGC
```

# PARALLEL



## NOTES

-  Default garbage collector from JDK 5 to JDK 7
-  Young Generation algorithm: Copy Collector
-  Old Generation algorithm: Mark Sweep Compact





# CMS

Concurrent Mark and Sweep

**DEPRECATED**

AVAILABILITY

JDK 1.4 - 13

PARALLEL

YES

CONCURRENT

PARTIALLY

GENERATIONAL

YES

HEAP SIZE

MEDIUM - LARGE

PAUSE TIMES

MODERATE

THROUGHPUT

MODERATE

LATENCY

MODERATE

CPU OVERHEAD

MODERATE (5-15%)

## CHOOSE WHEN



Response time is more important than throughput



Pause time must be kept shorter than 1 sec

## BEST SUITED FOR

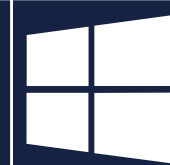


Web applications



Mediums sized enterprise systems






OS SUPPORT



JVM SWITCH

```
> java -XX:+UseConcMarkSweepGC
```

## NOTES

-  Deprecated as of JDK 9
-  Removed from JDK 14
-  Young Generation algorithm: Copy Collector
-  Old Generation algorithm: Concurrent Mark and Sweep
-  Full GC algorithm: Mark Sweep Compact



# G1

Garbage First



# Heap-Layout

Region size 1 - 32 MB


Max no. of region <= 2048

Heap	Region
< 4 GB	1 MB
< 8 GB	2 MB
< 16 GB	4 MB
< 32 GB	8 MB
< 64 GB	16 MB
> 64 GB	32 MB

Example 8GB Heap:

8 GB Heap = 8192 MB

8192 MB / 2048 = 4 MB region size





# Heap-Layout

Region size 1 - 32 MB

Max no. of region <= 2048

Heap	Region
< 4 GB	1 MB
< 8 GB	2 MB
< 16 GB	4 MB
< 32 GB	8 MB
< 64 GB	16 MB
> 64 GB	32 MB

Example 8GB Heap:

8 GB Heap = 8192 MB

8192 MB / 2048 = 4 MB region size



# Heap-Layout

Region size 1 - 32 MB



Max no. of region <= 2048

Heap	Region
< 4 GB	1 MB
< 8 GB	2 MB
< 16 GB	4 MB
< 32 GB	8 MB
< 64 GB	16 MB
> 64 GB	32 MB


Example 8GB Heap:

8 GB Heap = 8192 MB


8192 MB / 2048 = 4 MB region size




Unassigned region




Eden region



Survivor region



Tenured region



Humongous region  
(> 0.5 \* Region size)

Young Gen  
5 - 60%

Old Gen

Example:  
6 Eden Regions  
3 Survivor Regions

2 Regions with most garbage will  
be collected/promoted



AVAILABILITY

JDK 7U4+

PARALLEL

YES

CONCURRENT

PARTIALLY

GENERATIONAL

YES

HEAP SIZE

MEDIUM - LARGE

PAUSE TIMES

SHORT - MEDIUM

THROUGHPUT

HIGH

LATENCY

LOWER

CPU OVERHEAD

MODERATE (5-15%)

## CHOOSE WHEN



Response time is more important than throughput



Pause time must be kept shorter than 1 sec

## BEST SUITED FOR



Mixed workloads




Large sized enterprise systems



Responsive in medium to large heaps

OS SUPPORT







JVM SWITCH

> `java -XX:+UseG1GC`



## NOTES

-  Default collector from JDK 9 onwards
-  Young Generation algorithm: Evacuating Collector (Mark and Compact)
-  Old Generation algorithm: Concurrent Mark and Compact
-  Full GC algorithm: Mark and Compact





# EPSILON

# EPSILON



AVAILABILITY	JDK 11+
PARALLEL	-
CONCURRENT	-
GENERATIONAL	-
HEAP SIZE	-
PAUSE TIMES	-
THROUGHPUT	-
LATENCY	-
CPU OVERHEAD	VERY LOW

## CHOOSE WHEN

-  Testing performance or memory pressure
-  Highest performance is needed and nearly no garbage is created

## BEST SUITED FOR

-  Extremely short lived jobs
-  Last drop latency improvements
-  Last drop throughput improvements

OS SUPPORT



JVM SWITCH

> `java -XX:+UseEpsilonGC`

## NOTES

 Only in builds of OpenJDK






# SHENANDOAH

# SHENANDOAH






AVAILABILITY	JDK 11.0.9+
PARALLEL	YES
CONCURRENT	FULLY
GENERATIONAL	NO
HEAP SIZE	MEDIUM - LARGE
PAUSE TIMES	SHORT
THROUGHPUT	VERY HIGH
LATENCY	VERY LOW
CPU OVERHEAD	MODERATE (10-20%)

## CHOOSE WHEN

-  Response time is a high priority
-  Using a very large heap (100GB+)
-  Predictable response times needed

## BEST SUITED FOR

-  Latency sensitive applications
-  Large scale systems
-  Highly concurrent applications

OS SUPPORT	  
------------	---

JVM SWITCH	> <code>java -XX:+UseShenandoahGC</code>
------------	--

# SHENANDOAH



## NOTES

- 🗑 Not available in Oracle JDK
- 🗑 A bit reduced throughput due to concurrent GC
- 🗑 Makes use of new barrier concept, load reference barrier



# ZGC


Zero Garbage Collector





Heap-Layout


HEAP



EMPTY REGION



YOUNG GEN REGION



OLD GEN REGION

AVAILABILITY

JDK 15 / 21+

PARALLEL

YES

CONCURRENT

FULLY

GENERATIONAL

NO / YES

HEAP SIZE

LARGE

PAUSE TIMES

SHORT

THROUGHPUT

VERY HIGH


LATENCY

VERY LOW




CPU OVERHEAD

MODERATE (10-20%)


## CHOOSE WHEN

-  Response time is a high priority
-  Using a very large heap (100GB+)
-  Predictable response times needed

## BEST SUITED FOR

-  Low latency sensitive applications
-  Large scale systems
-  Highly concurrent applications

OS SUPPORT



JVM SWITCH

```
> java -XX:+UseZGC -XX:+ZGenerational*
```

\* Not needed in the future, because generational ZGC will become the default

## NOTES

-  Will become the default collector in the future
-  Non-generational version will be deprecated



# C4

Concurrent Continues Compacting Collector

## NOTES

 Part of Azul Zing JVM

 Makes use of Loaded Value Barrier everywhere  
(Test + Jump which only takes 1 cpu cycle -> very fast)

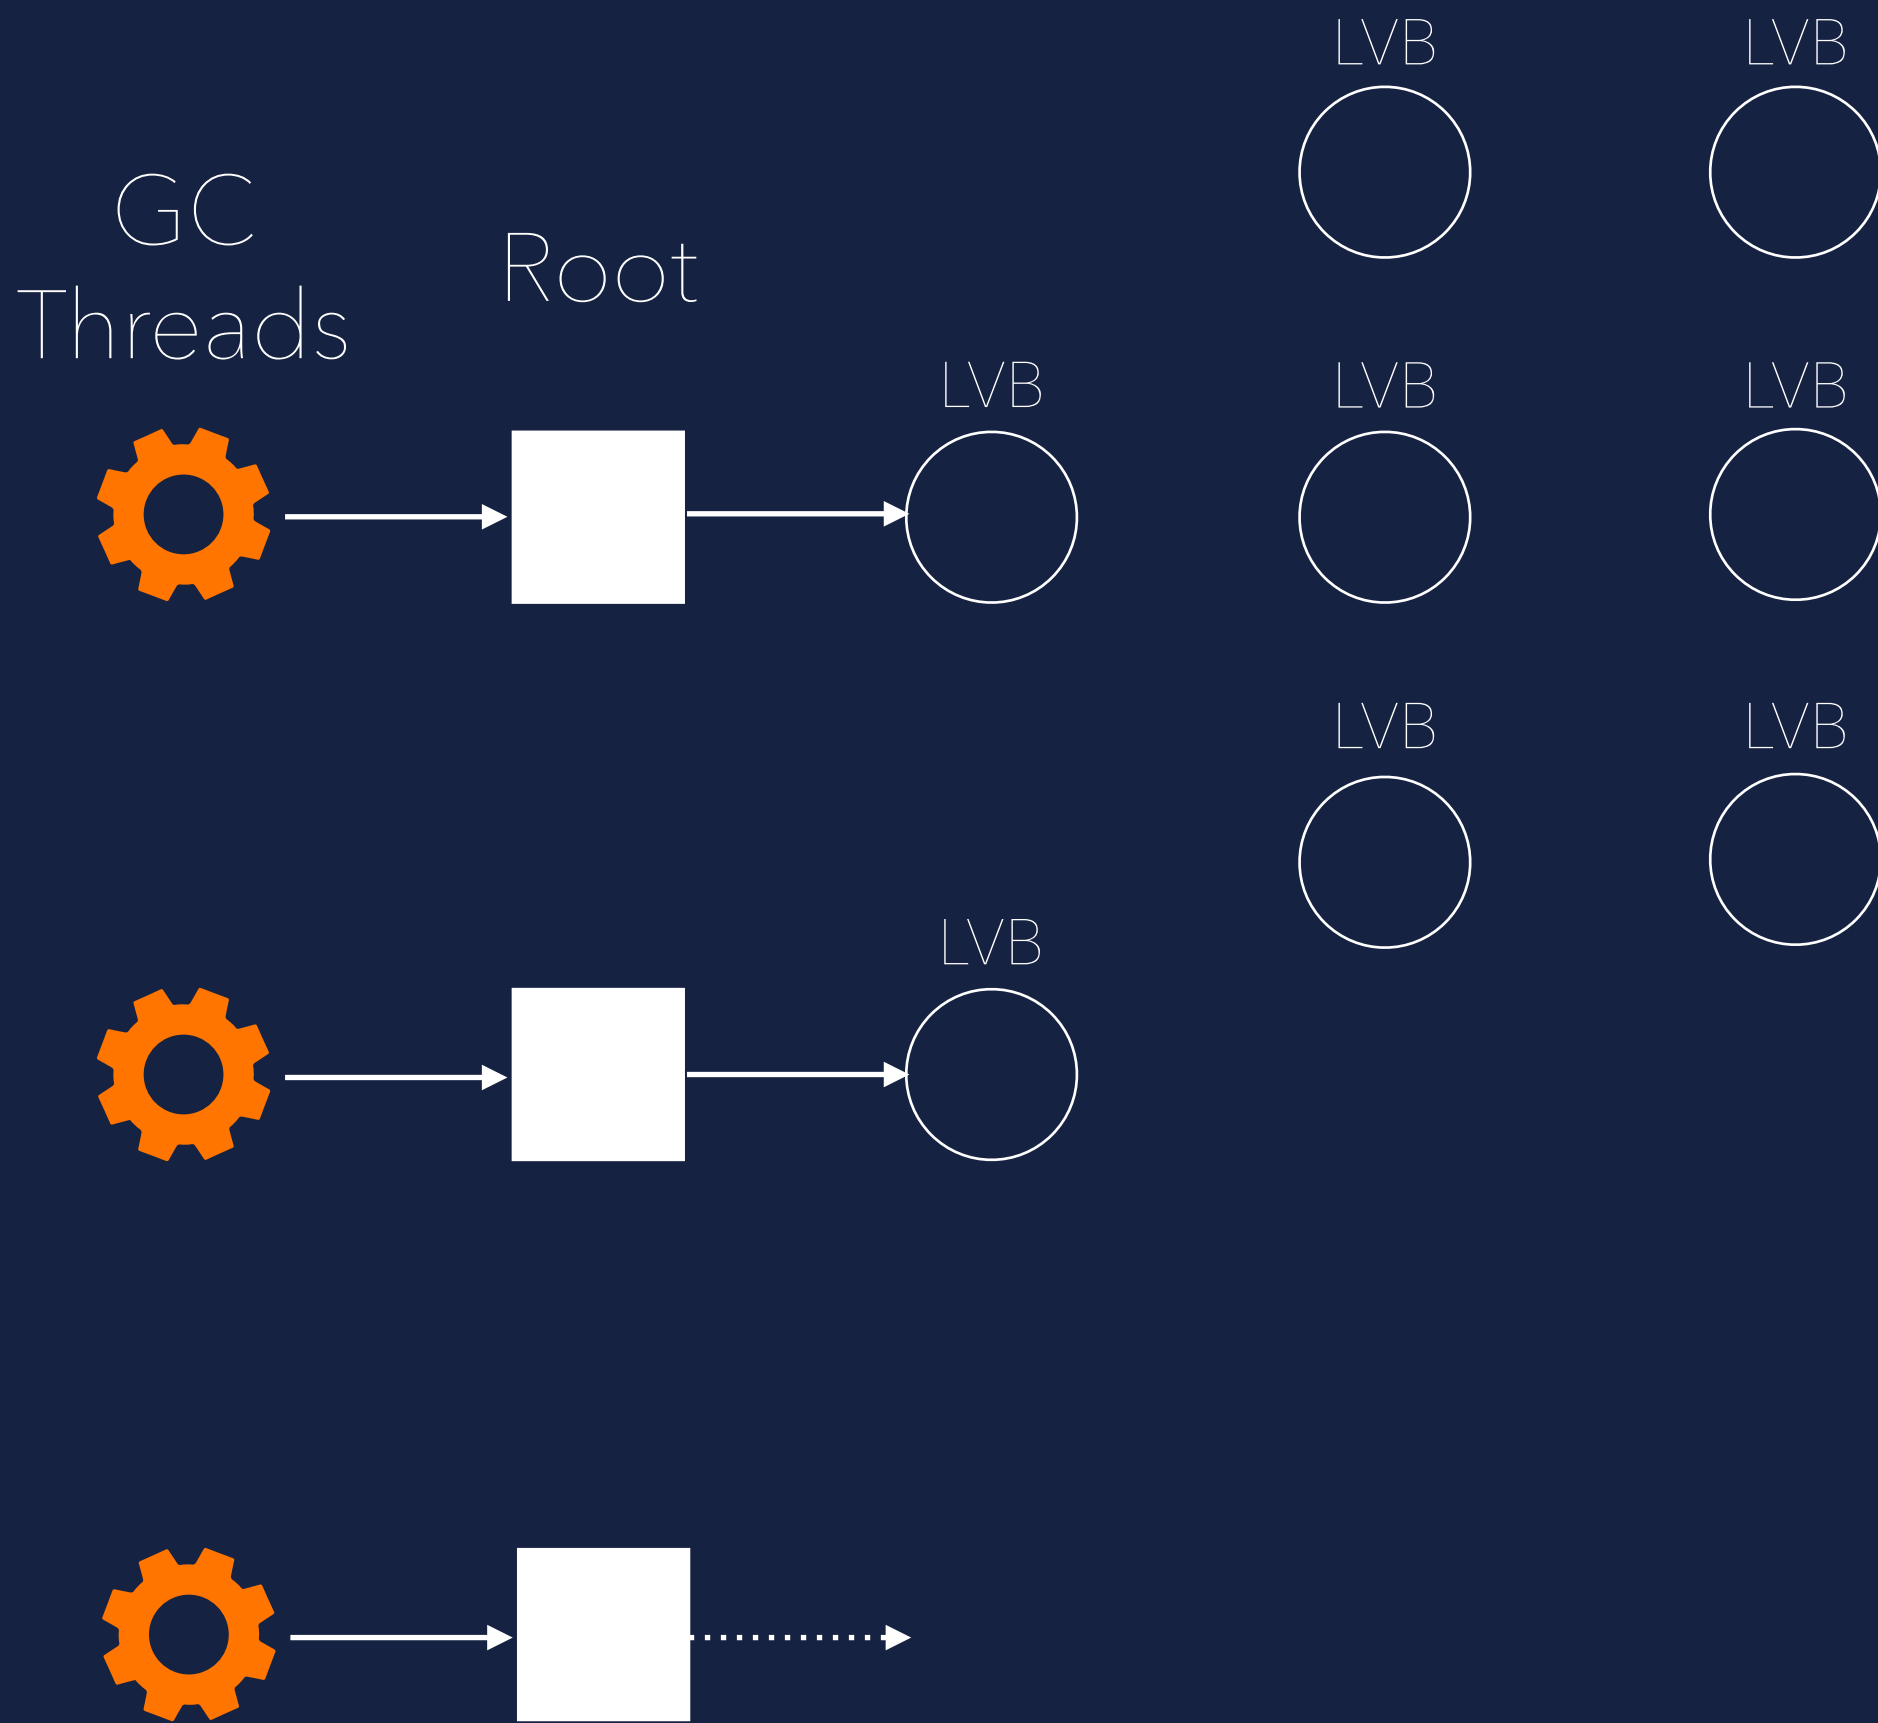
 Best performance by using Transparent Huge Pages  
(Normal page size 4kB, THP size 2MB)

C4

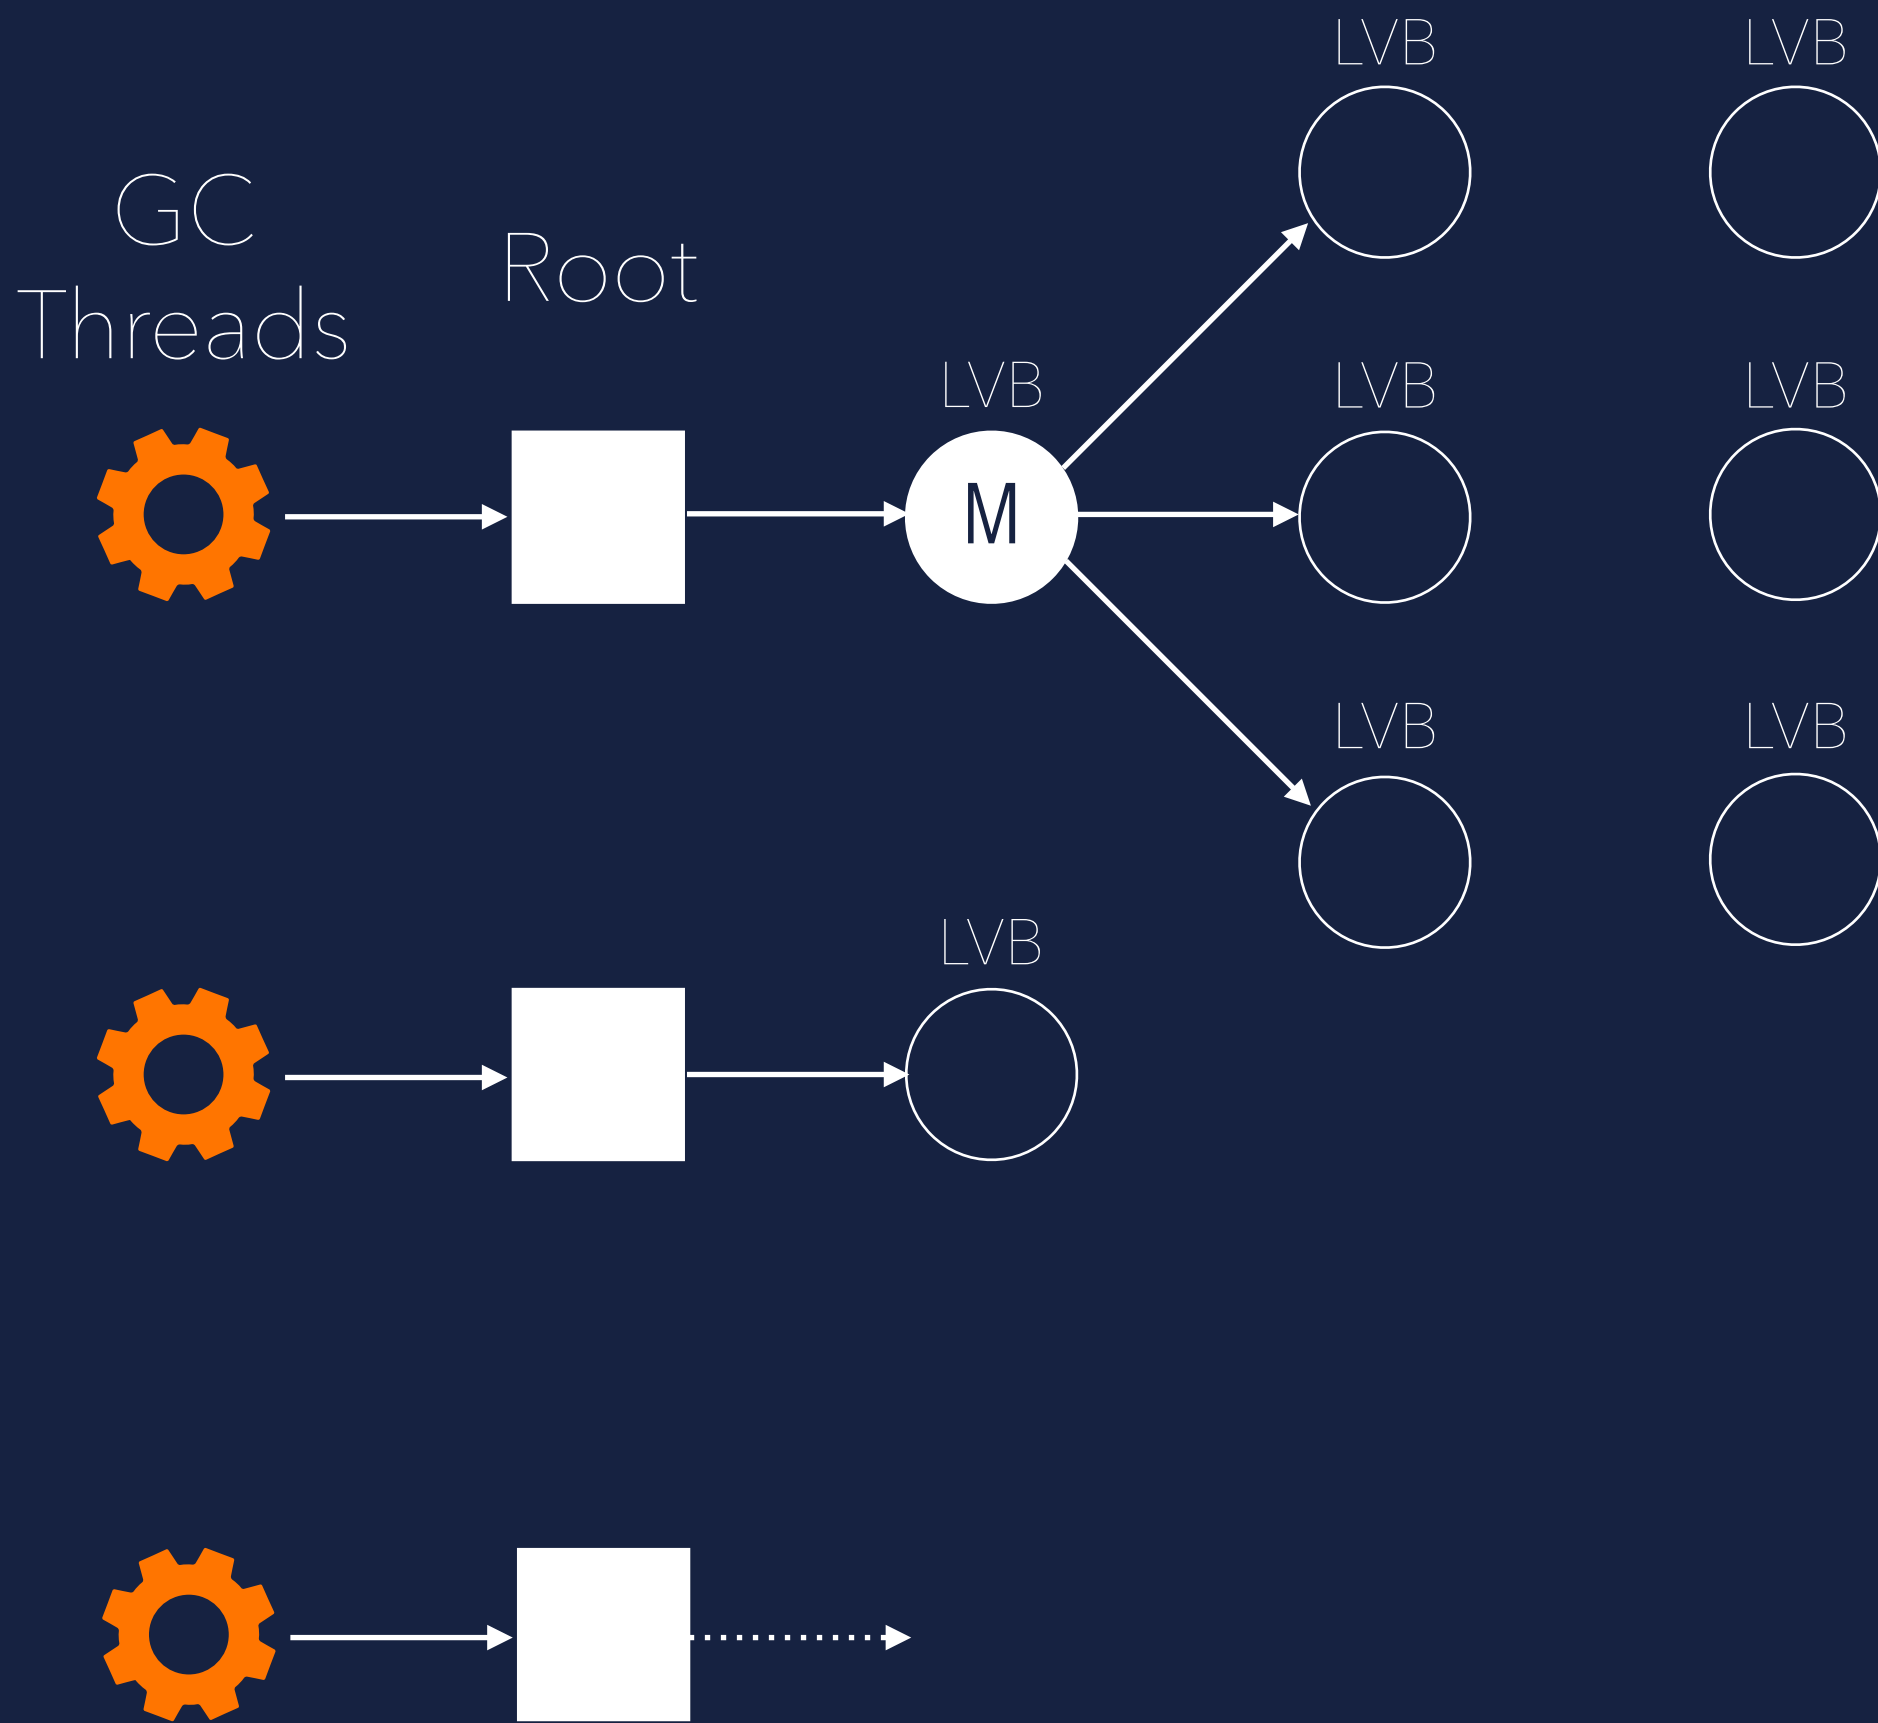


# MARKING PHASE

## Marking Phase




## Marking Phase

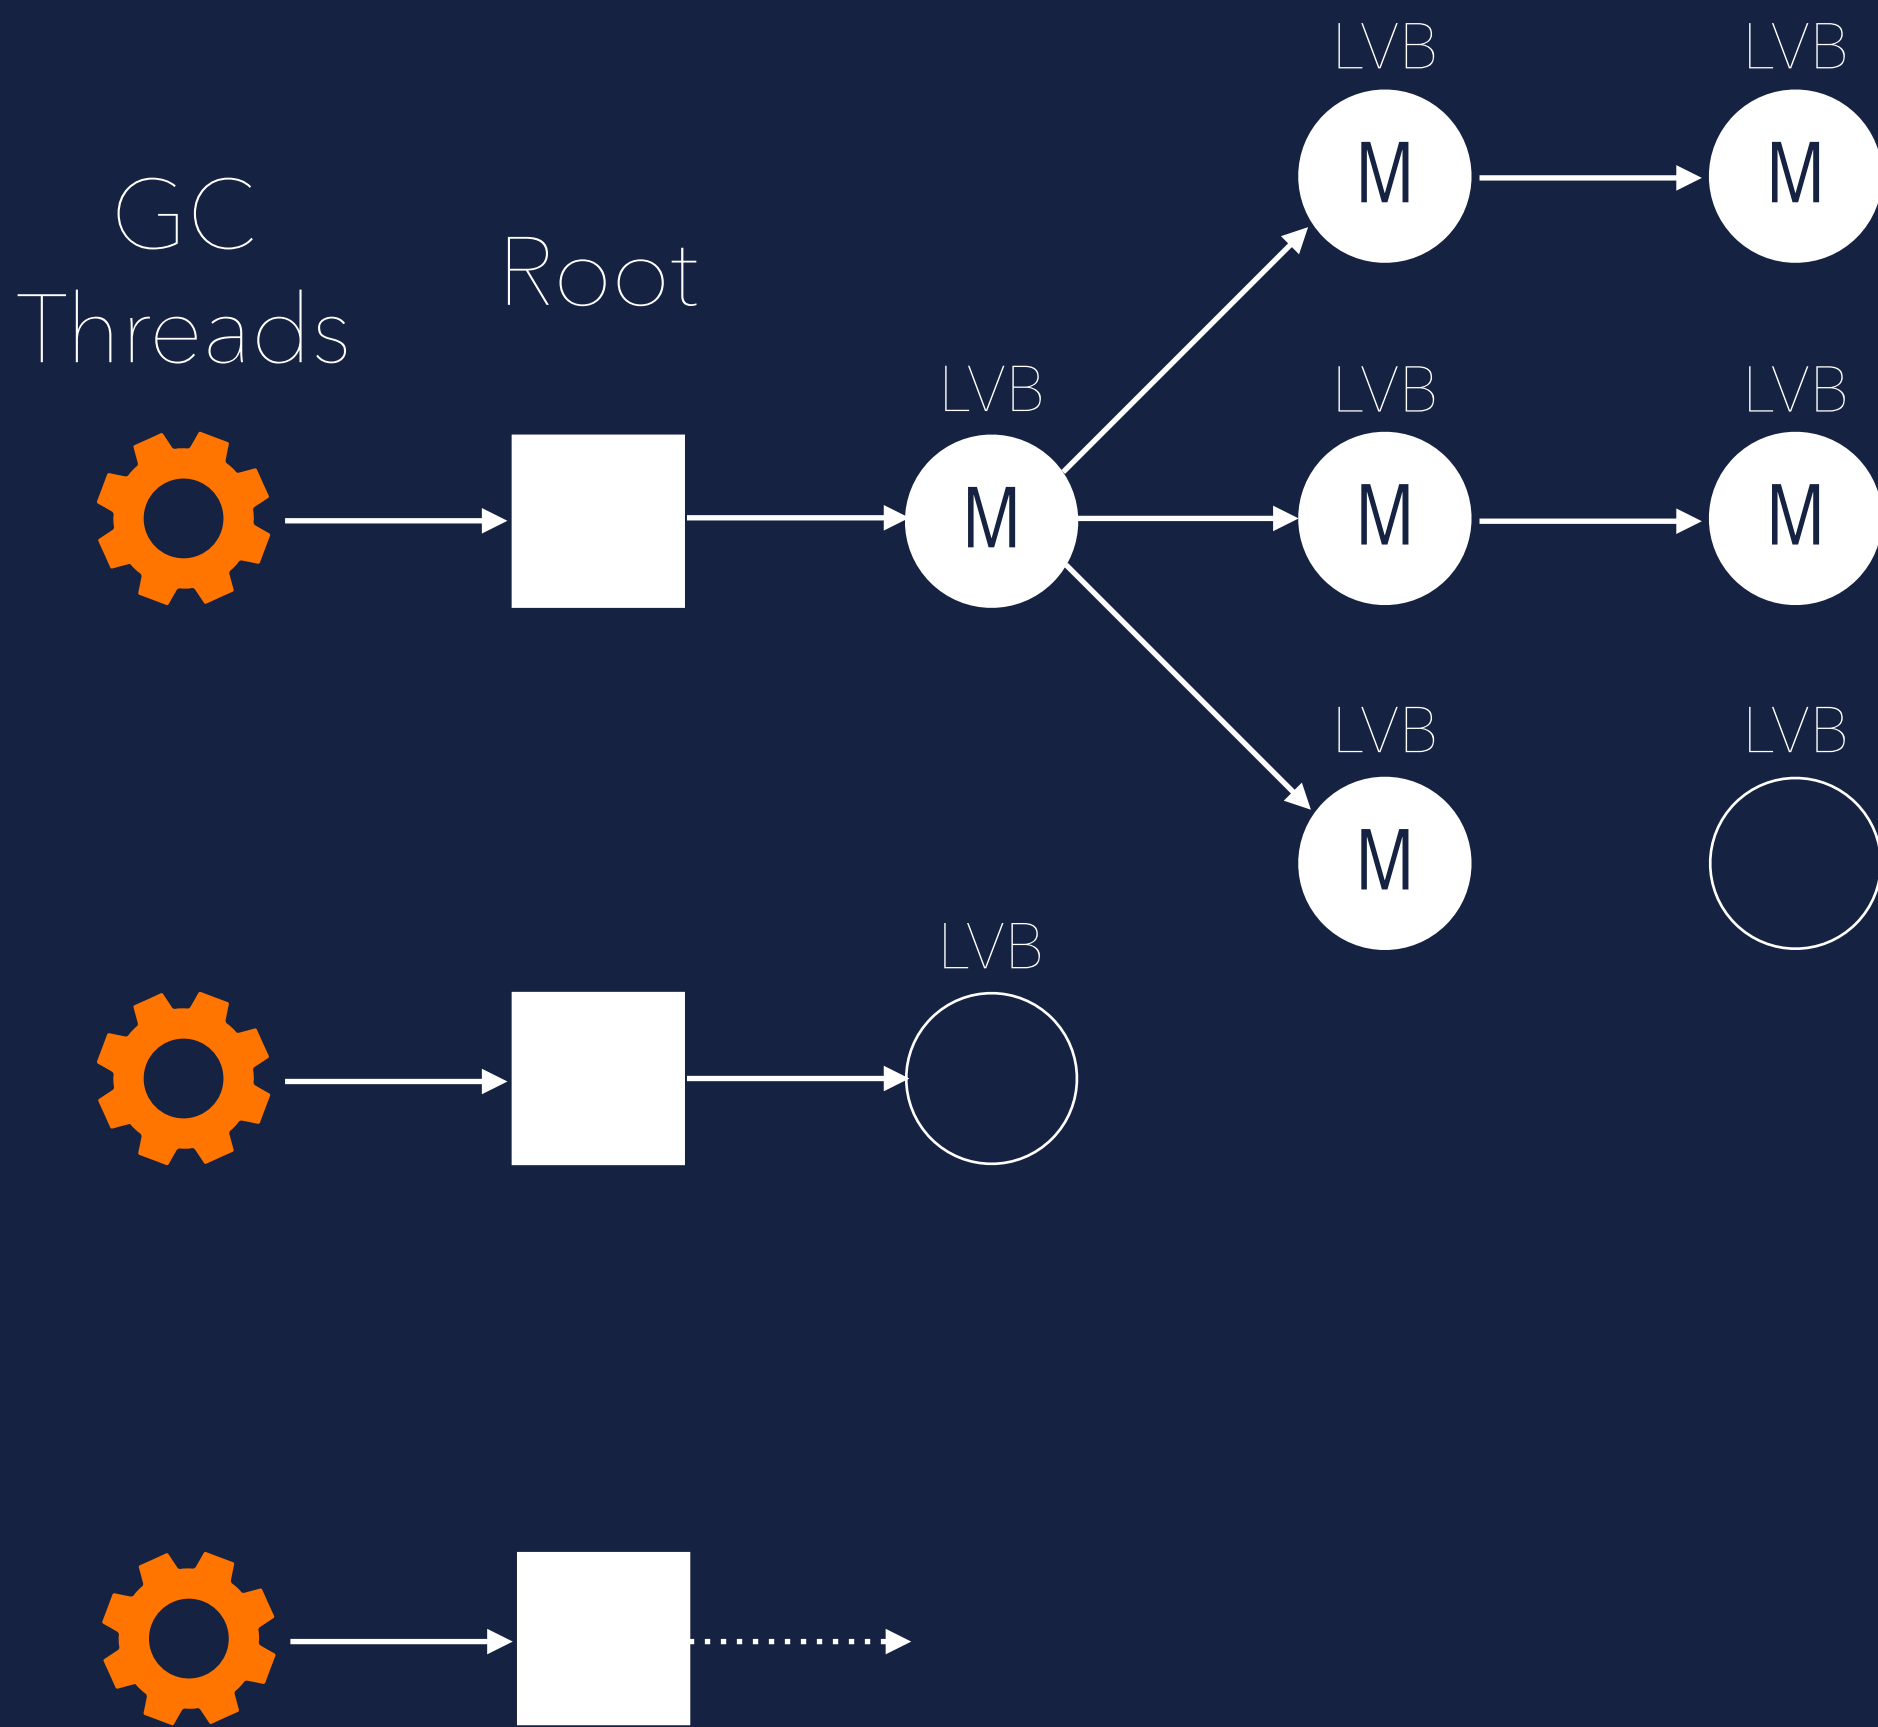





## Marking Phase



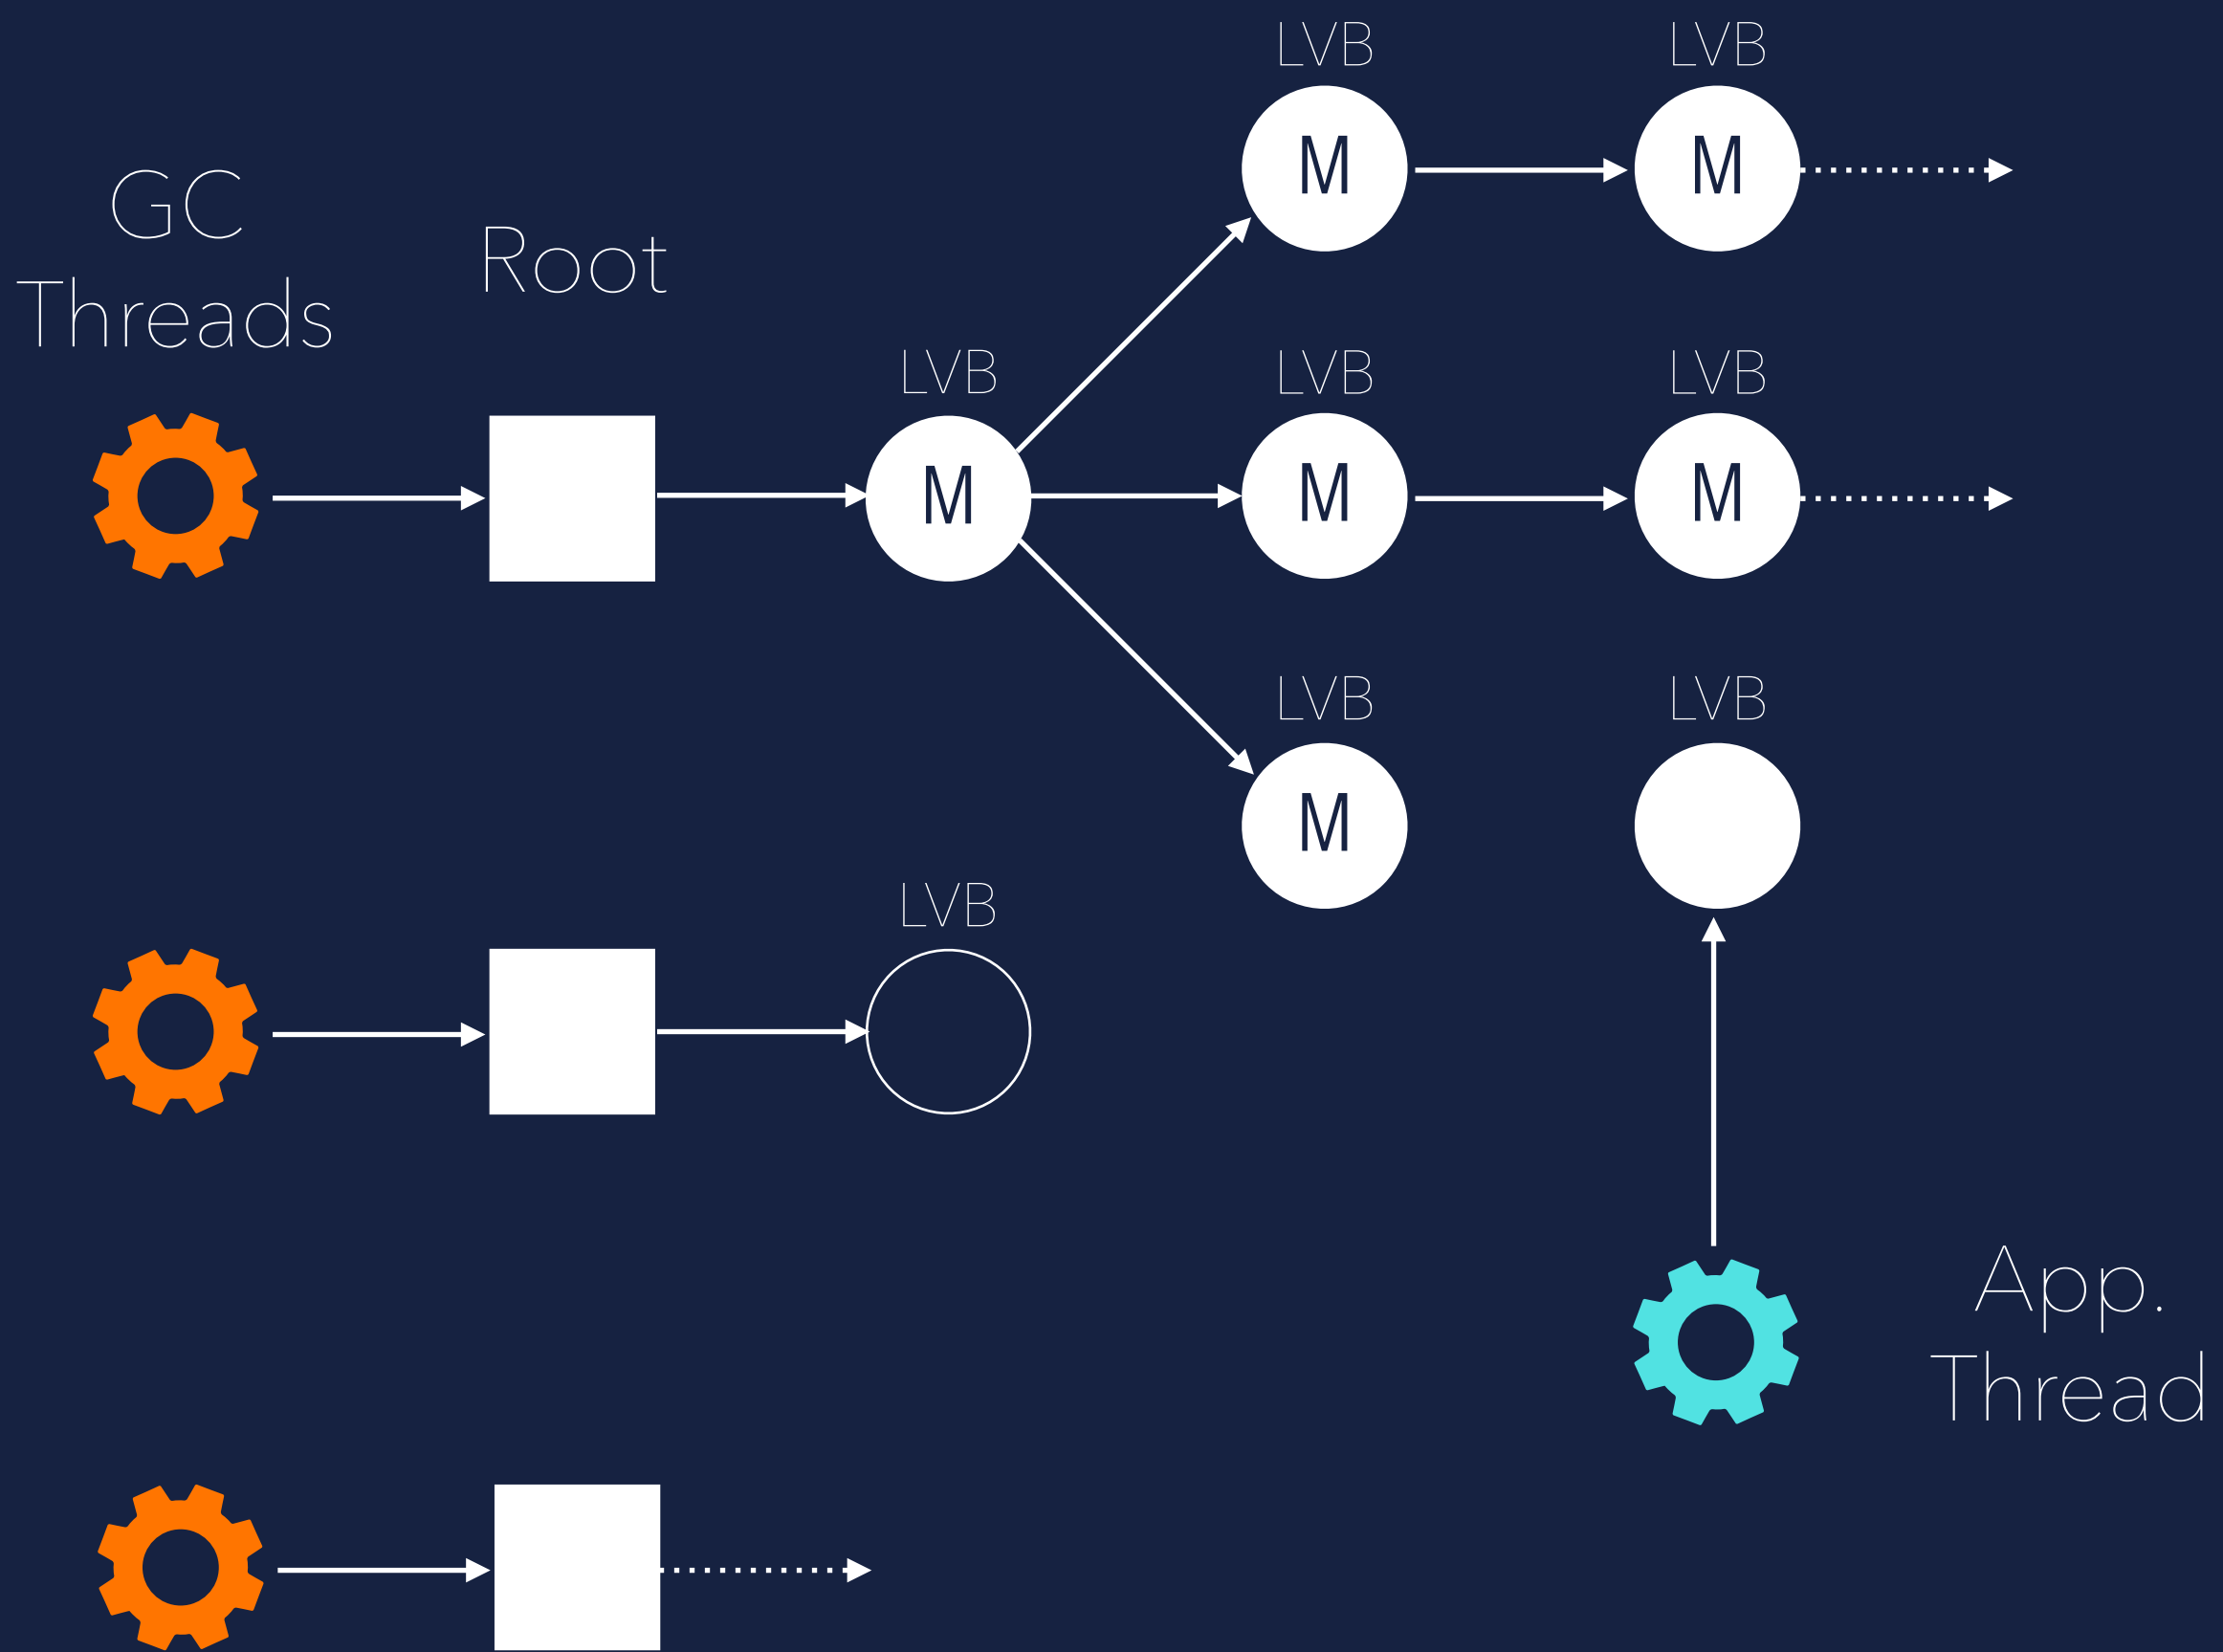
## Marking Phase



## Marking Phase




# Marking Phase




Trigger LVB

# Marking Phase



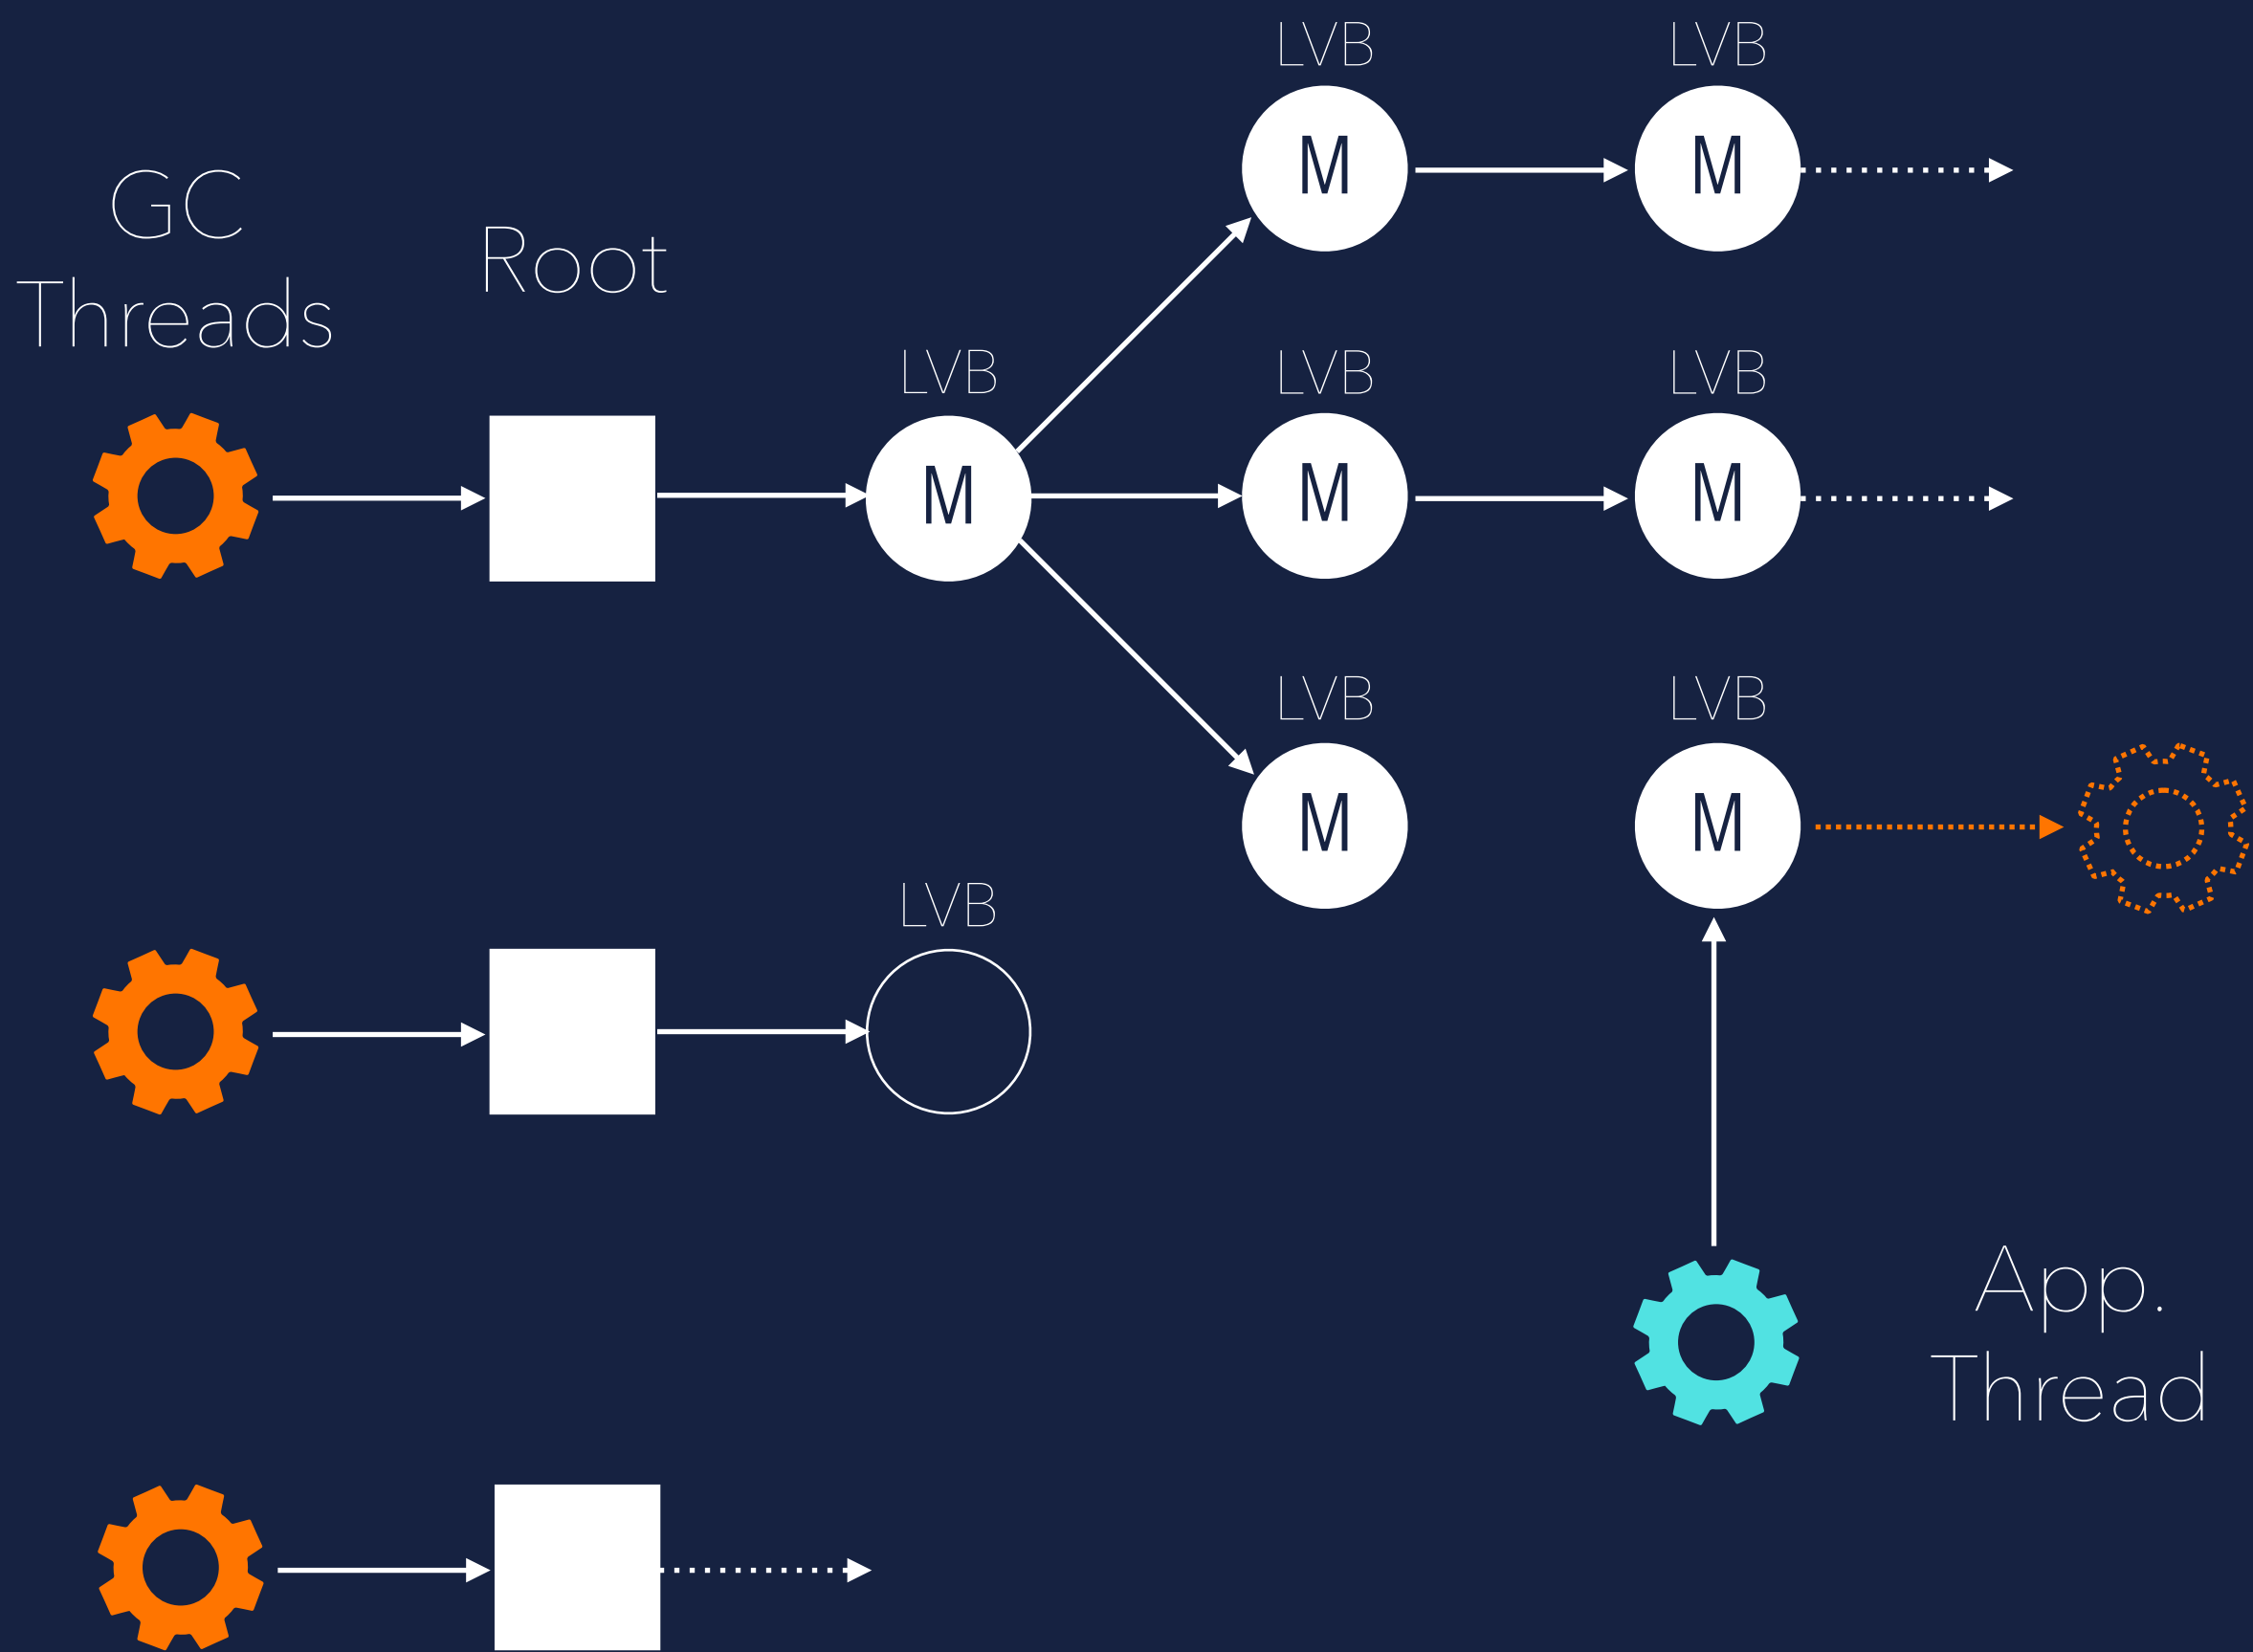
Test+Jump

## Marking Phase




Mark

# Marking Phase



Hand over  
to GC

# Marking Phase



No need to mark again by GC !





C4




# QUICK RELEASE

# Remapping phase with quick release



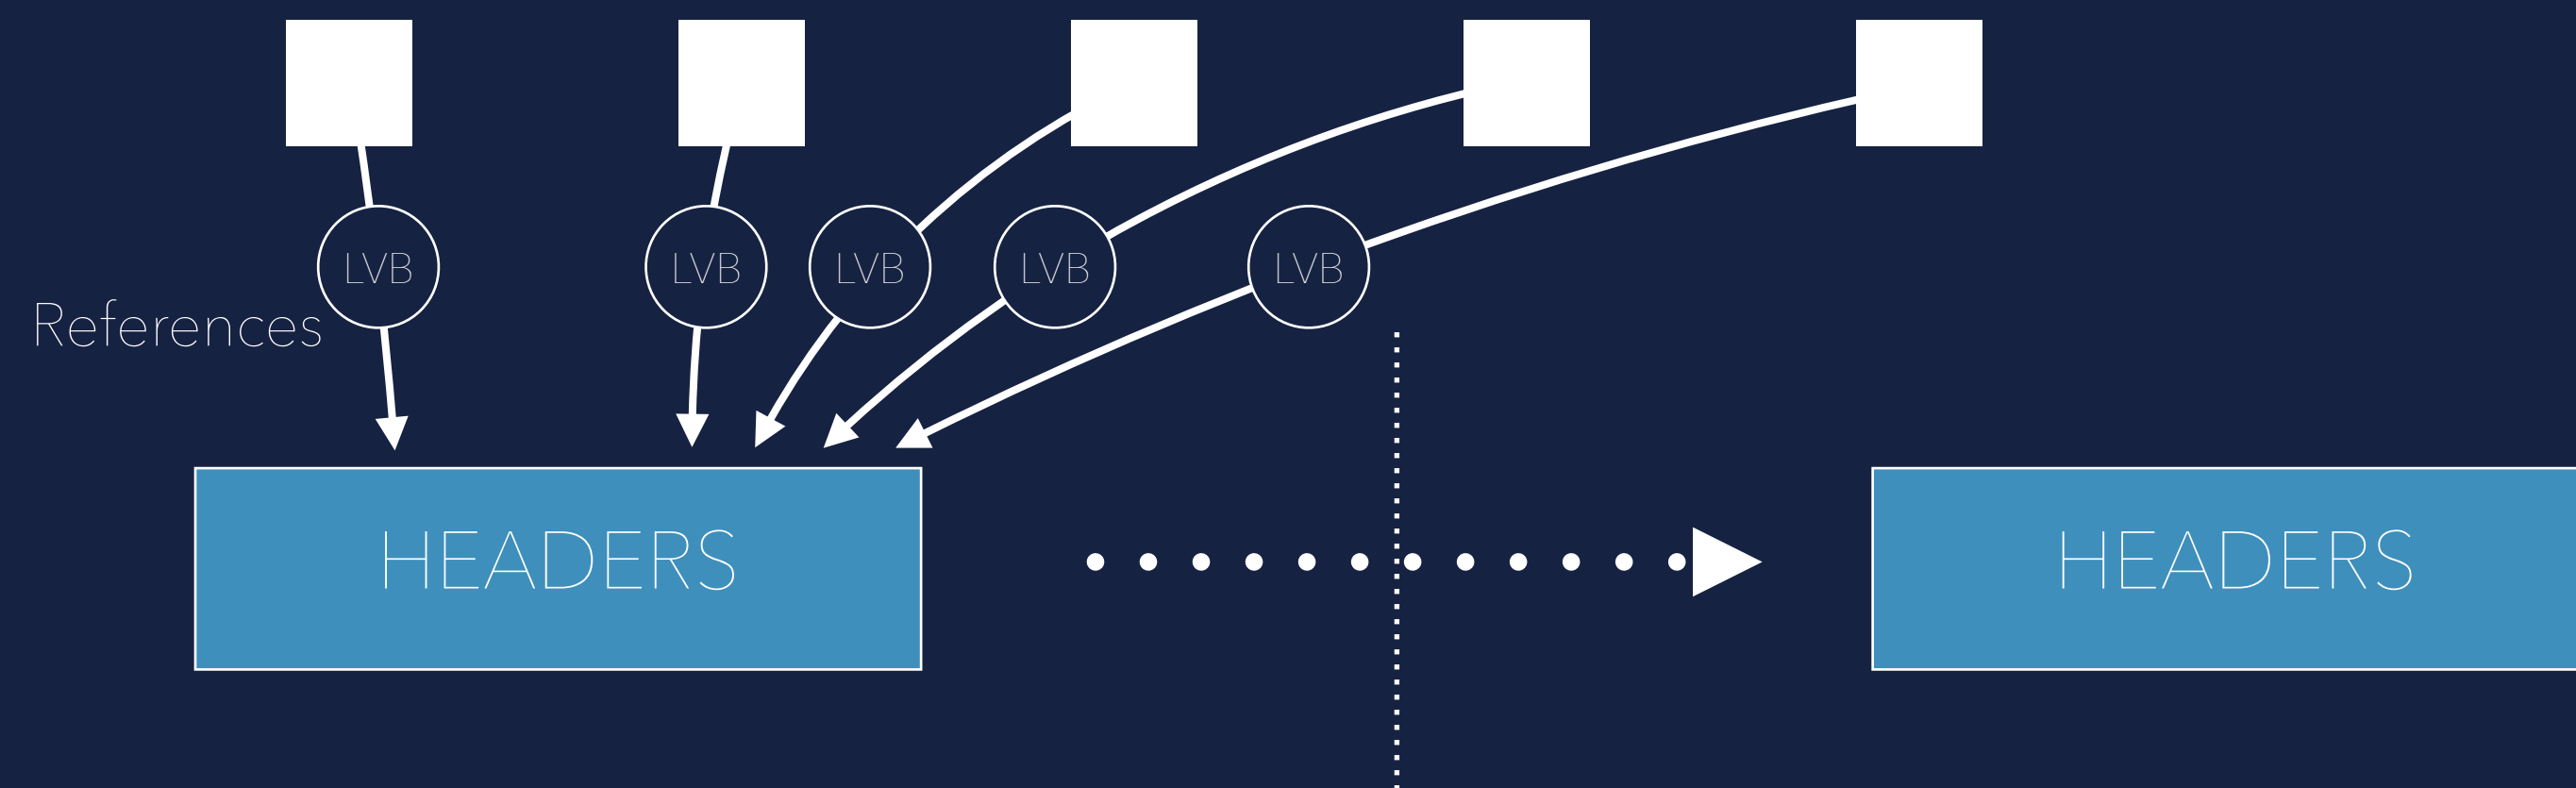
# Remapping phase with quick release

Other GC



Copy the Object


C4



Copy the Object

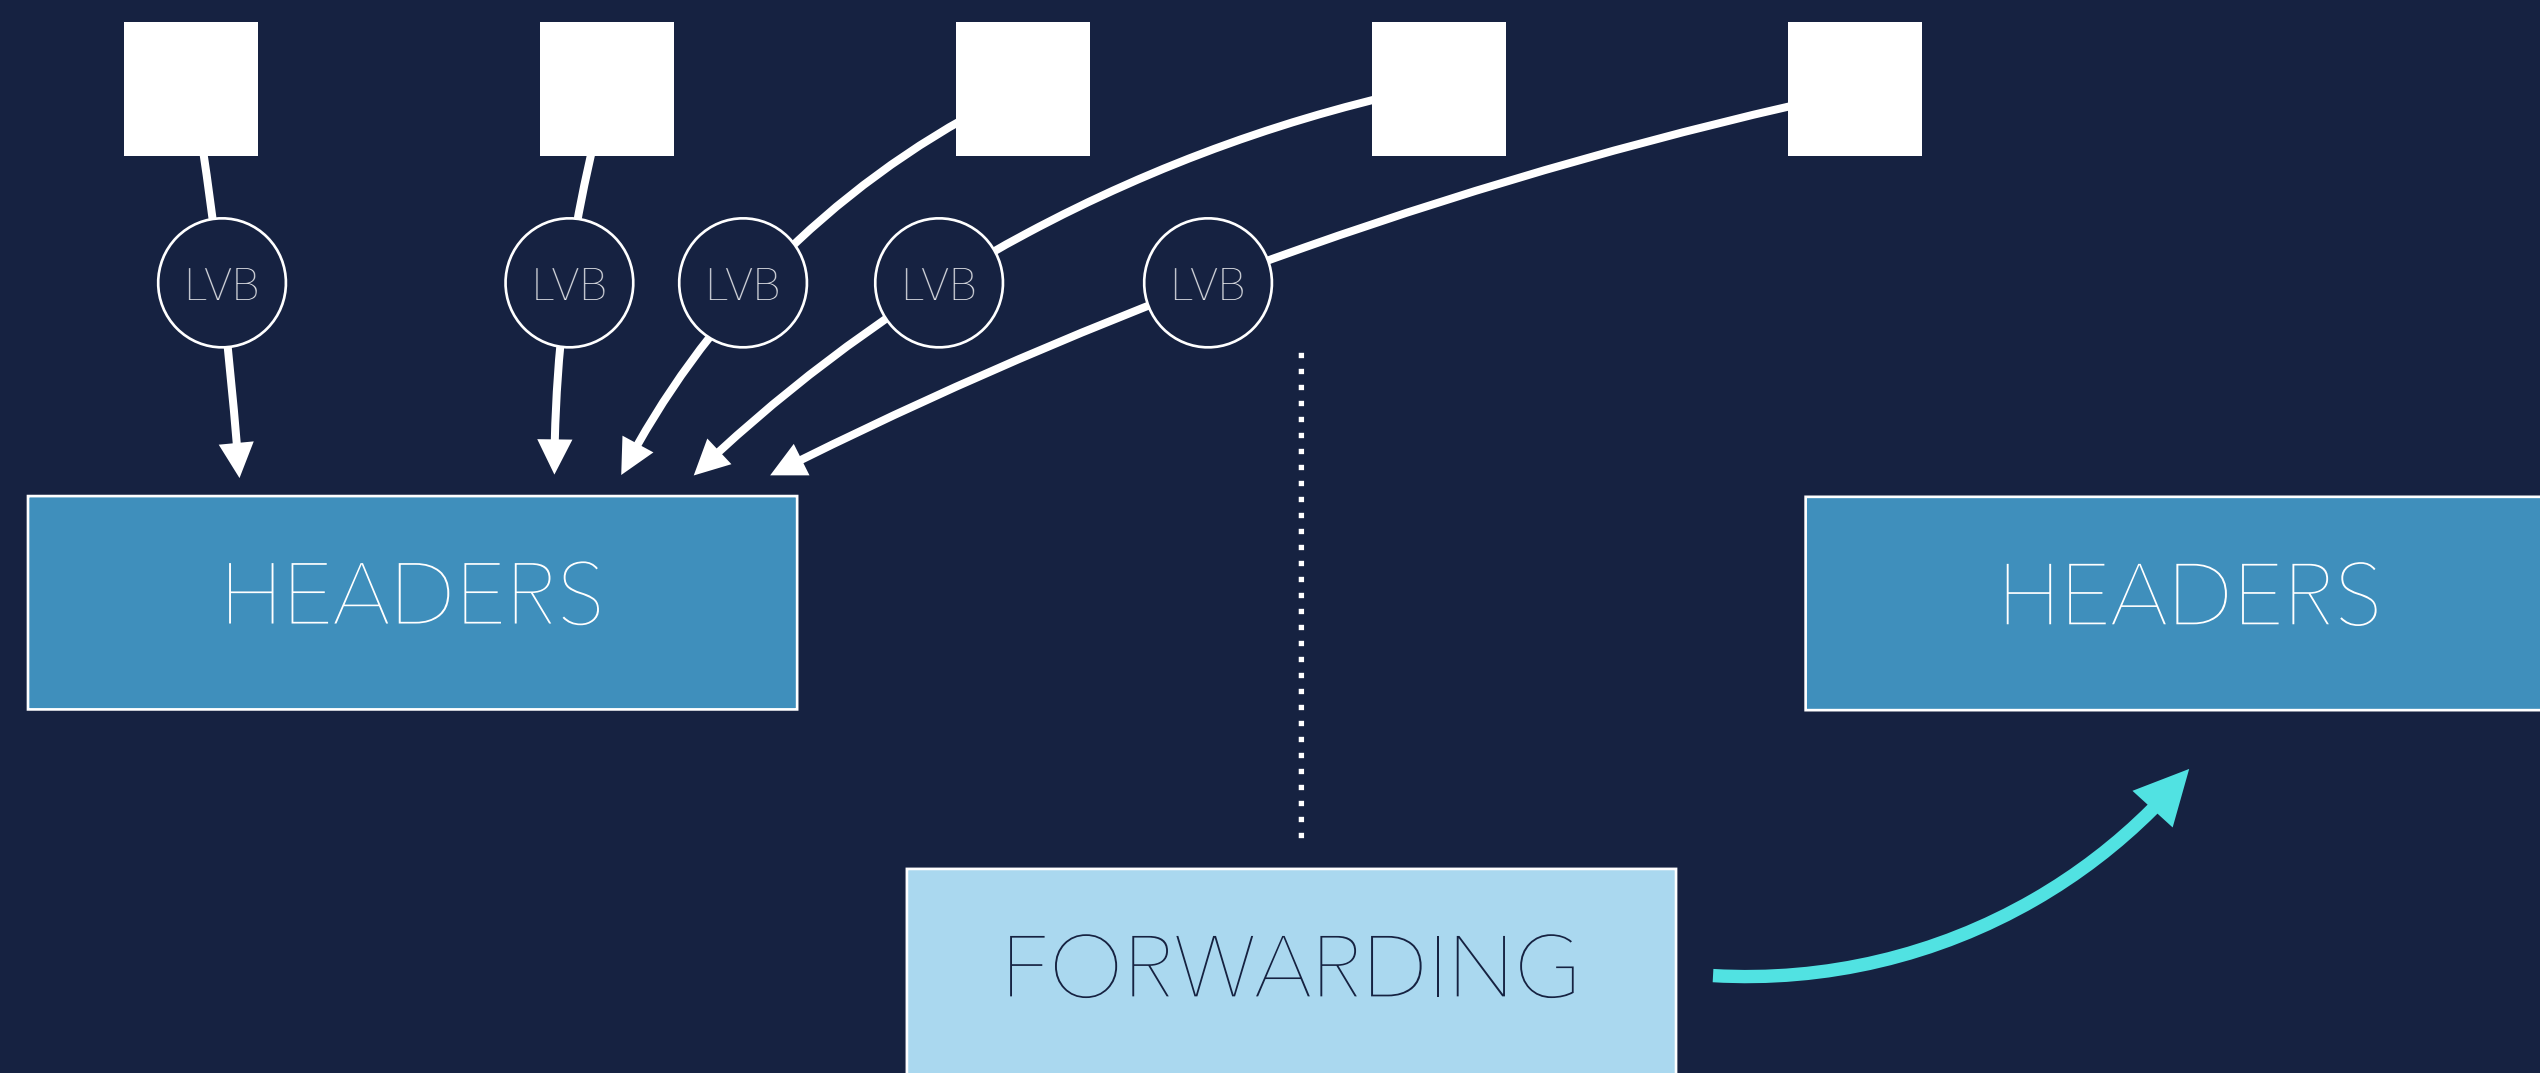
# Remapping phase with quick release

Other GC



Save forwarding info  
to Object header


C4



Save forwarding info  
to Off Heap page

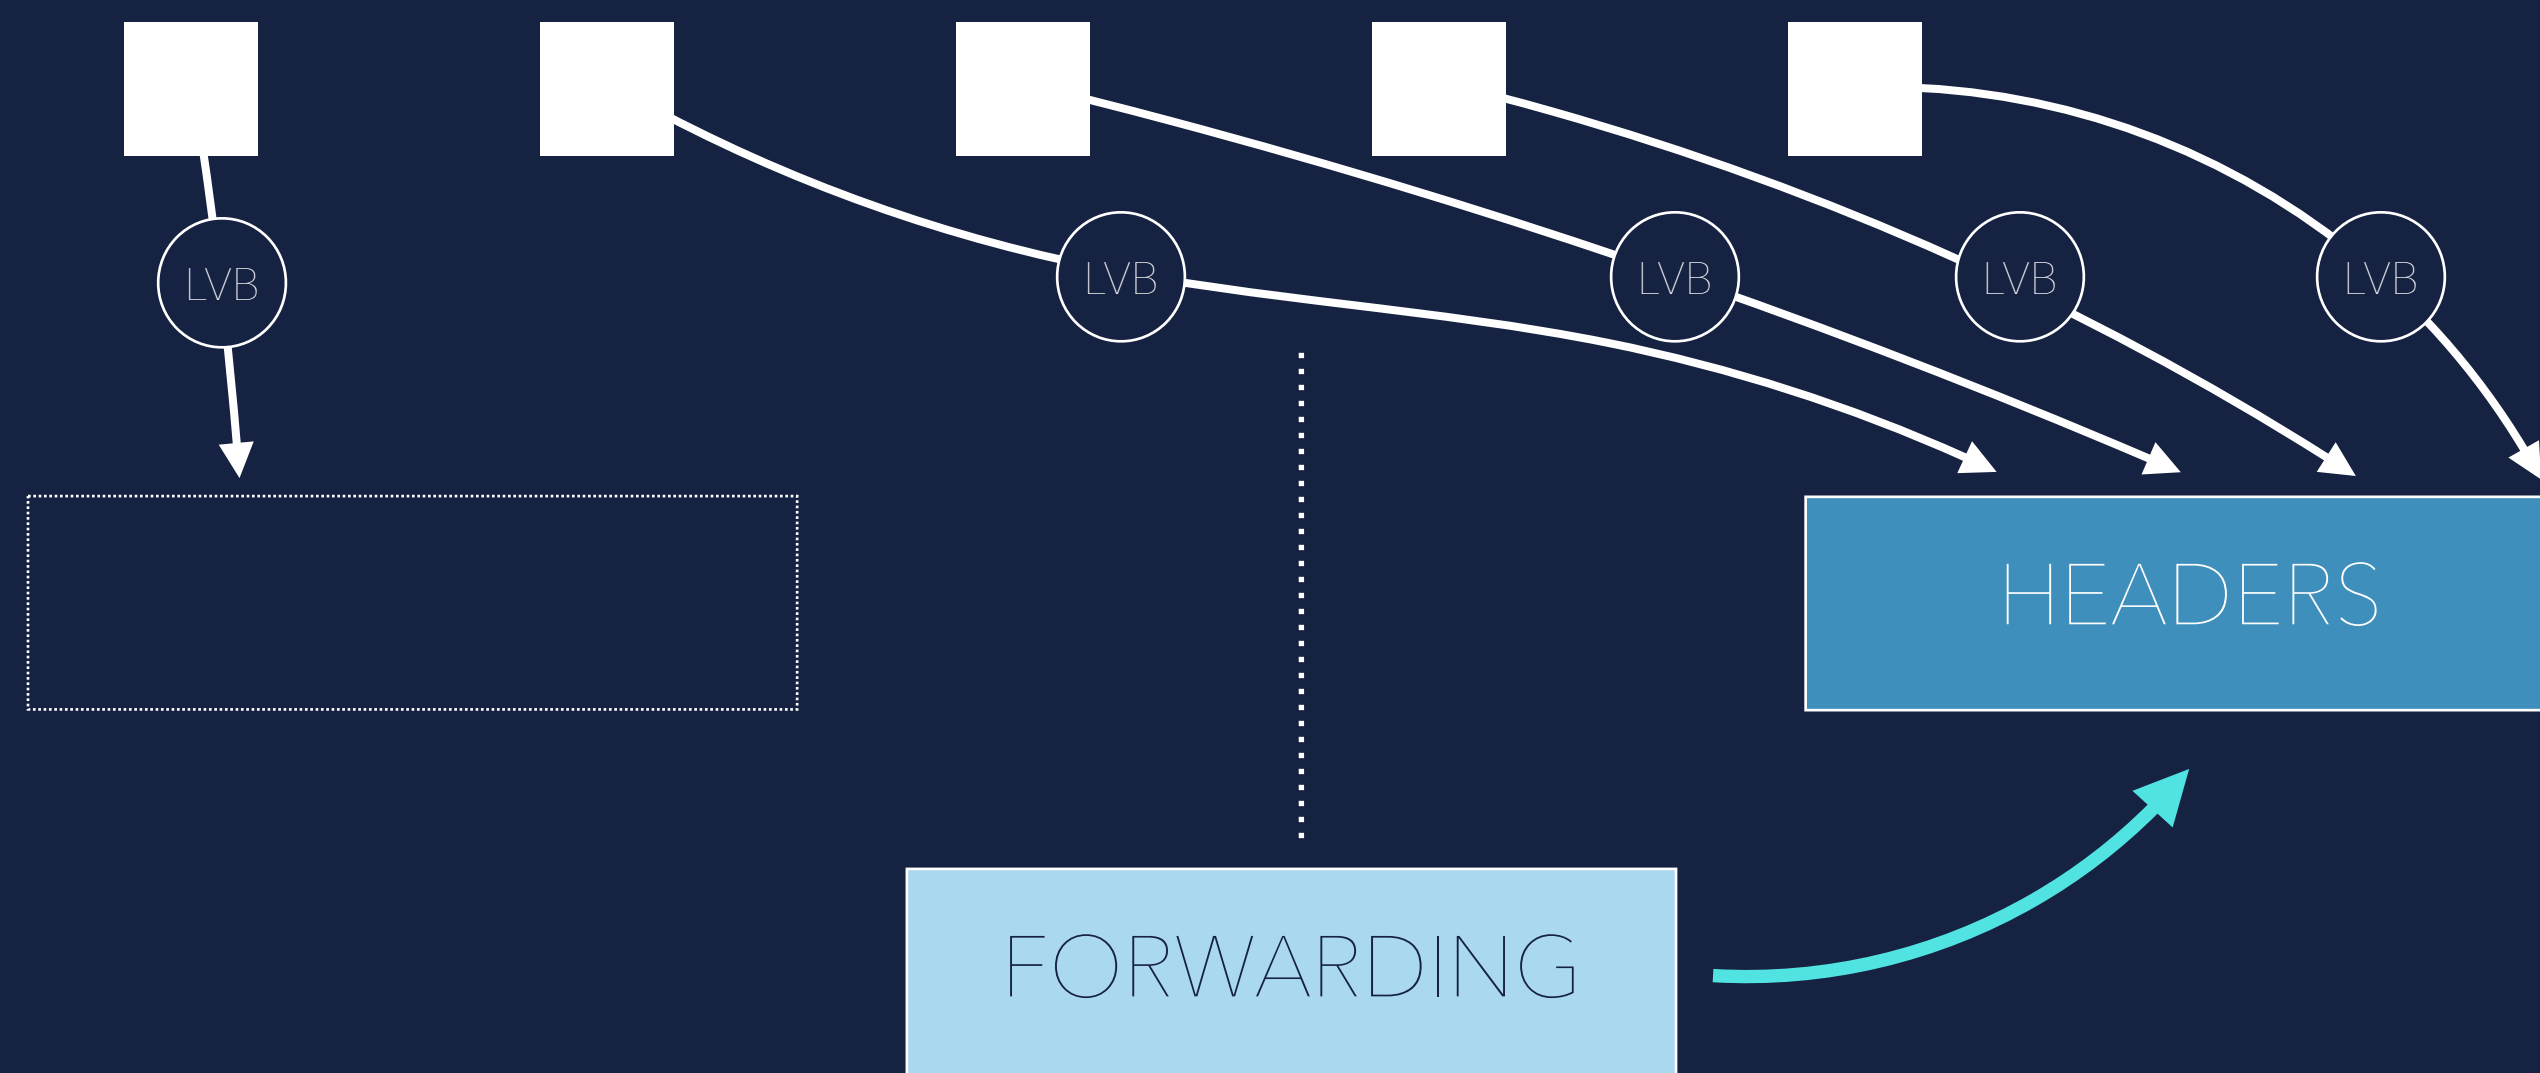
# Remapping phase with quick release

Other GC



Update all references  
(Collector walks the heap and replaces  
all references to new location)


C4



Original reference is freed, and  
Collector updates references  
(No pressure because LVB is self healing)

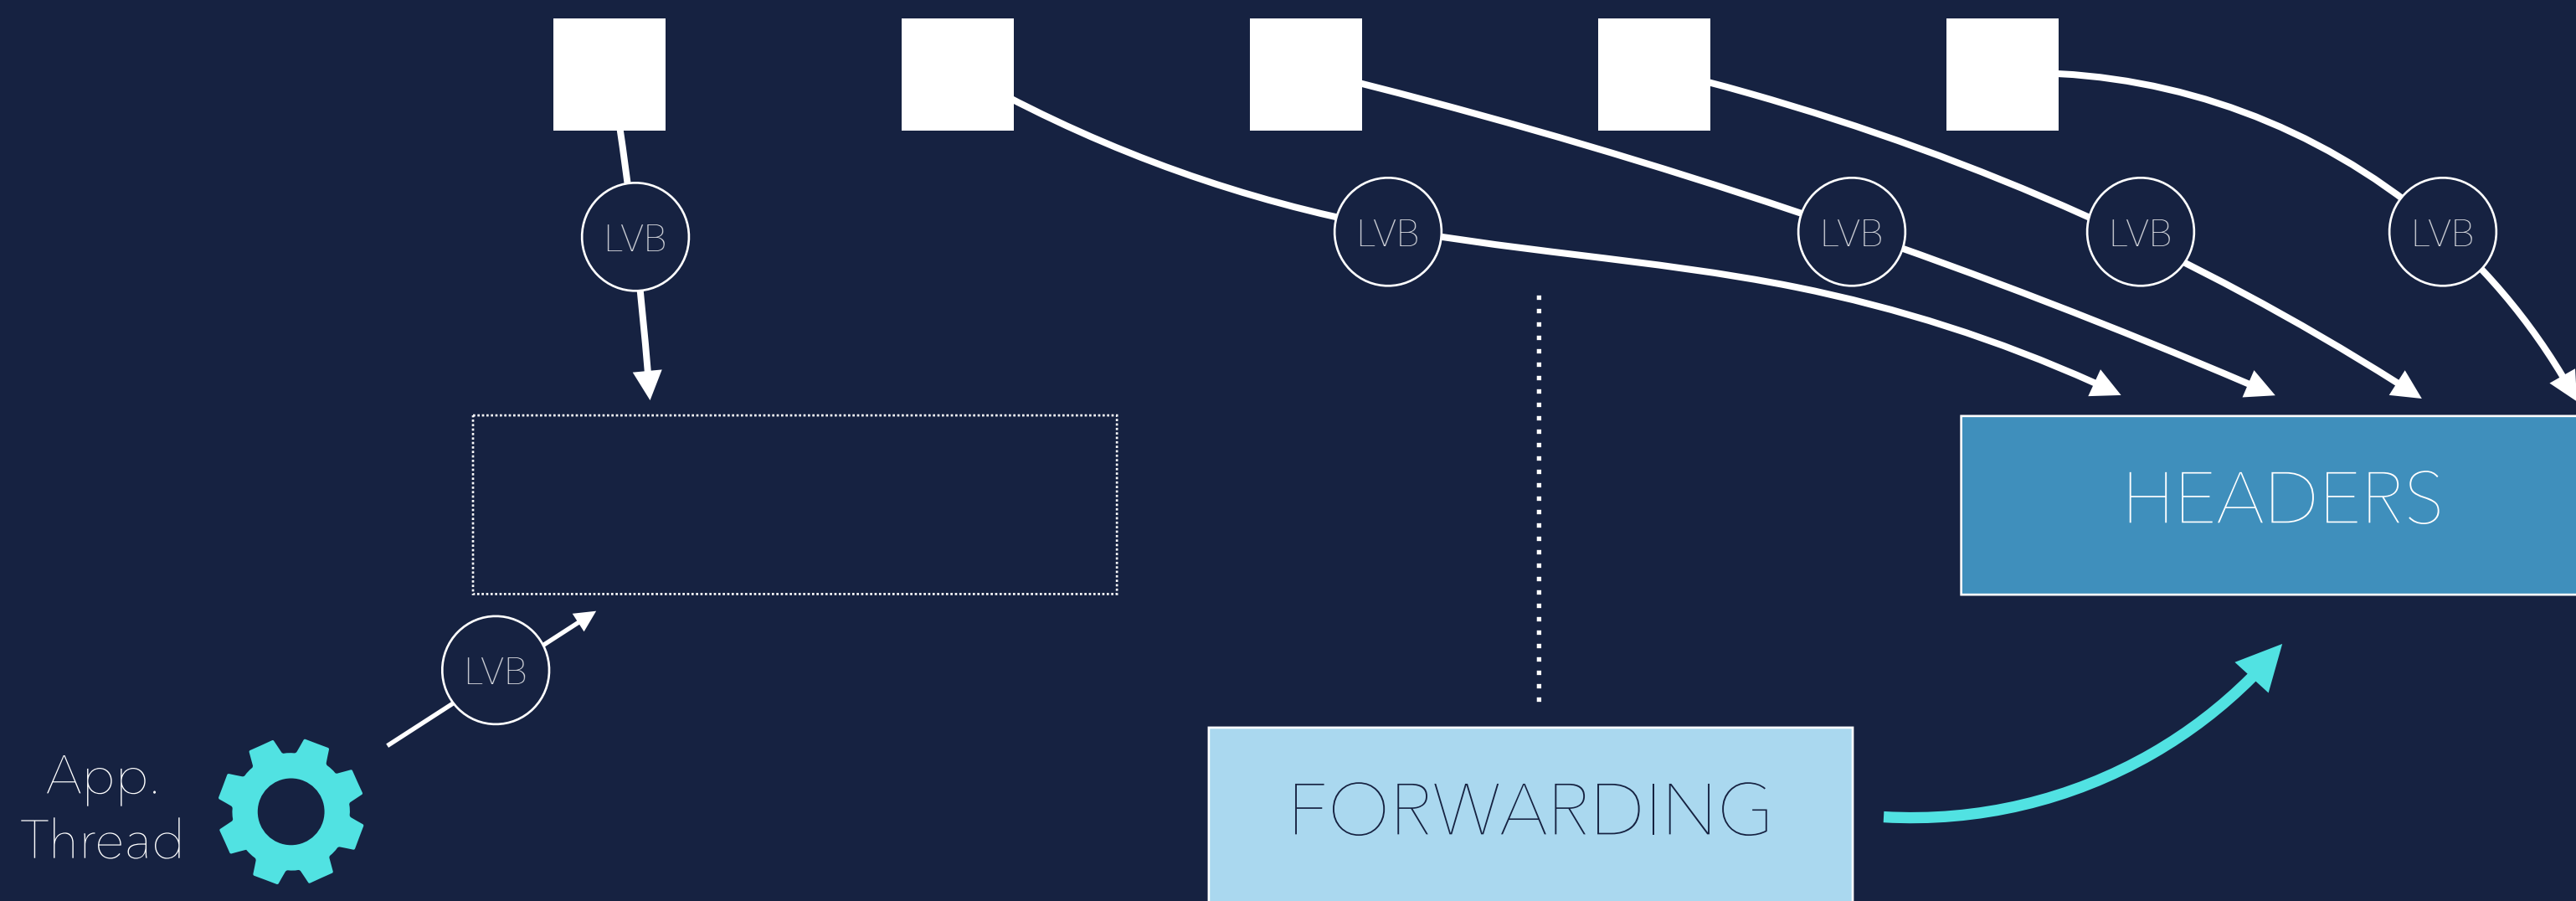
# Remapping phase with quick release

Other GC



Update all references  
(Collector walks the heap and replaces  
all references to new location)


C4



App Thread triggers LVB  
(Test + Jump)

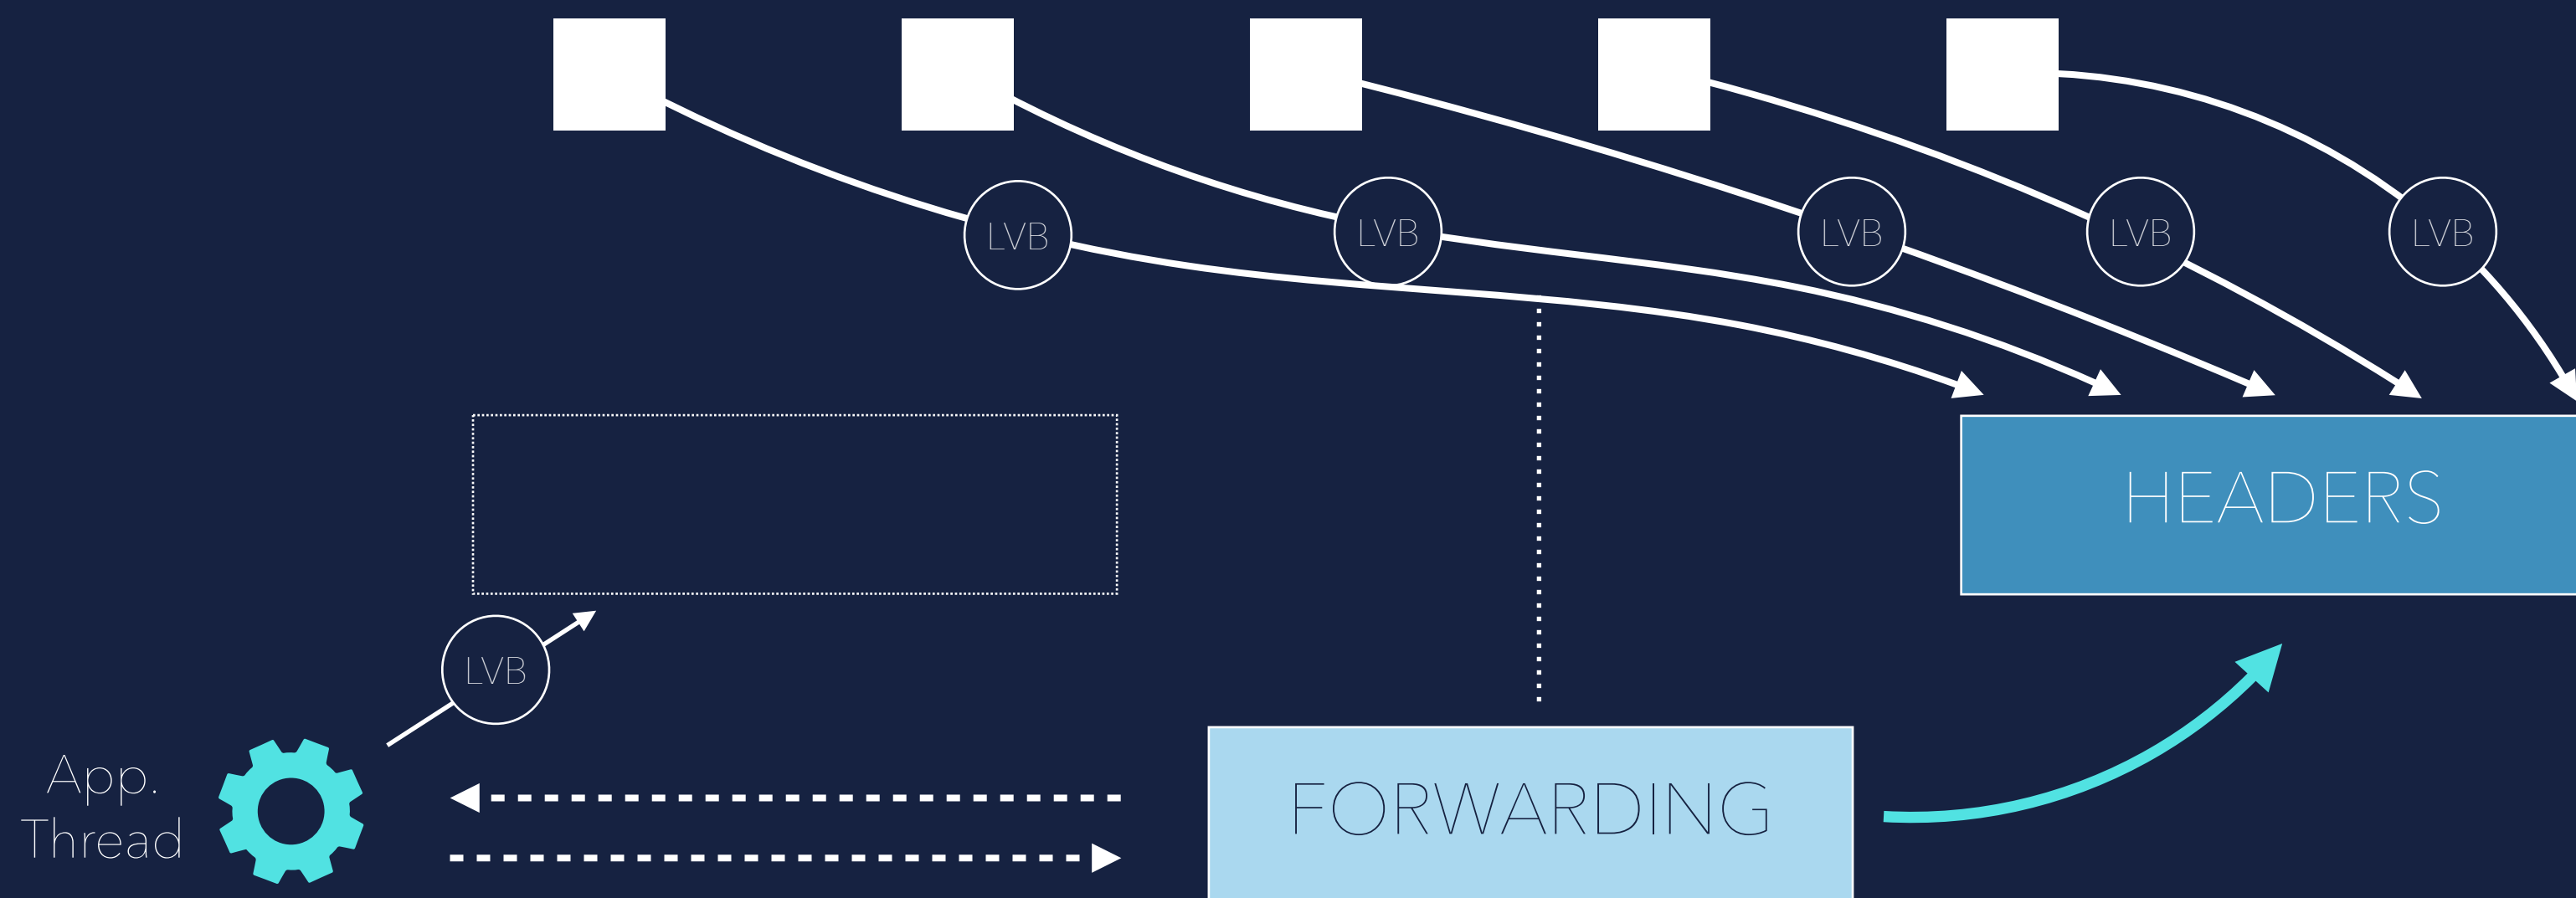
# Remapping phase with quick release

Other GC



Update all references  
(Collector walks the heap and replaces  
all references to new location)


C4



App Thread updates ref  
(Using the info from the off heap page)

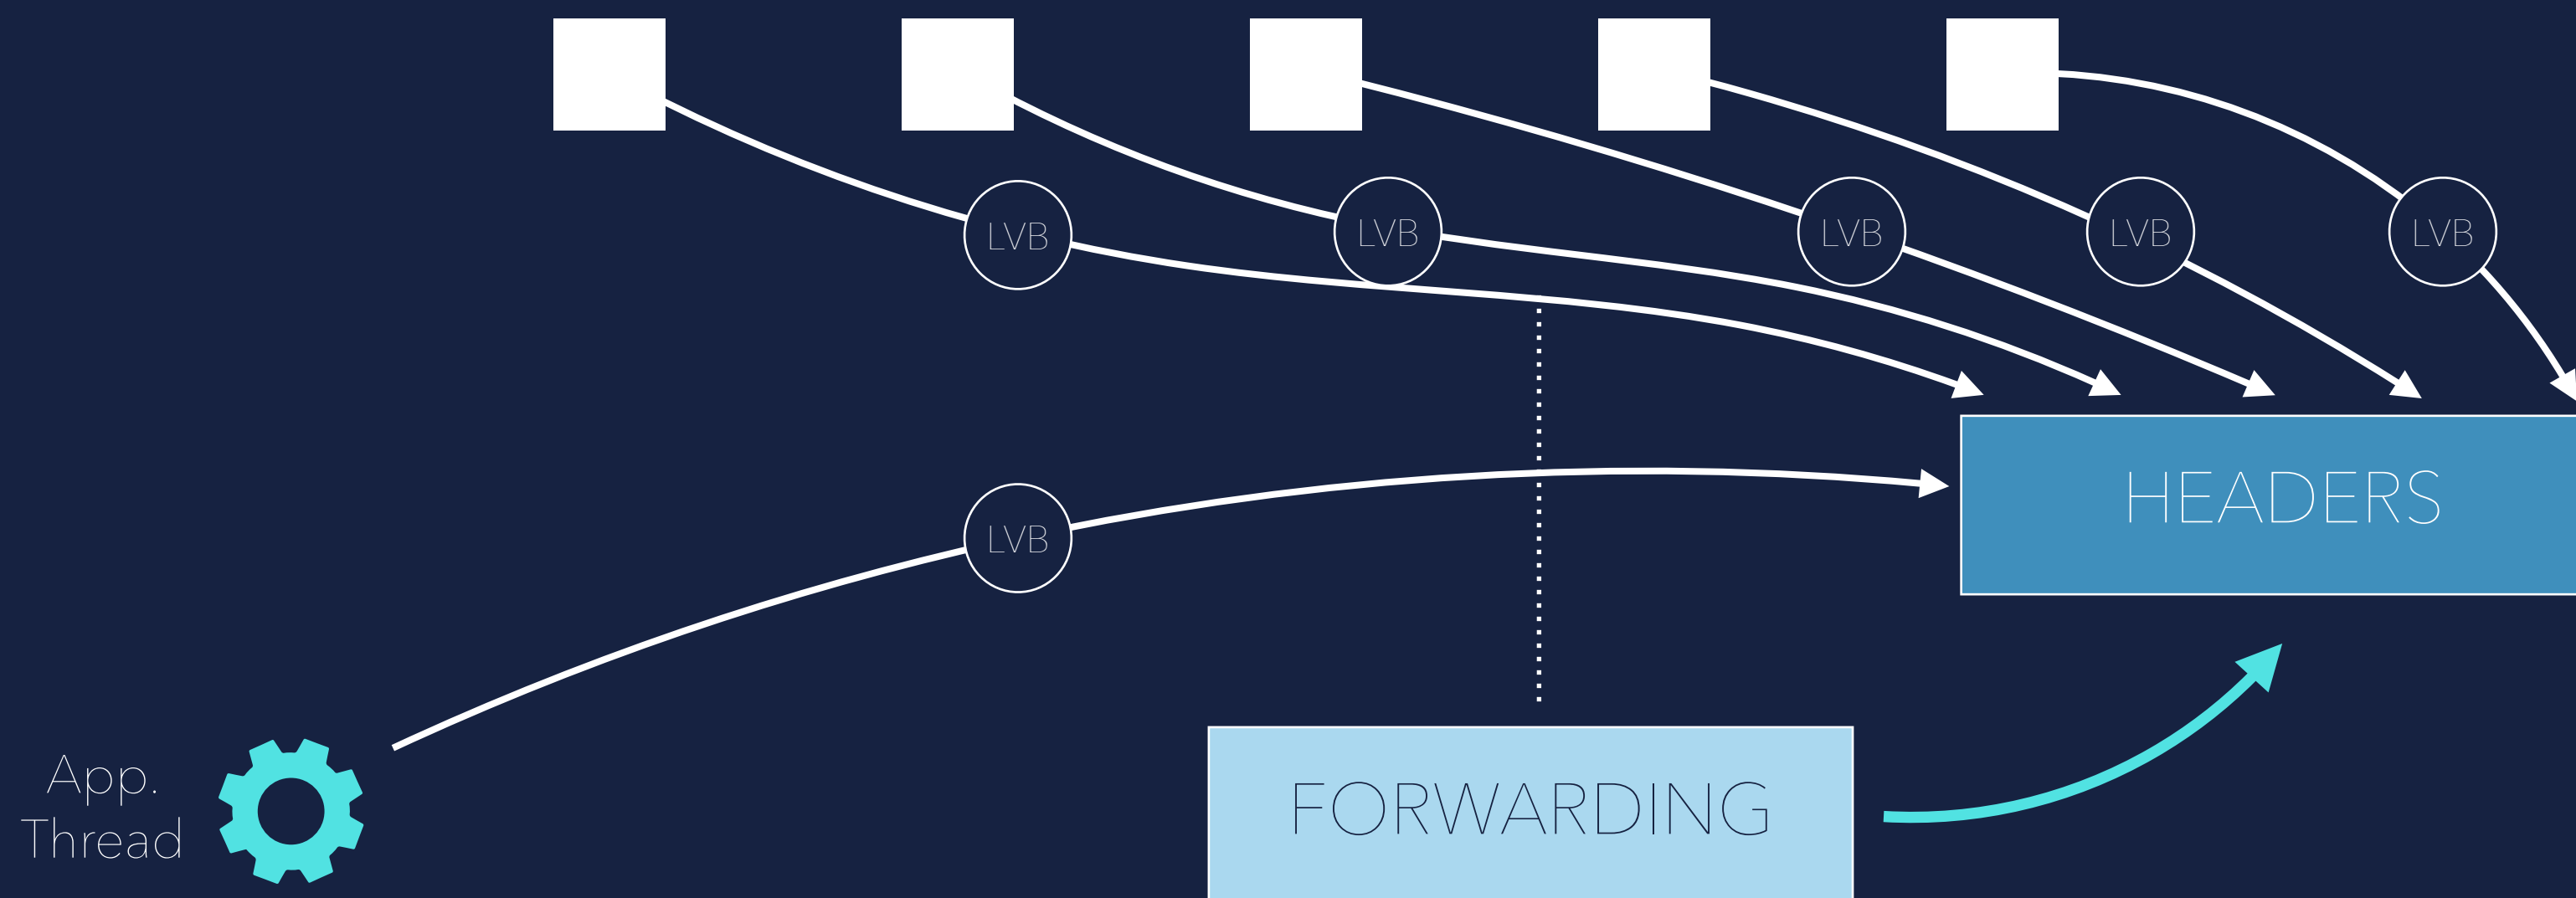
# Remapping phase with quick release

Other GC



Update all references  
(Collector walks the heap and replaces  
all references to new location)

C4




Self healing through LVB



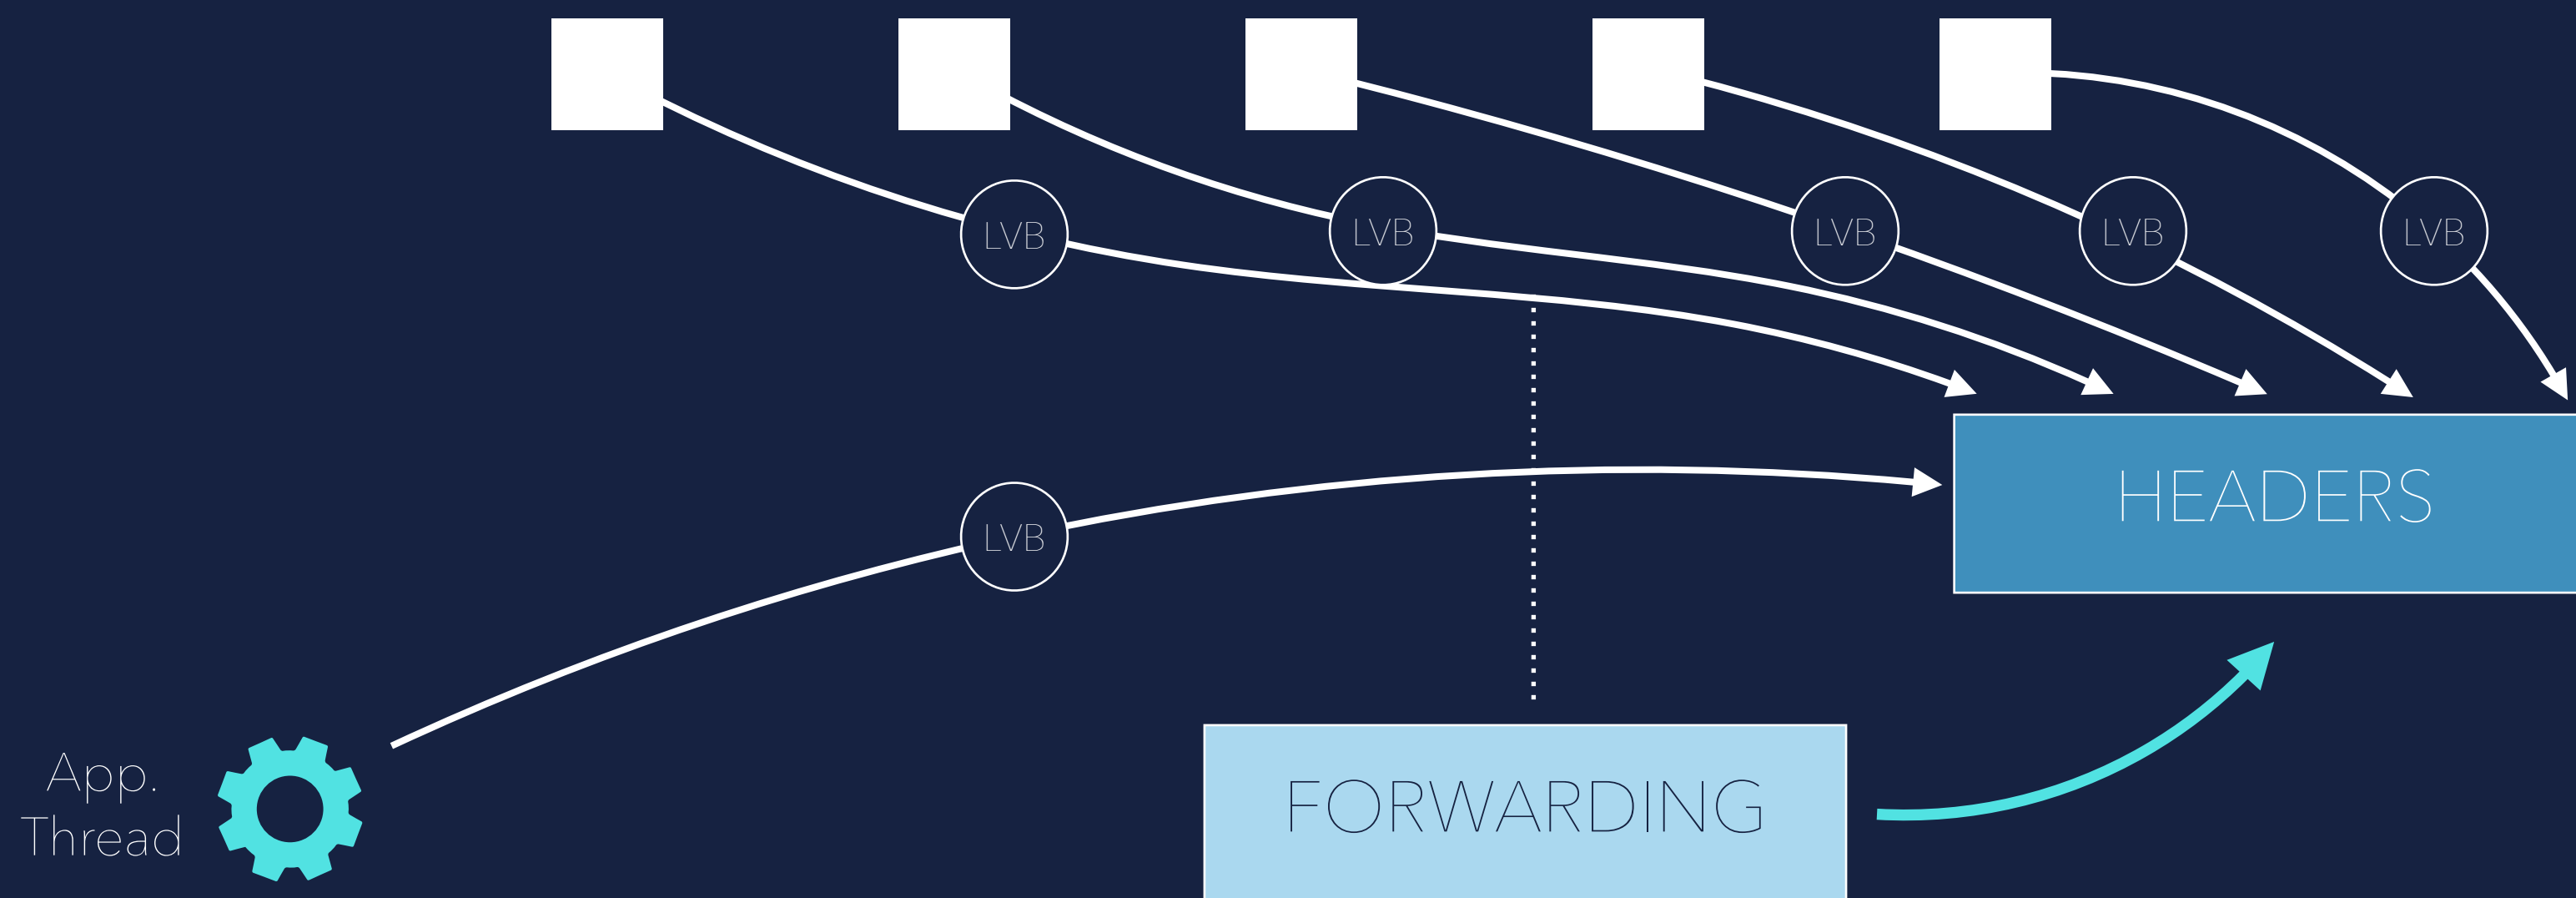
# Remapping phase with quick release

Other GC



Update all references  
(Walk the heap and replace all references  
with forwarding pointer to new location)


C4



Memory can directly be freed  
(ReadBarrier takes care about updating refs.)


# Remapping phase with quick release

Other GC



Remove original object




C4






Off Heap page enables  
Quick Release

AVAILABILITY	AZUL ZING JVM
PARALLEL	YES
CONCURRENT	FULLY
GENERATIONAL	YES
HEAP SIZE	LARGE
PAUSE TIMES	SHORT
THROUGHPUT	VERY HIGH
LATENCY	VERY LOW
CPU OVERHEAD	MODERATE (10-20%)

## CHOOSE WHEN

-  Response time is a high priority
-  Using a very large heap (100GB+)
-  Predictable response times needed

## BEST SUITED FOR

-  Low latency sensitive applications
-  Large scale systems
-  Highly concurrent applications

OS SUPPORT			
------------	---	--	--

JVM SWITCH	> -
------------	-----

## NOTES

-  Only available in Azul Zing JVM
-  No performance overhead because of faster Falcon compiler

**WHICH ONE...?**

# WHICH ONE...?

## Essential Criteria



### Throughput

Percentage of total time spent in application vs. memory allocation and garbage collection

# WHICH ONE...?

## Essential Criteria



### Throughput

Percentage of total time spent in application vs. memory allocation and garbage collection



### Latency

Application responsiveness, affected by gc pauses

# WHICH ONE...?

## Essential Criteria



### Throughput

Percentage of total time spent in application vs. memory allocation and garbage collection



### Latency

Application responsiveness, affected by gc pauses



### Resource usage

The working set of a process, measured in pages and cache lines



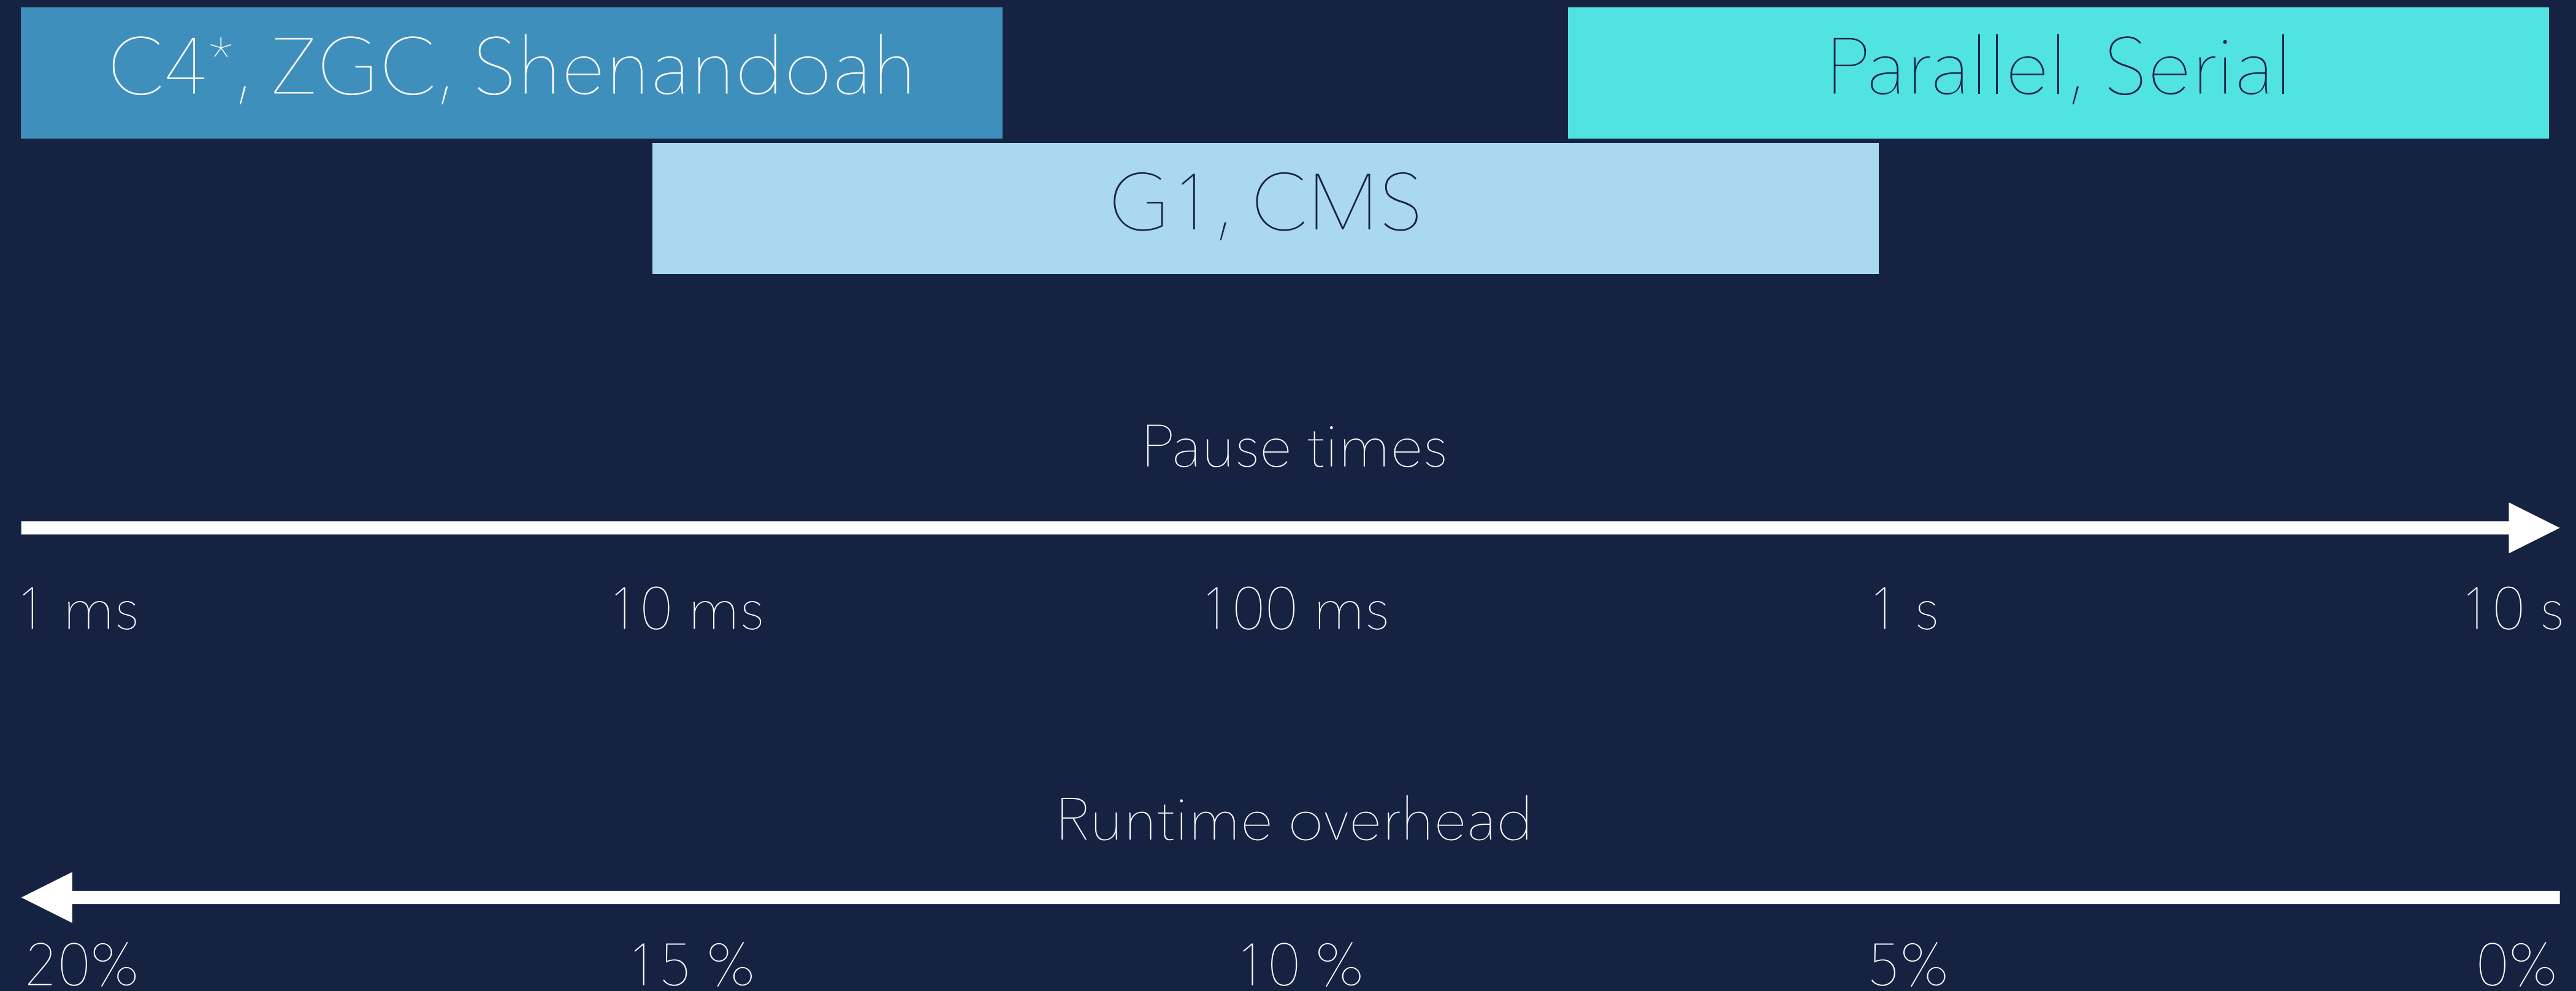
# WHICH ONE...?

## Essential Criteria



# WHICH ONE...?









Choose dependent on your workload



\* C4 has less overhead due to faster Falcon compiler

# OVERVIEW


# OVERVIEW

	Serial GC	Parallel GC	CMS GC	G1	Epsilon	Shenandoah	ZGC	C4
								
Availability	ALL JDK's	ALL JDK's	JDK 1.4-13	JDK 7u4+	JDK 11+	JDK 11.0.9+	JDK15 / 21+	Azul Prime
Parallel	NO	YES	YES	YES		YES	YES	YES
Concurrent	NO	NO	PARTIALLY	PARTIALLY		FULLY	FULLY	FULLY
Generational	YES	YES	YES	YES		NO	NO / YES	YES
Heap Size	SMALL - MEDIUM	MEDIUM - LARGE	MEDIUM - LARGE	MEDIUM - LARGE		LARGE	VERY LARGE	VERY LARGE
Pause Times	LONGER	MODERATE	MODERATE	SHORT - MEDIUM		VERY SHORT (<10ms)	VERY SHORT (<1ms)	VERY SHORT (<1ms)
Throughput	LOW	HIGH	MODERATE	HIGH		VERY HIGH	VERY HIGH	VERY HIGH
Latency	HIGHER	LOWER	MODERATE	LOWER		VERY LOW	VERY LOW	VERY LOW
Performance	LOWER	HIGHER	MODERATE	HIGHER	VERY HIGH	VERY HIGH	VERY HIGH	VERY HIGH
CPU Overhead	LOW	LOWER	MODERATE	MODERATE	VERY LOW	LOW - MODERATE	LOW - MODERATE	LOW - MODERATE
Tail latency	HIGH	HIGH	HIGH	HIGH		MODERATE	LOW	LOW

**WANNNA KNOW  
MORE ?**

# WANNA KNOW MORE ?

R. Jones et al. "The Garbage Collection Handbook". Chapman & Hall/CRC, 2012





# THANK YOU

